



DIOGO HENRIQUE SILVESTRE MATOS CEBOLA

SENTRIS: A SEARCHABLE ENCRYPTED TRIPLESTORE SHARING PROTOCOL

A PRACTICAL AND SECURE SOLUTION FOR THE COLLABORATIVE
DEVELOPMENT, AND DISTRIBUTION, OF SEMANTIC DATA
KNOWLEDGE BASES

MASTER IN COMPUTER SCIENCE

NOVA University Lisbon

Draft: September 25, 2024

SENTRIS: A SEARCHABLE ENCRYPTED TRIPLESTORE SHARING PROTOCOL

A PRACTICAL AND SECURE SOLUTION FOR THE COLLABORATIVE
DEVELOPMENT, AND DISTRIBUTION, OF SEMANTIC DATA KNOWLEDGE BASES

DIOGO HENRIQUE SILVESTRE MATOS CEBOLA

Advisers: Matthias Knorr
Associate Professor, NOVA University Lisbon
João Leitão
Associate Professor, NOVA University Lisbon
Henrique Domingos
Associate Professor, NOVA University Lisbon

ABSTRACT

Ontologies provide a shared understanding of a domain. They are often used in the medical field, alongside electronic health records and clinical decision systems. These systems heavily manipulate personal data, which is sensitive.

Additionally, due to the investment of time and money in their development, or due to the nature of their data, ontologies themselves can have an inherent private, and monetary, value, and need to be secured.

With ontologies gradually growing in dimension, the industry is currently evolving towards the use of collaborative solutions and the outsourcing of heavy computations, relying on cloud infrastructures or other shared computing platforms. These approaches, albeit inevitable, present new security concerns.

In this thesis, we investigate existing secure computation techniques, encryption protocols and semantic data storage systems with the finality of developing a practical and secure solution for the collaborative development and distribution of semantic data knowledge bases.

Our main contributions from this research are: (i) SPARQL-EVAL and SPARQL-EVAL-II, two [SPARQL Protocol and RDF Query Language \(SPARQL\)](#) query evaluation protocols, in the former we apply query rewriting techniques; (ii) SSE-TP, a triple pattern [Symmetric Searchable Encryption \(SSE\)](#) protocol, where encrypted [Resource Description Language \(RDF\)](#) data remains searchable via triple patterns; (iii) SSE-SPARQL-EVAL, an adaptation of the SPARQL-EVAL-II protocol that works with SSE-TP encrypted data; (iv) [Searchable Encrypted Triplestore Sharing \(SEnTriS\)](#), an adaptation of SSE-SPARQL-EVAL for a [Multiple Writer/Multiple Reader \(M/M\)](#) setting which offers data confidentiality and [Role-Based Access Control \(RBAC\)](#).

The solution has been evaluated using the [Lehigh University Benchmark \(LUBM\)](#). T¹

1) *TODO*: Summarize experimental evaluation results

Keywords: Ontologies, Knowledge Bases, SPARQL, Searchable Encryption, Homomorphic Encryption

RESUMO

As ontologias fornecem conhecimento partilhado referente a um domínio. Fazem-no através da especificação dos conceitos relevantes, e suas relações, num dado domínio.

Estas são frequentemente usadas no campo da medicina, juntamente com sistemas de apoio à decisão clínica e repositórios de dados médicos. Sistemas estes que manipulam dados pessoais e de carácter sensível.

Adicionalmente, devido aos elevados investimentos monetários e laborais que podem estar associados ao seu desenvolvimento, ou devido à natureza dos seus conteúdos, as ontologias podem ter um valor privado, e monetário, associado, necessitando de serem protegidas.

Atualmente, com o gradual aumento da dimensão das ontologia, a indústria está evoluir no sentido da utilização de ferramentas colaborativas e *outsourcing* de computações pesadas, recorrendo a infraestruturas *cloud* e outras plataformas de computação partilhada. Estas soluções, embora inevitáveis, apresentam novas questões de segurança.

Nesta tese, investigamos técnicas existentes de computação seguras, protocolos de encriptação e sistemas de armazenamento de dados semânticos com a finalidade de desenvolver uma solução prática e segura para o desenvolvimento colaborativo e distribuição de fontes de conhecimento semânticas.

As principais contribuições da investigação são: (i) SPARQL-EVAL e SPARQL-EVAL-II, dois protocolos para a avaliação de queries [SPARQL](#), no segundo utilizamos técnicas de rescrita das queries; (ii) SSE-TP, um protocolo de encriptação simétrica pesquisável de padrões de triplos [RDF](#); (iii) SSE-SPARQL-EVAL, uma adaptação do protocolo SPARQL-EVAL-II que funciona com dados encriptados via o protocolo SSE-TP; (iv) [SEnTriS](#), uma adaptação do protocolo SSE-SPARQL-EVAL para um contexto de multiutilizador, com escritas e leituras concorrentes, e que oferece confidencialidade de dados e [RBAC](#).

O sistema foi avaliado com o *benchmark* [LUBM](#).

Palavras-chave: Ontologias, Encriptação Pesquisável, Computação Segura, SPARQL

CONTENTS

List of Figures	viii
List of Tables	ix
List of Listings	x
Acronyms	xii
1 Introduction	1
1.1 Context	1
1.2 Problem	3
1.3 Contributions	4
1.4 Document Organization	5
2 Related Work	6
2.1 Semantic Web: A layered architecture	6
2.1.1 RDF: Resource Description Framework	7
2.1.1.1 Overview	7
2.1.1.2 Syntax	8
2.1.2 RDFS: RDF Schema	9
2.1.2.1 Overview	9
2.1.2.2 RDFS primitives	9
2.1.3 OWL2: Web Ontology Language	10
2.1.3.1 Overview	10
2.1.3.2 OWL Language	11
2.1.3.3 OWL2 DL	14
2.1.4 SPARQL	15
2.1.4.1 Overview	15
2.1.4.2 SPARQL Queries	15
2.1.4.3 Preserving reasoning	18
2.1.5 Industry solutions	19
2.1.5.1 Development Software & Tools	19
2.1.5.2 Triplestores	19
2.2 Secure Computations	19
2.2.1 Homomorphic Encryption	20
2.2.1.1 Overview	20

2.2.1.2	Types of HE schemes	20
2.2.2	Property-preserving Encryption	21
2.2.2.1	Overview	21
2.2.2.2	Deterministic Encryption	21
2.2.2.3	Order-Preserving Encryption	21
2.2.3	Searchable Encryption	22
2.2.3.1	Overview	22
2.2.3.2	General Index-Based Model	22
2.2.3.3	Client/Server Model	22
2.2.3.4	Information Leakage	23
2.2.4	Symmetric Searchable Encryption	23
2.2.4.1	Overview	23
2.2.4.2	Security Definitions	23
2.2.4.3	SSE Schemes	24
2.3	Secure Semantic Data Solutions	25
2.3.1	Self-Enforcing Access Control for Encrypted RDF	26
2.3.2	CryptGraphDB	26
2.3.3	GOOSE	28
2.3.4	Discussion	29
2.4	Summary	30
3	SEnTriS: A Searchable Encrypted Triplestore Sharing Protocol	31
3.1	Solution Models	31
3.1.1	Overview	31
3.1.2	System Model	31
3.1.3	Security Model	32
3.2	Solution Definition	32
3.2.1	Overview	32
3.2.2	Core elements	33
3.2.3	Cryptographic Tools	36
3.2.3.1	Notation	36
3.2.3.2	DGK Encryption	36
3.2.4	Evaluation of SPARQL queries	38
3.2.4.1	Storage of triple patterns	38
3.2.4.2	SPARQL Query Evaluation Protocol	39
3.2.4.3	SPARQL Query Evaluation Protocol with Reasoning	41
3.2.4.4	Supported Queries	42
3.2.5	Encryption of triple patterns	43
3.2.5.1	Triple Pattern SSE scheme	43
3.2.5.2	Trapdoor generation	44
3.2.5.3	Initial setup & re-initialization of scheme	45
3.2.5.4	Triple Pattern encryption and decryption	46
3.2.5.5	Storage, searches & updates	46
3.2.6	Evaluation of SPARQL queries over encrypted data	47
3.2.6.1	Query Preparation Phase	48
3.2.6.2	Query Evaluation Phase	49
3.2.6.3	Decryption of SPARQL query results	49
3.2.7	The M/M setting	49
3.2.7.1	Authentication Procedure & Session Validation	54

3.2.7.2	Lock-based Concurrent Access & Permission Enforcement	55
3.2.7.3	Creation and Sharing of Triplestores	55
3.2.7.4	Sharing of Secrets	56
3.2.7.5	Management of Roles	56
3.2.8	Definition of SEnTriS	56
3.3	Summary	59
4	Implementation	60
4.1	General Considerations	60
4.1.1	Apache JENA Limitations	60
4.1.2	Memory Constraints	60
4.1.3	Design & Architecture Choices	60
4.2	Component Specific Implementation Details	60
4.2.1	Client	60
4.2.1.1	Exposed API	60
4.2.1.2	Architecture	60
4.2.1.3	Implementation Details	60
4.2.2	Triplestore	60
4.2.2.1	Exposed API	60
4.2.2.2	Architecture	60
4.2.2.3	Implementation Details	60
4.2.3	IAM Provider	61
4.2.3.1	Exposed API	61
4.2.3.2	Architecture	61
4.2.3.3	Implementation Details	61
4.2.4	Proxy	61
4.2.4.1	Exposed API	61
4.2.4.2	Architecture	61
4.2.4.3	Implementation Details	61
4.2.5	Vault	61
4.2.5.1	Exposed API	61
4.2.5.2	Architecture	61
4.2.5.3	Implementation Details	61
4.2.6	Discussion	61
4.3	Summary	61
5	System Evaluation	62
5.1	Experimental Evaluation	62
5.1.1	Setup & Methodology	62
5.1.1.1	LUBM Benchmark adaptation	62
5.1.1.2	Experiments	62
5.1.2	Experimental Results	62
5.1.2.1	Repository Size & (Up)load Time	62
5.1.2.2	Query Response Time	62
5.1.2.3	Query Completeness & Soundness	62
5.1.2.4	Combined Metric	62
5.1.2.5	Result Analysis	62
5.2	Security Analysis	62
5.2.1	Security Properties	62

5.2.2	Leakage Analysis	63
5.2.2.1	SEnTriS V1	63
5.2.2.2	SEnTriS V2	63
5.2.3	Security Properties Proofs	63
5.3	Complexity Analysis	63
5.3.1	Storage Complexity	63
5.3.1.1	SEnTriS V1	63
5.3.1.2	SEnTriS V2	63
5.3.2	Query Evaluation Complexity	63
5.3.2.1	SEnTriS V1	63
5.3.2.2	SEnTriS V2	63
5.4	Final Discussion	63
6	Conclusion & Future Work	65
	Bibliography	66
	Appendices	
A	SPARQL-EVAL Listings	75
B	SSE-TP Listings	96
C	SSE-SPARQL-EVAL Listings	109
D	IAM Listings	122

LIST OF FIGURES

2.1	The <i>Semantic Web</i> stack [120].	6
2.2	A simple RDF graphical example.	8
2.3	A simple RDFS graphical example.	9
2.4	The Web Ontology Language (OWL2) Structure [100].	11
2.5	Interaction between client and server.	22

LIST OF TABLES

2.1	Resource Description Framework Schema (RDFS) Core Classes and respective description.	10
2.2	RDFS Core Properties and respective description.	10

LIST OF LISTINGS

2.1	Simple RDF example, written in Terse RDF Triple Language (TURTLE) syntax.	8
2.2	Simple OWL2 example, written in TURTLE syntax.	12
2.3	SPARQL query with a conjunction pattern	17
2.4	SPARQL query with filters and modifiers	17
2.5	SPARQL query with a disjunction pattern	18
A.1	Triple Pattern Storage Engine	75
A.2	Bindings Table	77
A.3	SPARQL Query Engine	80
A.4	SPARQL Query Planner (without reasoning capabilities)	80
A.5	Query Execution	85
A.6	Query Worker	85
A.7	SPARQL-EVAL client	87
A.8	SPARQL-EVAL triplestore	88
A.9	Storage of schema triples	89
A.10	SPARQL-EVAL-II triplestore (extra operations)	89
A.11	Ontology	90
A.12	Query Plan Generation with Query Rewriting	92
A.13	SPARQL-EVAL-II client (modified operations)	95
B.1	Base logic for the encryption scheme	96
B.2	Triple Pattern Encryption for SEnTriS V1	98
B.3	Triple Pattern Encryption for SEnTriS V2	99
B.4	SSE-TP Storage Engine	100
B.5	SSE-TP Triplestore Component	101
B.6	Searches in SSE-TP Client (SEnTriS V1)	102
B.7	Searches in SSE-TP Client (SEnTriS V2)	103
B.8	Write operations in SSE-TP Client (SEnTriS V1)	103
B.9	Write operations in SSE-TP Client (SEnTriS V2)	106
C.1	SSE-SPARQL-EVAL ontology	109
C.2	SSE-SPARQL-EVAL SPARQL Planner	110
C.3	Query Preparation in SSE-SPARQL-EVAL Client (SEnTriS V1)	113
C.4	Query Preparation in SSE-SPARQL-EVAL Client (SEnTriS V2)	114
C.5	SSE-SPARQL-EVAL Triplestore	114
C.6	SSE-SPARQL-EVAL Triplestore Storage Engine	114
C.7	Query Preparation in SSE-SPARQL-EVAL Proxy	115
C.8	Persisting Bindings in SSE-SPARQL-EVAL Proxy Storage Engine	115

C.9 Query Evaluation in SSE-SPARQL-EVAL Proxy	115
C.10 SSE-SPARQL-EVAL Query Worker	116
C.11 Constructing Bindings Tables (SEnTriS V1)	116
C.12 Keyed set operations in SSE-SPARQL-EVAL Query Worker (SEnTriS V2)	117
C.13 Constructing Bindings Tables (SEnTriS V2)	117
C.14 Bindings Tables (SEnTriS V2)	118
C.15 Generating Final Results in SSE-SPARQL-EVAL (SEnTriS V1)	119
C.16 Generating Final Results in SSE-SPARQL-EVAL (SEnTriS V2)	121
D.1 Authentication in Client	122
D.2 Session Validation in IAM Provider	122
D.3 User Data Storage	123
D.4 Locks Client	124
D.5 Permission Checks	125
D.6 ACL & Access Tokens Data Storage	127
D.7 Access Management in Client	129
D.8 Access Management in IAM Provider	131
D.9 Storage of Access Requests	135
D.10 Secret Management in Vault	136
D.11 Role Management in Client	136
D.12 Role Management in IAM Provider	137
D.13 Storage of Role Requests	138

ACRONYMS

ACL	Access Control List 33 , 50 , 54 , 55
AES	Advanced Encryption Standard 26
AES-CBC	Advanced Encryption Standard in cipher-block chaining mode 28
API	Application Programming Interface 19 , 60
BGP	Basic Graph Pattern 15 , 40 , 42
BSBM	Berlin SPARQL Benchmark 19
DBMS	Database Management Systems 27
DGK	Damgård, Geisler and Krøigaard 5 , 36 , 38 , 46 , 49
FHE	Fully Homomorphic Encryption 3 , 20
GO	Gene Ontology 2
GOOSE	Graph Outsourcing and SPARQL Evaluation 4 , 28 , 29 , 30
HE	Homomorphic Encryption 3 , 6 , 20 , 21 , 30
HTML	HyperText Markup Language 8
HTTPS	Hyper Text Transfer Protocol Secure 60
IAM	Identity and Access Management 33 , 55 , 59
IND-CKA2	Indistinguishability Against Adaptive Chosen Keyword Attack 23 , 24
IND-CKA1	Indistinguishability Against Non-Adaptive Chosen Keyword Attack 23 , 24
IND-CPA	Indistinguishability Against Chosen Plaintext Attacks 23
IRI	Internationalized Resource Identifier 6 , 7 , 8 , 12 , 15
LFHE	Leveled-Fully Homomorphic Encryption 3 , 20
LUBM	Lehigh University Benchmark ii , iii , 19
M/M	Multiple Writer/Multiple Reader ii , 22 , 31 , 50 , 56 , 59
M/S	Multiple Writer/Single Reader 22
NCI	National Cancer Institute 2

NIST	National Institute of Standards and Technology 55
OPE	Order-Preserving Encryption 21, 28
ORAM	Oblivious RAM 3, 20, 30
OWL2	Web Ontology Language viii, x, 1, 7, 11, 12, 13, 14, 15, 18, 19, 25, 30
PBKDF2	Password-based Key Derivation Function 2 55
PEKS	Public Key Encryption with Keyword Searchable 3, 22, 30
PHE	Partially Homomorphic Encryption 3, 20
PRF	Pseudo-Random Function 24
PRP	Pseudo-Random Permutation 24
RBAC	Role-Based Access Control ii, iii, 31, 32
RDF	Resource Description Language ii, iii, xiii, 1, 5, 7, 8, 9, 10, 11, 15, 16, 17, 18, 19, 28, 29, 30, 31, 33
RDFS	Resource Description Framework Schema ix, 1, 7, 9, 10, 11, 12, 18, 30
REST	Representation State Transfer 60
RSA	Rivest-Shamir-Adleman 36, 37
S/M	Single Writer/Multiple Reader 22
S/S	Single Writer/Single Reader 22, 30
SE	Searchable Encryption 3, 6, 20, 22, 23, 30, 31
SEnTriS	Searchable Encrypted Triplestore Sharing ii, iii, 5, 31, 32, 33, 34, 35, 36, 39, 41, 42, 43, 44, 45, 46, 47, 48, 49, 55, 56, 58, 59, 62, 63
SHE	Somewhat Homomorphic Encryption 3, 20
SLA	Service Level Agreement 2
SNOMED CT	Systematized Nomenclature of Medicine Clinical Terms 2
SPARQL	SPARQL Protocol and RDF Query Language ii, iii, x, 1, 4, 5, 7, 15, 16, 17, 18, 19, 28, 29, 30, 31, 33, 34, 35, 36, 38, 39, 40, 41, 42, 43, 47, 48, 49, 56, 57, 59, 63
SQL	Structured Query Language 16, 26, 27, 28
SSE	Symmetric Searchable Encryption ii, 3, 4, 5, 22, 23, 24, 25, 30, 31, 33, 43, 44, 59
STE	Structured Symmetric Encryption 3, 25
SWRL	Semantic Web Rule Language 7
TLS	Transport Layer Security 60
TURTLE	Terse RDF Triple Language x, 8, 12, 17
UCRPQ	Unions of Conjunctions of Regular Path Queries 4, 28, 29
UUID	Universally Unique Identifier 33, 34, 35, 45, 46, 48
W3C	World Wide Web Consortium 7
WWW	World Wide Web 1
XML	Extensible Markup Language 7, 8, 11

INTRODUCTION

In this chapter we will contextualize and define the research problem. In Section 1.1 we introduce the research context: explaining the concept of the *Semantic Web*, analyzing the industry landscape, enumerate security concerns, that the usage of cloud technologies poses, and show mitigation approaches. In Section 1.2 we define the problem and research questions. Finally, In Section 1.3 we present the main contributions from our work.

1.1 Context

The Semantic Web In 2001, the term *Semantic Web* [7] was first used to define what would be the evolution of the *World Wide Web (WWW)*. The traditional *Web* solved the challenge of distributing information at scale and globally, however it was designed for human use. Documents are easily accessible and readable by humans, but for automated agents little more than document retrieval is possible. Contrary to humans, machines cannot reason over the information they access, if they do not have access to the *semantics of the data*.

The *Semantic Web* addresses this issue, by redesigning the representation of information, in a way that not only preserves meaning and relationships but is also machine-readable. A concept is mapped to a unique identifier, and its meaning is encoded in sets of triples, each triple consisting of a subject, predicate and object (similar to an elementary sentence). *Ontologies* provide a shared understanding of a domain, by formalizing the specification of relevant concepts, and their relations, in such domain.

This novel method of knowledge representation differs from the previous, centralized way, and greatly simplifies the process of interconnecting different knowledge bases. Centralized representations tend to be more expensive to develop, usually needing to be built from scratch [61]. Furthermore, for sharing, both parties need to agree on the semantics. In contrast, using ontologies, information is semantically linked by default. In the *Semantic Web*, machines can truly communicate, share and reason about distinct knowledge bases.

Semantic data is specified using a stack of different language specifications. The basic data model is specified using *RDF* [107], these resources can be organized in hierarchies with *RDFS* [111] and to define ontologies we use *OWL2* [100]. To query semantic data, the standard language is *SPARQL* [115, 116, 117].

Industry Landscape Knowledge representation is essential in the industry [65] and ontologies help remove ambiguity in terminology. Nowadays, they are being developed

and applied in a variety of domains [113] ranging from the health sector¹, to linguistics² and even being used by companies like NASA³ [6].

To develop and maintain ontologies, advanced semantic editors are the standard tool for professionals, with Protégé [60] being a very popular example. These provide numerous functionalities, usually based on plugins (e.g. automatic verification of errors and consistency in the design, debugging tools, query answering and graphic visualization of data and knowledge).

With datasets gradually growing in dimension, the industry is currently evolving towards the use of collaborative solutions [96, 95] and outsourcing heavy computations, relying on cloud infrastructures or other shared computing platforms. These approaches, albeit inevitable, present new security concerns and engineering challenges.

By 2021, 50 percent of all enterprise workloads were already hosted in *public clouds* [82]. As of 2023, worldwide end-user spending on public cloud services is forecast to grow 20.7% more than the growth forecast for 2022 of 18.8% [33].

Research on secure cloud computing [14] studies such adoption but also discusses security concerns. Adoption can be explained with the cost efficiency⁴, high scalability, flexibility and availability that cloud provides. Additionally, the pay-per-use model is an easy entry point to the ecosystem.

Cloud Security Concerns While delegation of responsibilities is beneficial for costs, it is prejudicial to control. Using cloud services implies trusting its owner. A [Service Level Agreement \(SLA\)](#) defines, in accordance to metrics, the service levels of the system that the vendor should provide (e.g. percentage of availability), while also defining penalties in case of not achieving what was agreed-on. However, the risk of loss of privacy and data breaches still exists. A problem that is especially concerning when sensitive data is being stored by such services. This is much more relevant in *public clouds* where, compared to *private clouds*, trust is distributed between various actors and not on a single organization, namely the cloud operator [14].

In cloud storage services, administrators, or hacker with root privileges, might have access to the data [10]. Additionally, these cloud services might outsource computations to third parties whom we do not trust. In this setting, we usually consider the *honest-but-curious*, or *passive*, adversary model. This adversary will follow the protocol properly but keep a record of all its intermediate computations [35]. Furthermore, it cannot do any other action except attempting to learn private information while observing views of the protocol execution [27].

Security Approaches A first approach to secure data in the cloud would be to use *Encryption at Rest*, which naively encrypts all data stored in the untrusted servers. If a user wants to update or query the data, they will have to download all the data, and

¹The NCI Thesaurus, [84], developed by the [National Cancer Institute \(NCI\)](#), the [Gene Ontology \(GO\)](#) [39], the world's largest source of information on the functions of genes, or the [Systematized Nomenclature of Medicine Clinical Terms \(SNOMED CT\)](#) [119], the most comprehensive, multilingual clinical healthcare terminology in the world, are some existing applications.

²The WordNet [57] is a lexical database for English where related words and concepts are semantically linked.

³Observing the visual representation provided by the Linked Open Data Cloud [55] is a good exercise to perceive the dimension of adoption of this new paradigm.

⁴Cloud services reduce implementation, maintenance, power, cooling, floor space, storage and operational costs.

decrypt it. This incurs in high network requirements and is computationally expensive on the client-side, hindering all the benefits of using a cloud service to store the data.

Although we can observe some major database services, like MongoDB, starting to adopt [Searchable Encryption](#) [59], the majority of cloud and storage providers only support the former solution.

Currently, active research is being done in the fields of [Searchable Encryption](#) (SE) and [Homomorphic Encryption](#) (HE).

[Homomorphic Encryption](#) is based on the concept of algebra homomorphisms, where both decryption and encryption can be seen as a homomorphism of the relations between the plaintext and the ciphertext [3, 31]. [Homomorphic Encryption](#), based on the type and number of operations supported (e.g. addition or multiplication of encrypted values), can be classified as: [Fully Homomorphic Encryption](#) (FHE) supporting any number of operations with no restriction on the number of times; [Somewhat Homomorphic Encryption](#) (SHE) permitting some types of operations, a limited number of times; [Partially Homomorphic Encryption](#) (PHE) allowing one type of operation with no restriction on the number of times; and [Leveled-Fully Homomorphic Encryption](#) (LFHE) [11] supporting any type of operations, but with a pre-determined expected depth- L (can be executed up to L times).

[Searchable Encryption](#) uses classical encryption schemes to perform encrypted search queries, sequentially over the ciphertext [86], or through the generation of searchable *encrypted indexes* accessible with *tokens*, generated via *trapdoor* functions. [Searchable Encryption](#) can be divided in [SSE](#), when it is based on symmetric key cryptography, and [Public Key Encryption with Keyword Searchable](#) (PEKS), when asymmetric key algorithms are used. [SSE](#) is often more practical than [PEKS](#), as symmetric cryptographic schemes tend to be more efficient than their asymmetric counterparts [71]. These techniques may leak some information about the *indexes*, *search pattern*, and *access pattern* during execution. Most approaches leak, at least, the *search* and *access patterns* [10]. Recently, to mitigate information leakage, *hybrid* solutions combining the software approach of [SE](#) with *trusted hardware* have been proposed [29].

Besides [HE](#) and [SE](#), [Oblivious RAM](#) (ORAM) [36] could theoretically be used as an alternative. However, [ORAM](#) and [HE](#) continue to be too inefficient for real world applications, when considered as an exclusive solution [10].

1.2 Problem

Ontologies are often used alongside electronic health records and clinical decision systems. These systems heavily manipulate personal data, which is sensitive. Additionally, ontologies themselves can have an inherent private value, due to the investment of time and money in their development, or the nature of their data. With all these potential security needs, a solution that supports the secure development, sharing and usage of ontologies is essential.

Taking into consideration the efficiency, security trade-offs and ease of implementation from the various alternative techniques, a pure software [SSE](#) approach seems to be the most practical and easier to adopt solution, not requiring excessive computational power or special trusted hardware. Recently, Chase and Kamara, on their work on [Structured Symmetric Encryption](#) (STE), generalized index-based [SSE](#) to support queries on *structured* data [15].

Semantic data can be represented as a graph and [STE](#) supports various types of query operations in graph-structured data. Some existing schemes, adapted to conceptual

graphs, need to pre-compute all possible query answers *a priori*, which is not practical for large datasets [74]. Furthermore, there can be leakage of information about the topology of the data while trying to pre-compute all query answers.

Recent developments do not require pre-computation of all query answers. The *CryptGraphDB* [2] system is an adaptation of *CryptDB* [77, 76, 75] capable of answering queries over conceptual graphs, but it was developed for the cypher [63] language and not *SPARQL*.

Another work, focused on access control for RDF datasets, did also implement a solution for retrieving simple triple patterns over an encrypted dataset, but, due to performance needs, hashes were used for indexing the encrypted triples. This leaks the structure of the dataset, which they tried to mitigate by creating dummy triples [28].

The *Graph Outsourcing and SPARQL Evaluation (GOOSE)* [17] framework, demonstrates how to use multiparty computation techniques to answer *Unions of Conjunctions of Regular Path Queries (UCRPQ)* over encrypted graphs, only relying on a standard *SPARQL* engine. But deterministic data obfuscation techniques are used, and therefore the structure of the graph is also leaked. This framework only focuses on *SPARQL* read queries.

Although, a lot of progress has been made in this field of research there are still various open questions. In this thesis, we will focus on answering the following questions:

1. *How to support secure and expressive read and write *SPARQL* queries over encrypted data?*

A practical solution ought to support more than basic triple pattern retrieval, it must be capable of querying complex patterns. As such it should be able to execute secure set operations without the client needing to decrypt the intermediate values (obtained from simple triple patterns). Furthermore, it must be able to execute write queries, so that datasets can be updated.

2. *How can we add reasoning capabilities to the solution?*

A simple *SPARQL* engine does not possess reasoning capabilities. The query results may or may not possess inferred values, depending on the underlying storage solution. When working with ontologies, a system should have reasoning capabilities to take full advantage of the semantics preserved within the data.

3. *How can we minimize the information leakage and enforce access control?*

A data outsourcing solution should support granular access control over the shared datasets. *SSE* solutions always leak some information about the data, we will investigate how to minimize such leakage during query execution.

1.3 Contributions

In this thesis we aim to develop a practical and secure solution for the collaborative development, distribution and usage of semantic data knowledge bases. Our main contributions are the following:

SPARQL-EVAL & SPARQL-EVAL-II - Distributed two-party *SPARQL* query evaluation protocol: We developed two *SPARQL* query evaluation protocols which support a subset of queries: queries target solely the default graph; the allowed modifier are *DISTINCT*, *ORDERBY* and *LIMIT*; graph patterns must not contain property

paths, nested subqueries queries, graph selection, nor triple patterns composed solely by variables. In SPARQL-EVAL-II we use query rewriting techniques to retain reasoning capabilities.

SSE-TP - A distributed two-party triple pattern SSE protocol: SSE-TP is protocol for encrypting [RDF](#) data in a way that the data remains searchable via the specification of triple patterns.

SSE-SPARQL-EVAL - A distributed three-party SPARQL query evaluation protocol for SSE-TP encrypted triplestores: SSE-SPARQL-EVAL is an adaptation of the SPARQL-EVAL-II protocol that works with SSE-TP encrypted triplestores.

SEnTriS - A multiparty searchable encrypted triplestore sharing protocol:

SEnTriS is an adaptation of the SSE-SPARQL-EVAL protocol for a multiuser setting. Encrypted triplestores are shareable, remain queryable, possess reasoning capabilities and support role based access control. Access to triplestores is controlled with the usage of ephemeral access tokens. Concurrent writes are resolved with a lock-based strategy. Secrets are distributed among user with access to the triplestores. The protocol is implemented in two versions:

- [SEnTriS V1](#): Data encryption retains elements from *CryptDB* [77, 76, 75]. Subjects, predicates and objects are encrypted in two layers. The outer layer is probabilistically encrypted, whereas the second is deterministically encrypted. When equality checks are required, we peel the outer layer of the encryption onions and compare the values.
- [SEnTriS V2](#): Data is encrypted using the [Damgård, Geisler and Krøigaard \(DGK\)](#) cryptosystem [23, 22, 32, 52], an additive homomorphic scheme developed for secure integer comparison. The value of a subject, predicate or object is mapped to an *equality tags*. Trapdoors index the *equality tags* and equality tags reference the probabilistically encrypted subjects, predicates and objects. During query evaluation, we mask the values of the equality tags, via re-encryption. When doing equality checks, we decrypt the equality tags and compare the values. These remain deterministic, but different for every query evaluation.

1.4 Document Organization

The remaining of the document will be organized as follows:

Chapter 2 discusses related work on *Semantic Web* and *Secure Computations* and compares existing solutions for secure semantic data outsourcing and [SPARQL](#) evaluation.

Chapter 3 presents the solution.

Chapter 4 describes the implementation details of each of the system's components.

Chapter 5 presents the experimental evaluation procedures, describes and discusses the obtained results.

Chapter 6 summarizes main conclusions from the research and suggests future work.

RELATED WORK

In this chapter, we will discuss the architecture and design of the *Semantic Web* as well as analyzing work already done in the fields of *Secure Computations*. The remaining of the chapter is organized as follows: in Section 2.1 we discuss the *Semantic Web* stack, languages and tools; Section 2.2 describes work on *Secure Computation*, with particular attention regarding [HE](#) and [SE](#); finally, in Section 2.4 we summarize and discuss which work is relevant.

2.1 Semantic Web: A layered architecture

In Section 1.1, we have contextualized the main philosophy and design principles of the *Semantic Web*:

1. Usage of *labeled graphs* as the data model where concepts are represented as triples consisting of a subject, predicate and object (with nodes being objects and relations between such objects modelled as the edges);
2. All data items are identified by [Internationalized Resource Identifier \(IRI\)](#) [25];
3. Usage of ontologies to model semantics within a domain.




Figure 2.1: The *Semantic Web* stack [120].

Additionally, the *Semantic Web* can also be seen as a stack of various layers [5] (Figure 2.1). Most of these layers are standardized due to the collective work of the [World Wide Web Consortium \(W3C\)](#).

Briefly describing this layered architecture, at the root level, we have the [IRIs](#) and a character set providing the baseline of communication, with distinguishable identification of resources and an agreed encoding.

At the syntax level, we find the [Extensible Markup Language \(XML\)](#) [98]. [XML](#) supports the definition of unlimited *tags*, that can be placed in the text of a document. These *tags* facilitate the creation of hierarchy and give *structure* to the information.

From structured data we then define the basic data model using [RDF](#) [107]. With [XML](#) like syntax, we can write statements about resources, thus embedding *semantics* within the information, as *metadata*.

[RDFS](#) [111] is a vocabulary description language to organize [RDF](#) resources into hierarchies, based on primitives such as *classes* and *properties*. It can be seen as a rudimentary *ontology language* [5].

Ontologies are defined with [OWL2](#) [100]. Like [RDFS](#), it is also a vocabulary description language, but provides more expressiveness (e.g. relations between classes, *cardinality*, *equality*, richer properties with characteristics).

[RDF](#) triples can be interpreted as logical binary predicates, making reasoning possible over [RDF](#)-based data. This allows us to infer and validate ontologies, and also provide explanations on operations' results. The rules layer is a reference to this specific part of the architecture stack. There exist various languages, such as the [Semantic Web Rule Language \(SWRL\)](#) [118], that allow us to express rules.

To query and retrieve any data based on [RDF](#) we use [SPARQL](#).

The Proof and Unifying Logic layers include the notion of interconnectivity of terminologies between *knowledge bases*, with explainable proofs to operations.

Only with secure cryptographic schemes can we reach the trust layer. The security layer is still not standardized.

The top layer, references all the applications and interfaces that use this new paradigm of data representation.

In the following of this section we will describe, in more detail, some languages and tools. We will focus on [RDF](#) in Section 2.1.1; on [RDFS](#), in Section 2.1.2; on [OWL2](#), in Section 2.1.3; on [SPARQL](#), in Section 2.1.4; on common tools and software, in Section 2.1.5.

2.1.1 RDF: Resource Description Framework

2.1.1.1 Overview

[RDF](#) [107] is a flexible *domain-independent* data model. The [RDF](#) basic primitives are *resources*, *properties* and *literals*. We can construct *statements* about resources by organizing these basic primitives in sets of *triples*. A triple is composed of a *subject*, a *predicate* and an *object*. In a statement (triple) about a subject (always a resource) we define the value (object) of a property (predicate), which can be another resource or a literal.

Literals are atomic values (e.g. strings, numbers). Resources and properties, respectively, model the "things" and relationships between such "things". Both are identified using [IRIs](#). Resources without an [IRI](#) are called *blank nodes* (or *bnodes*).

To make a statement *B* about a statement *A* we can use a technique called *reification*. Reification consists in: firstly, defining a statement *B* where the object *objA* represents the statement *A*; secondly, creating the properties *subject*, *predicate* and *object* mapped from

objA to the original subject, predicate (now a resource), and original object of statement *A* [5].

2.1.1.2 Syntax

Graphical RDF being a data model, can be written using different syntaxes. Graphically we can represent statements as a labeled and directed graph. The *nodes* are labelled with the resources' IRIs, a literal value or left blank, whereas *arcs* are labelled with the properties' IRIs. The arcs are directed from subject to object of a statement. Figure 2.2 provides an example of a RDF graph.




Figure 2.2: A simple RDF graphical example.

Text-Based There are multiple text-based syntaxes for RDF. Some of the most commonly used are **TURTLE** [109] (Listing 2.1), and its extension for named graphs *TriG* [108]. In **TURTLE**, IRIs are enclosed by angle brackets. The subject, property, and object are written in this order. Statements are finished by a period. Literals are surrounded in quotes with the type of data specified with an IRI. The data value and type are separated by ^^ (if no data type is mentioned, the value is inferred as being a string). We can also define IRI abbreviations (using @prefix, followed by the prefix to substitute the IRI, a colon, the IRI and a period) and simplify statements, for easier readability (e.g. using a semicolon to make more than one statement about the same subject, defining statements with nested blank nodes using square brackets). Comments start with #.

```

1 # FCT is located in Almada
2 <http://www.my-custom-domain/faculties.ttl#FCT> <http://www.my-custom-domain/ontology/
   ↳ location> <http://www.my-custom-domain/resource/Almada> .
3
4 # FCT has 8.000 students
5 <http://www.my-custom-domain/faculties.ttl#FCT> <http://www.my-custom-domain/faculties.
   ↳ ttl#asNumberOfStudents> "8000"^^<http://www.w3.org/2001/XMLSchema#integer> .

```

Listing 2.1: Simple RDF example, written in **TURTLE** syntax.

XML/RDF [110] is a standardized XML-based syntax and *RDFa* [112] permits RDF to be embedded in the **HyperText Markup Language (HTML)**¹.

¹We recommend the curious reader to read the W3C recommendations to better learn the syntaxes.

2.1.2 RDFS: RDF Schema

2.1.2.1 Overview

RDF is a flexible data-model. However, for interconnectivity of different RDF knowledge bases, semantics for restricting meaning in a specific domain are needed.

RDFS [111] does that by providing vocabulary to describe groups of related resources and the relationships between these resources. This vocabulary is a set of RDF resources identified under <http://www.w3.org/2000/01/rdf-schema#>.




Figure 2.3: A simple RDFS graphical example.

2.1.2.2 RDFS primitives

Classes group resources together, the *instances* of the class. *Properties* are defined by specifying to which classes of resource they may apply, using the mechanism of *domain* and *range*. This allows for easy extension, by not having to redefine classes wherever a new property, that applies to an existing class, is needed.

To state that a resource is an instance of a class we use the *rdf:type* property.

The *rdfs:subClassOf* is used to define that all the instances of a class are instances of its *superclasses*: if *A* *rdfs:subClassOf* *B*, then instances of *A* are also instances of *B*.

The *rdfs:subPropertyOf* states that instances related by a property are also related by its *superproperties*: if A *rdfs:subPropertyOf* B , then instances related to A are also related to B . The properties *rdfs:subClassOf* and *rdfs:subPropertyOf* are *transitive*, and there are no restrictions on the number of superclasses or superproperties that, respectively, a class or a property can have.

The *rdfs:domain* property defines that any resource with a given property is an instance of the domain classes that the property applies to. The *rdfs:range* property specifies that the values of the property are instances its range classes.

In tables 2.1 and 2.2, we summarize the core RDFS classes and properties. In Figure 2.3, we extend the RDF graphical Example 2.2 with RDFS.

RDFS Core Classes	
Class	Description
rdfs:Resource	The class resource, everything.
rdfs:Class	The class of classes.
rdfs:Literal	The class of literal values, e.g. textual strings and integers.
rdf:Property	The class of RDF properties.
rdfs:Statement	The class of RDF statements.

Table 2.1: RDFS Core Classes and respective description.

Core RDFS Properties			
Property	Description	Domain	Range
Relationship Definition			
rdf:type	The subject is an instance of a class.	rdfs:Resource	rdfs:Class
rdfs:subClassOf	The subject is a subclass of a class.	rdfs:Class	rdfs:Class
rdfs:subPropertyOf	The subject is a subproperty of a property.	rdf:Property	rdf:Property
Property Restriction			
rdfs:domain	A domain of the subject property.	rdf:Property	rdfs:Class
rdfs:range	A range of the subject property.	rdf:Property	rdfs:Class
Reification			
rdf:subject	The subject of the subject RDF statement.	rdf:Statement	rdfs:Resource
rdf:predicate	The predicate of RDF statements.	rdf:Statement	rdfs:Resource
rdf:object	The object of RDF statements.	rdf:Statement	rdfs:Resource
Utility			
rdfs:seeAlso	Further information about the subject resource.	rdfs:Resource	rdfs:Resource
rdfs:isDefinedBy	The definition of the subject resource.	rdfs:Resource	rdfs:Resource
rdfs:comment	A description of the subject resource.	rdfs:Resource	rdfs:Literal
rdfs:label	A human-readable name for the subject.	rdfs:Resource	rdfs:Literal

Table 2.2: RDFS Core Properties and respective description.

2.1.3 OWL2: Web Ontology Language

2.1.3.1 Overview

An *ontology* was defined by Studer et al. as "a formal, explicit specification of a shared conceptualisation"[88]. An ontology should specify consensual and machine-readable knowledge by means of a formal definition of concepts, where concepts have explicit types and restrictions. Therefore, an *ontology language* should provide a well-defined

syntax, formal semantics, convenient and sufficient expressiveness and efficient reasoning support.

RDFS can be considered a very limited *ontology language*. **OWL2** [100] provides a number of additional features (e.g. cardinality constraints, class equivalence, intersection and disjunction). **OWL2** is divided in two sublanguages: **OWL2 Full** [104], a direct extension of **RDFS** with all **OWL2** semantics, which is undecidable [5]; **OWL2 DL** [99], a language with some restrictions on the way **OWL2**, **RDF** and **RDFS** primitives may be used.

The correspondence theorem of [99] states that: given an **OWL2 DL** ontology, inferences drawn using its *Direct Semantics* will still be valid if the ontology is mapped into an **RDF** graph and interpreted using the *RDF-Based Semantics*, the semantics of **OWL2 Full**.

OWL2 can be expressed with any valid **RDF** syntax, however there are various specific syntaxes: the Functional-Style Syntax [105], which is very compact and was the language used in the formal specification of the language; the Manchester Syntax [101], designed to be human-readable (it is very popular in user interfaces of semantic editors); and the **OWL/XML** Syntax [106], which allows easier processing with **XML** tools.

Figure 2.4 details the structure of **OWL2**. Finally, **OWL2** can have different specifications, called profiles [102], which provide reasoning in polynomial time. The coverage of **OWL2** profiles is not in the scope of this work.




Figure 2.4: The **OWL2** Structure [100].

2.1.3.2 OWL Language

In **OWL2** the term individual is equivalent to instance in **RDFS**. An *assertion* is a statement about a resource. *Expressions* are a combination of classes, properties and instances. An

axiom relates such expressions to classes². Listing 2.2 provides a simple [OWL2](#) example.

Individual Assertions In [OWL2](#) *assertions* about individual resources can be used to define their type, class membership, properties and even provide ways to define identity and negative statements. More specifically on the identity statements, because [OWL2](#) has an open world assumption, we can never differentiate two individuals just because they have different [IRIs](#). To do so we use `owl:differentFrom`. To explicitly state negative facts we use `owl:NegativePropertyAssertion` [5].

```
1 @prefix owl: <http://www.w3.org/2002/07/owl#> .
2 @prefix rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#> .
3 @prefix rdfs: <http://www.w3.org/2000/01/rdf-schema#> .
4 @prefix xsd: <http://www.w3.org/2001/XMLSchema#> .
5
6 # The base namespace of the FCT ontology
7 @base<http://www.my-custom-domain/fct.ttl> .
8
9 <http://www.my-custom-domain/fct.ttl>
10   rdf:type owl:Ontology ;
11   rdfs:label "FCT_Ontology"^^xsd:string ;
12   rdfs:comment "An example OWL2 ontology"^^xsd:string ;
13   owl:versionIRI <http://www.my-custom-domain/fct.ttl#1.0> .
14
15 # Class Expression
16 _:x rdf:type owl:Class ;
17     owl:unionOf ( :Mandatory :Optional ) .
18
19 # Class Axiom
20 :Course owl:equivalentClass _:x .
21
22 # Class Assertion
23 :class101 rdf:type :Course .
```

Listing 2.2: Simple [OWL2](#) example, written in [TURTLE](#) syntax.

Classes In [OWL2](#) every resource that is of type `owl:Class` is a class. The two classes defined by default are *owl:Thing* and *owl:Nothing*. Respectively, these two specify the superclass of all classes and the subclass of all classes: every individual is a member of *owl:Thing* and no individual can be member of *owl:Nothing*. This is very useful for reasoning, because inconsistency is detected when classes are equivalent to *owl:Nothing* (have no individuals in their membership).

Class Axioms In [OWL2](#) we can relate *expressions* to classes using *axioms*. These can define:

- *Subclass Relations*, with equivalent usage like in [RDFS](#);
- *Class Equivalence*, in equivalent classes every individual must be a member of each class;
- *Enumerations*, a type of class definition via explicitly stating all individuals who belong to the class membership;

²For an overview of the language we recommend the reader to consult the reference guide [103].

- *Intersection*, in a class described as the intersection of various classes, all of its members must be also individuals belonging to the membership of each of those classes;
- *Union and Disjoint Union*, in a class described as the union of various classes, every individual belonging to its membership must also belong to the membership of at least one of those classes. In the particular case of disjoint union, the intersection of the classes is equivalent to owl:Nothing;
- *Disjoint Classes*, in disjoint classes no individual can be a member of both classes (the intersection of the classes is equivalent to owl:Nothing);
- *Complement*, complement classes contain all the individuals such that the union of both classes is equivalent to owl:Thing.

Properties OWL2 has three types of properties: *object* properties, which relate individuals to other individuals; *datatype* properties, which relate individuals to literals; and *annotation* properties; which are used to specify human-readable labels, comments and explanations.

The *top* and *bottom* properties specify, respectively, properties that relate to all and no individuals: owl:topObjectProperty is the superproperty of all owl:ObjectProperties and owl:bottomObjectProperty relates to no individual; owl:topDataProperty is the superproperty of all owl:DatatypeProperties, relating an individual to any possible literal value, and owl:bottomDataProperty relates no individual to any literal value.

Properties can have different characteristics:

- *Transitive Properties*: if *a* relates to *b* via property *p* and *c* relates to *a* via property *p*, then *c* relates to *b* via property *p*;
- *Symmetric Properties*: if *a* relates to *b* via property *p*, then inverse is true;
- *Asymmetric Properties*: if *a* relates to *b* via property *p*, then the inverse is $\perp$;
- *Functional Properties*: if for each instance the property can only have one unique value;
- *Inverse-functional Properties*: if the object of a property statement relates to a single subject;
- *Reflexive Properties*: if every individual relates to itself via the property;
- *Irreflexive Properties*: if no individual relates to itself via the property.

Property Axioms Just like class axioms, we can specify *property axioms* to describe the way properties relate to classes and other properties. These allow us to define:

- *Domain and Range*: used to determine class membership of individuals;
- *Inverse Relationship*: used to specify the inverse of the property;
- *Equivalence*: used to specify equivalent properties. Individuals related by one property will always relate by all equivalent properties;

- *Disjointness*: used to specify disjoint properties, where individuals related by one property can never be related by the other;
- *Property Chains*: used to aggregate connected properties under another property.

More Granular Class Axioms We can write more granular class definitions in [OWL2](#) than by simply defining class axioms. This is achieved by combining class and property axioms. It is done by defining a class axiom that relates to an anonymous class (owl:Restriction) which contains axioms on its properties. The anonymous class definition captures all individuals that satisfy such restrictions. As such we can define:

- *Universal and Existential Restrictions*: which the members of the class must satisfy. The class is defined as a subclass of an anonymous class with property restrictions on, respectively, owl:allValuesFrom and owl:someValuesFrom;
- *Necessary and Sufficient Conditions*: relating a class to a property, using rdfs:subClassOf, does not provide sufficient information to conclude whether an individual belongs to the class membership or not (it only states necessary, not sufficient, conditions for class membership). With owl:equivalentClass, it becomes clear that an individual belongs to the class membership, if it satisfies the restriction (the conditions are necessary and sufficient).
- *Value, Cardinality, Data Range and Datatype Restrictions* on the anonymous class: limiting the value, cardinality of the value, range of possible values and the datatype of the property value relating to members of the superclass;
- *Self Restrictions*: on the individuals captured by the anonymous class.

2.1.3.3 OWL2 DL

Contrary to *OWL2 Full*, *OWL2 DL*, being less expressive, retains *decidability*. This is achieved by imposition of various restrictions [5]:

- The application of [OWL2](#) primitives to each other is not allowed; only classes and non-literals resources may be defined, with all classes being individuals of owl:Class;
- Properties are disjoint individuals of either owl:ObjectProperty or owl:DatatypeProperty;
- Annotation properties are ignored by reasoners;
- Properties, for which range includes non-literal resources, are separated from those which relate to literal values;
- A resource can only be a class, property, or instance at the same time. Classes, properties or instances with the equal names will be treated as distinct (a mechanism called punning);
- Only functional properties can be assigned to datatype properties;
- Only object properties can have an inverse;
- Property chains can only involve object properties;
- Self restrictions on datatype properties are not allowed.

2.1.4 SPARQL

2.1.4.1 Overview

In the previous section we have discussed the structure of [OWL2](#). [OWL2](#) documents have a direct mapping to [RDF](#) documents. As such, we can publish and query [OWL2](#) ontologies using the standard approach for [RDF](#) data.

[RDF](#) triples are stored in *triple stores*, specialized databases for this type of data. The underlying architecture of such databases varies, from relational to NoSQL approaches [20, 58]. The recommended query language to use is [SPARQL](#) [115]. It provides a range of operations over published [RDF](#) graphs, accessible via SPARQL endpoints: from read queries [117] to write operations [116, 114].

2.1.4.2 SPARQL Queries

In this section we will mathematically define core components of *SPARQL*, briefly discuss the basis of functioning for this type of queries, explain how they are evaluated, given the semantics of *graph patterns*, and cover their syntax.

Definition 1 (RDF dataset). Let I, B, L be disjoint sets of [IRIs](#), blank nodes, and literals, respectively. [RDF](#) data is a set of subject, predicate and object triples $t = (s, p, o)$, where $s \in I \cup B$, $p \in I$ and $o \in I \cup B \cup L$.

Let V be a set of nodes formed by subjects and objects and E be a set of edges, labeled by the corresponding predicate. A set of [RDF](#) triples can be represented as a directed, edge-labeled, graph $G = (V, E)$. An [RDF](#) dataset is a set of *named* graphs and the *default* unnamed graph. We will define such dataset as $D = \{G, \gamma\}$, where γ is a function which assigns a named graph $\gamma(i)$ to each i in a finite set $\text{dom}(\gamma) \in I \cup B$ [73, 78].

Definition 2 (SPARQL Query). Let $T = \{\text{SELECT}, \text{ASK}, \text{CONSTRUCT}, \text{DESCRIBE}\}$, a [SPARQL](#) query can be defined as a tuple $q = (t_q, d, m, P_q)$ where $t_q \in T$ is the *query type*, $d \in I \cup B$ is an optional term identifying the [RDF](#) graph to be queried, m is a set of *solution modifiers* and P_q is a *graph pattern*.

Definition 3 (Graph pattern). Let $V = \{?v_0, \dots, ?v_n\}$, where $n \in \mathbb{Z}^+$, be a set of variables, disjoint from I, B , and L . We will denote the set of variables in P as $\text{vars}(P)$. A triple pattern is a tuple $tp = (s', p', o')$ such that $s' \in I \cup V$, $p' \in I \cup V$ and $o' \in I \cup L \cup V$. A set of triple patterns is usually referred as a [Basic Graph Pattern \(BGP\)](#). A *property path pattern* pp can be inductively defined as an element of $(I \cup B \cup V) \times pp \times (I \cup B \cup L \cup V)$. Finally, a *graph pattern* is an expression specified by the following grammar:

$$P ::= tp \mid pp \mid q \mid P_1 \cdot P_2 \mid P \text{ FILTER } R \mid P_1 \text{ UNION } P_2 \mid P_1 \text{ OPTIONAL } P_2 \mid P_1 \text{ MINUS } P_2 \mid \text{Graph } i \mid \text{Graph } v?P \quad (2.1)$$

Here tp is a triple pattern, pp is a *property path pattern*, q is a subquery, $i \in I$, $?v \in V$ and R is a [SPARQL](#) filter constraint, $P_1 \cdot P_2$ represents an equality join operation between two patterns³. To remove ambiguities parentheses can be added to a pattern. Finally, we will define the set of variables that occur in a pattern P as $\text{vars}(P)$ [73].

³For a detailed explanation of the syntax and semantics, we recommend the reader to consult section 17 of [117]

How do SPARQL queries work? An isomorphism of two graphs is an adjacency preserving bijection between their vertex sets [56].

The basis of SPARQL queries is finding matches with *graph patterns* in the RDF dataset [49, 20]. This problem is equivalent to the *subgraph isomorphism problem*, a generalization of the graph isomorphism problem: given two graphs G_1 and G_2 , we try to find if G_1 contains a subgraph that is isomorphic to G_2 .

In the context of a SPARQL query q , G_1 is a published RDF graph identified by d and G_2 corresponds to a *graph pattern* P_q . From this match we obtain a set of *bindings*: a multiset of mappings between $\text{vars}(P)$ and data within G_d . The final result of a query will also depend on t_q , the *query type*, and m , the specified *solution modifiers* (i.e. aggregation, grouping, sorting, duplicate removal).

Definition 4 (Semantics of SPARQL graph patterns). A mapping μ is a partial function $\mu: V \rightarrow I \cup B \cup L$ that assigns values to a finite set of variables. The domain of μ , denoted by $\text{dom}(\mu)$, is the subset of V where μ is defined.

Two mappings μ_1 and μ_2 are compatible, denoted by $\mu_1 \sim \mu_2$, if $\mu_1(?v) = \mu_2(?v)$ for all common variables $?v \in \text{dom}(\mu_1) \cap \text{dom}(\mu_2)$.

Abusing notation, we define $t = \mu(tp)$ as the triple obtained after replacing the variables in the *triple pattern* tp according to μ [51, 73]. Finally, let $\mu \models R$ represent a mapping that satisfies a filter condition R .

We define the *join*, *union*, *difference*, and *left outer join* operations on two sets of mappings Ω_1 and Ω_2 as follows:

$$\begin{aligned} \Omega_1 \bowtie \Omega_2 &= \{\mu_1 \cup \mu_2 \mid \mu_1 \in \Omega_1, \mu_2 \in \Omega_2, \mu_1 \sim \mu_2\} \\ \Omega_1 \cup \Omega_2 &= \{\mu \mid \mu \in \Omega_1 \vee \mu \in \Omega_2\} \\ \Omega_1 - \Omega_2 &= \{\mu_1 \in \Omega_1 \mid \forall \mu_2 \in \Omega_2, \mu_1 \not\sim \mu_2\} \\ \Omega_1 \Join \Omega_2 &= (\Omega_1 \bowtie \Omega_2) \cup (\Omega_1 - \Omega_2) \end{aligned} \tag{2.2}$$

The semantic evaluation of P over a dataset D is a set of mappings $\llbracket P \rrbracket_D$. Inductively, we define it as:

$$\begin{aligned} \llbracket tp \rrbracket_D &= \{\mu \mid \text{dom}(\mu) = \text{vars}(tp) \wedge \mu(tp) \in G_D\} \\ \llbracket pp \rrbracket_D &= \{\mu \mid \text{dom}(\mu) = \text{vars}(pp) \wedge \mu(pp) \in G_D\} \\ \llbracket q \rrbracket_D &= \{\mu \mid \text{dom}(\mu) = \text{vars}(P_q) \wedge \mu(P_q) \in G_D\} \\ \llbracket P_1 \cdot P_2 \rrbracket_D &= \llbracket P_1 \rrbracket_D \bowtie \llbracket P_2 \rrbracket_D \\ \llbracket P_1 \text{ UNION } P_2 \rrbracket_D &= \llbracket P_1 \rrbracket_D \cup \llbracket P_2 \rrbracket_D \\ \llbracket P_1 \text{ OPTIONAL } P_2 \rrbracket_D &= \llbracket P_1 \rrbracket_D \Join \llbracket P_2 \rrbracket_D \\ \llbracket P_1 \text{ MINUS } P_2 \rrbracket_D &= \llbracket P_1 \rrbracket_D - \llbracket P_2 \rrbracket_D \\ \llbracket P_1 \text{ FILTER } R \rrbracket_D &= \{\mu \in \llbracket P \rrbracket_D \mid \mu \models R\} \\ \llbracket \text{Graph } i \ P \rrbracket_D &= \{\mu \in \llbracket P \rrbracket_{(\gamma(i), \gamma)} \mid i \in \text{dom}(\gamma)\} \\ \llbracket \text{Graph } ?v \ P \rrbracket_D &= \bigcup_{i \in \text{dom}(\gamma)} \{?v \rightarrow i\} \bowtie \llbracket P \rrbracket_{(\gamma(i), \gamma)} \end{aligned} \tag{2.3}$$

Syntax SPARQL queries have a similar structure to the *Structured Query Language (SQL)* queries: we define the operation to execute using a keyword, followed by the projection of variables to be shown in the result, the keyword *WHERE*, the *graph pattern*,

enclosed in brackets, and, optionally, some modifiers of the result (i.e. *ORDER BY* to sort results and *DISTINCT* to remove duplicates). Additionally, we may also specify the *RDF* graph to consider for the query evaluation with the keywords *FROM* or *FROM NAMED*. Furthermore, the *SPARQL* syntax is very similar to *TURTLE*, so we can specify prefixes for namespaces. The basic *SPARQL 1.1 Query Language* [117] provides the following query types:

- *SELECT*: used to extract values from a *RDF* graph. The result is a set of bindings;
- *ASK*: returns a boolean, true if there was a match between the *graph pattern* and the *RDF* graph, false otherwise;
- *CONSTRUCT*: we specify two *graph patterns*, one to fetch *bindings* and another to use as a template for constructing new triples, from the values in the *bindings*. The final result is the set of newly generated triples.
- *DESCRIBE*: results depend on implementation [73], but usually returns a set of *RDF* triples describing the matching patterns.

In the following listings 2.3, 2.4 and 2.5 we show a few query examples.

```

1 @prefix rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#> .
2 @prefix fct: <http://www.my-custom-domain/fct.ttl>
3
4 # Operation and variables
5 SELECT ?opt_course ?man_course ?popular_course
6 # Query Pattern
7 WHERE
8 {
9     ?opt_course rdf:type fct:Optional .
10    ?man_course rdf:type fct:Mandatory
11 }

```

Listing 2.3: *SPARQL* query with a conjunction pattern

```

1 @prefix rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#> .
2 @prefix fct: <http://www.my-custom-domain/fct.ttl>
3
4 # Operation and variables
5 SELECT DISTINCT ?popular_course
6 # Query Pattern
7 WHERE
8 {
9     ?popular_course fct:hasStudentEnrolled ?num_students .
10    FILTER (?num_students > 60)
11 }
12 # Optional Modifiers
13 LIMIT 20

```

Listing 2.4: *SPARQL* query with filters and modifiers

```
1 @prefix rdf: <http://www.w3.org/1999/02/22-rdf-syntax-ns#> .
2 @prefix fct: <http://www.my-custom-domain/fct.ttl>
3 @prefix nova: <http://www.my-custom-domain/nova.ttl>
4
5 # Operation and variables
6 SELECT DISTINCT ?optional_course
7 # Query Pattern
8 WHERE
9 {
10   {?optional_course rdf:type fct:Optional}
11   UNION
12   {?optional_course rdf:type nova:Elective}
13 }
```

Listing 2.5: SPARQL query with a disjunction pattern

Besides read operations, *SPARQL Update* [116], provides additional options for inserting, publishing and deleting data in triplestores:

- *INSERT/DELETE*: these operations add and delete, respectively, newly generated triples, based on a template, into a destination graph (if it does not exist it should be created). Like in a *CONSTRUCT* query we specify two *graph patterns*, one to fetch bindings and another to use as a template for constructing triples. We can join both operations, in this situation we specify three *graph patterns*, two to use as templates to generate the triples to be inserted or deleted and a third one to fetch bindings.
- *INSERT/DELETE DATA*: similar to their pattern-based variants, but in these operations we specify concrete data inline.
- *LOAD*: used to read the contents of a document that represents a graph and saves it into a destination graph in a triplestore.
- *CLEAR*: this operation removes all the triples in the specified graphs.

2.1.4.3 Preserving reasoning

In a *knowledge base* we can distinguish two sets of statements: the *TBox* and the *ABox*. The *TBox* set contains terminology data (e.g. ontologies) which we can define using *RDFS* or *OWL2*. The *ABox* set refers to assertions.

SPARQL was not designed for vocabularies like *RDFS* or *OWL2*. By default, a *SPARQL* engine solely considers explicit data, as if it were at the *RDF* layer. Query evaluations interpret all data as assertions, completely ignoring the reasoning capabilities of *RDFS* or *OWL2*. This results in answers that might be incomplete, not containing inferred statements (only derivable from the *ABox* set, using the *TBox* data).

Two popular techniques for preserving reasoning capabilities are *materialization* and *query rewriting*.

Materialization In this approach, we derive all the implicit information, from the original data, and persist it. This solution insures fast query evaluation that requires no inference, but suffers from high storage costs. Additionally, write operations are more complex, due to the problem of *truth maintenance*. After each update, we need to maintain the materialized data in a consistent state.

Insertions are fairly simple, we just need to compute new implicit triples from the data already saved in the system.

Deletions, however, are more complex. We need to compute the set of triples to be removed and assure that no triple in it is entailed without the removed assertions. We usually find three types of implementations for this operation: to always recompute the materialization after a deletion; removing the inferred triples and recomputing the closure of the remaining; maintaining a data-structure with the state of deletions. Neither of these options is perfect, the first two are not efficient for databases with high volume of write operations and the last one incurs in increased costs (associated with maintaining an additional structure).

Query Rewriting This technique consists in deriving a complete answer set at query time. Before evaluating a query, the system modifies it according to the statements in the *TBox*. Query evaluation is slower: rewriting the queries requires communication with a reasoner and the reformulated queries tend to be larger.

Contrary to the first solution, there is no problem with *truth maintenance*. Writes operations can be processed without any additional computations.

2.1.5 Industry solutions

2.1.5.1 Development Software & Tools

To develop and maintain ontologies, professionals use advanced semantic editors. *Protégé* [60] is a free, open-source, and the usual choice. It provides a rich ontology editing environment with full support for *OWL2*, available in a desktop version or on the web [95]. Many of its functionalities are based on plugins: automatic verification of errors and consistency in the design using reasoners like *Pellet* [85]; debugging tools; query answering; collaborative development [96]; graphic visualization of data and knowledge. Additionally, for developers, many *Application Programming Interface (API)*s [68] or full framework solutions, such as *Apache Jena* [92], are available for building *Semantic Web* applications.

2.1.5.2 Triplestores

To store *RDF* data, popular solutions are: *MarkLogic* [54], *Virtuoso* [67], *Amazon Neptune* [4], *GraphDB* [66], *Apache Jena - TDB* [93]. While *Apache Jena - TDB* is primarily a *RDF* triplestore, the rest are *multi-model* database systems.

To test their capability to handle *SPARQL* queries, the *Berlin SPARQL Benchmark (BSBM)* is a standard benchmark [16]. The *LUBM* is another benchmark, which is tailored for *OWL2* data [37]. It provides synthetic data generation for a single realistic ontology.

T²

2) *TODO*: Add queries in annex, explain evaluation metrics in experimental evaluation chapter

2.2 Secure Computations

In Section 1.1, we have contextualized the motivations behind the need for secure computation when sharing ontologies over cloud or other distributed networks.

In this setting we consider the *honest-but-curious*, or *passive*, adversary model. As such an adversary will follow the protocol properly but keep a record of all intermediate computations [35] in an attempt to learn private information while observing views of the protocol execution [27]. Security concerns should take into account the operation history, data stored in the server and also the memory access pattern, as the adversary will have access to all of this information.

The *Encryption at Rest* approach encrypts all data stored in the untrusted server. To perform computations we need to download the full dataset, decrypt it, and only then can we execute operations on the data. This is too inflexible and computationally expensive.

ORAM [36] fits well this adversary model, as it allows for the concealment of memory access patterns by shuffling and continuously re-encrypting data on memory, but it is still impractical.

When used exclusively, **Homomorphic Encryption** [3, 31] schemes are not efficient enough to be feasible in real-world applications [10]. However, they are applied in conjunction with other security schemes, such as in some **Searchable Encryption** schemes [10, 15]. For this reason, in the following Sections we will discuss related work in the fields of **HE**, in Section 2.2.1, and **SE**, in Section 2.2.3.

2.2.1 Homomorphic Encryption

2.2.1.1 Overview

The term homomorphism, in security, was used for the first time by Rivest et al. in 1978 [79]. **Homomorphic Encryption** is a cryptographic method based on the concept of algebra homomorphism, where both decryption and encryption are seen as homomorphisms of the relations between the plaintext and the ciphertext [3, 31]. With **HE**, computations done over encrypted data, are preserved after decryption of the ciphertext.

Definition 5 (Homomorphic Encryption). An encryption scheme is homomorphic over operation $\bullet_1$ if it satisfies the following equation, where E is the encryption algorithm, $\bullet_1$ and $\bullet_2$ are arbitrary operations and M is the set of all possible messages [3, 32]:

$$E(m_1) \bullet_2 E(m_2) = E(m_1 \bullet_1 m_2), \forall m_1, m_2 \in M \quad (2.4)$$

2.2.1.2 Types of HE schemes

Some **HE** schemes preserve additive properties [69, 23, 22], others maintain multiplicative properties [80, 26], and some conserve both additive and multiplicative properties [34]. The **HE** schemes that preserve both properties allow for any arbitrary operation, because boolean circuits can be represented using only *XOR* (addition) and *AND* (multiplication) gates [3].

HE is classified in four main types:

- **Fully Homomorphic Encryption**: which supports any number of operations with no restriction on the number of operations;
- **Somewhat Homomorphic Encryption**: which permits some types of operations, a limited number of times;
- **Partially Homomorphic Encryption**: that allows for one type of operation with no restriction on the number of times;
- **Leveled-Fully Homomorphic Encryption** [11]: which supports any type of operations for a limited, and expected, number of L times (referred to as depth- L).

Definition 6 (Additive Homomorphic Encryption). An encryption scheme is as an additive homomorphic scheme if Equation 2.4 holds for the $+$ operator. In this thesis, we will uniquely consider schemes that respect this condition, if $\bullet_1$ is an addition and $\bullet_2$ is a multiplication:

$$E(m_1) * E(m_2) = E(m_1 + m_2) \quad (2.5)$$

Multiplication with constants Multiplication by a constant $k > 0$ can be performed on an encrypted value, by raising it to the power of k . This is equivalent to performing $k - 1$ multiplications:

$$E(m)^k = \underbrace{E(m) * \dots * E(m)}_{k \text{ times}} = E(\underbrace{m + \dots + m}_{k \text{ times}}) = E(k * m) \quad (2.6)$$

Considering that the plaintext space is $M \in \mathbb{Z}_n$ and that messages are small in comparison to the modulus n (so that reductions modulo n never occur), we can represent a negative constant c as $n - c$. Multiplication by a negative value is achieved by raising the encrypted value to the power of $n - c$ (an overflow of module n will occur)[32]:

$$\begin{aligned} E(m)^{(n-c)} &= E((n-c) * m \bmod n) \\ &= E(n * m - c * m \bmod n) \\ &= E(-c * m) \end{aligned} \quad (2.7)$$

Subtraction Given two messages m_1 and m_2 , subtraction is computed by:

$$\begin{aligned} E(m_1 - m_2) &= E(m_1 + (-1) * m_2) \\ &= E(m_1) * E((-1) * m_2) \\ &= E(m_1) * E(m_2)^{-1} \end{aligned} \quad (2.8)$$

2.2.2 Property-preserving Encryption

2.2.2.1 Overview

Homomorphic Encryption allows arbitrary computations to be executed on encrypted data while not leaking any information about the original plaintext. For operations such as equality joins, where we require equality checks, we do need to preserve (leak) certain properties from the original plaintexts. This type of encryption is known as property-preserving encryption.

2.2.2.2 Deterministic Encryption

Deterministic encryption is type of property-preserving encryption scheme. Contrary to a probabilistic scheme, it will always generate the same ciphertext for each plaintext. This allows equality checks to be performed on the encrypted data. Naturally, its security guaranties are weakened. Deterministic ciphertexts are vulnerable to statistical analysis attacks.

2.2.2.3 Order-Preserving Encryption

Order-Preserving Encryption (OPE) schemes preserve the order guaranties of the plaintext data. Naturally, the order relations between the ciphertext are leaked. Some implementations can leak as much as half of the data bits of the ciphertexts [76].

2.2.3 Searchable Encryption

2.2.3.1 Overview

Searchable Encryption uses classical encryption schemes to perform encrypted search queries. While the first approach implemented it via a searchable ciphertext [86], most modern schemes do it using encrypted searchable indexes. **Searchable Encryption** can be divided in **SSE**, when it is based on symmetric key cryptography, and **PEKS**, when public key algorithms are used. We will focus on index-based **SE** schemes, more specifically in index-based **SSE** ones, due to their higher efficiency compared to **PEKS** approaches.

2.2.3.2 General Index-Based Model

A general model, for the index-based **SE** schemes was defined in [10]. Given a database $DB = (M_1, \dots, M_n)$ consisting in a set of n plaintext documents M , some data items (e.g, keywords in the documents) $W = (w_1, \dots, w_m)$ are extracted from the documents and encrypted using a key K , using an encryption algorithm Enc . The documents may be encrypted too, generating a collection of ciphertexts C , usually with a different key K' . A client generates an encrypted index I , for the database DB , using a $BuildIndex$ function. The client stores the following information in an *honest-but-curious* untrusted server:

$$\begin{aligned} I &= BuildIndex_K(DB = (M_1, \dots, M_n), W = (w_1, \dots, w_m)) \\ C &= Enc_{K'}(DB) \end{aligned} \quad (2.9)$$

To perform search queries, the client generates a trapdoor $T = Trapdoor_K(f(w_j))$, for a predicate f a keyword w_j . With trapdoor T , the server performs the search over I , using a $Search$ algorithm, by checking for an encrypted index $i = Search(I, T)$. If there is a keyword w that satisfies the predicate f , it returns to the client the corresponding ciphertexts associated with i or $\perp$ if no match was found. In Figure 2.5 we demonstrate the interaction between client and server.




Figure 2.5: Interaction between client and server.

2.2.3.3 Client/Server Model

In **SE** schemes, one or various clients, the *writers* store encrypted data in a server. To issue a query request to the server, one or multiple clients, the *readers*, generate trapdoors. There are four possible **SE** scheme architectures: **Single Writer/Single Reader (S/S)**; **Single Writer/Multiple Reader (S/M)**; **Multiple Writer/Single Reader (M/S)**; **M/M**.

The **S/S** architecture is usually used for *data outsourcing* and **M/S**, **S/M** and **M/M** employed in *data sharing* scenarios [10]. Symmetric key schemes are typically **S/S**, while public key schemes **M/S** (multiple users can write using a *public key*, but only the owner of

the *private* key can generate trapdoors). Extension of the single reader schemes to multiple reader schemes can be done with *key sharing* or with the usage of *proxy re-encryption*.

In proxy re-encryption cryptography, a proxy is used to re-encrypt a ciphertext, encrypted using key K , and generates a new ciphertext, that is decrypted using key K' , where $K \neq K'$ and the decrypted plaintext remains unchanged. Proxy re-encryption is usually employed with public key cryptography [8, 41, 24], but there are some schemes that support the symmetric key setting [91, 81, 72].

2.2.3.4 Information Leakage

Almost all SE techniques have privacy issues because they leak some information [10]. This leakage can be *index* information, the *search pattern*, or *access pattern* during execution:

- *Index Information*: refers to the information about the index keywords. It may include the number of keywords per document or database, the total of documents, document similarity, their length and metadata;
- *Search Pattern*: includes the information about keywords, and predicates, that is inferred by correlating equal results from different searches. The search pattern can be targeted using statistical analysis;
- *Access Pattern*: includes the information that may be derived from the query results, based on the memory accesses needed to retrieve the documents. (e.g. one can infer whether a query is more restrictive than another, based on the total of documents retrieved; or if some files are more accessed than others).

2.2.4 Symmetric Searchable Encryption

2.2.4.1 Overview

SSE is a type of SE based on symmetric key cryptography. We will start by specifying the security definitions used to evaluate such schemes and then describe a few reference schemes.

2.2.4.2 Security Definitions

Indistinguishability Against Chosen Plaintext Attacks (IND-CPA) is a property commonly used to evaluate encryption schemes, particularly in the context of public-key or symmetric-key cryptography. Informally, a scheme is **IND-CPA** secure if an adversary A , given two arbitrary plaintext messages and their respective ciphertexts, cannot distinguish which ciphertext corresponds to each message with a probability higher than that of random chance. This probability upholds even if A can adaptively query the encryption oracle.

This definition is not appropriate for SE schemes where the leakage of information originates from trapdoors and queries.

In 2006, Curtmola et al. [21] revisited the security definitions for SSE schemes, contributing with two new definitions, which are still the standard in use: **Indistinguishability Against Non-Adaptive Chosen Keyword Attack (IND-CKA1)** and **Indistinguishability Against Adaptive Chosen Keyword Attack (IND-CKA2)**.

Both require that only the outcome of search patterns may be leaked. Additionally, they guarantee that trapdoors do not leak information about the keywords, apart from

what can be inferred from the search and access pattern. Their main difference is whether the adversary can produce an adaptive or non-adaptive attack.

- *IND-CKA1*: guarantees security if the client generates all queries at once;
- *IND-CKA2*: guarantees security even if the client chooses the queries based on previously obtained trapdoors and search outcomes.

2.2.4.3 SSE Schemes

Song et al. In 2000, the first practical **SSE** scheme was proposed [87]. The general basis of their method is to encrypt each word of a document and embed hash values with a specific format in the ciphertext. During search, the server executes a *sequential scan* of the document, extracts the hashes and checks for matches on their values.

First, the plaintext is divided into *fixed-size* blocks w_i and each one is encrypted with a *deterministic*⁴ encryption algorithm Enc , using key k . The ciphertexts $X_i = Enc_k(w_i)$ are split into two parts $X_i = \langle l_i, r_i \rangle$. Using a *stream cipher*⁵ as a **Pseudo-Random Function (PRF)**, a pseudo-random string of bytes $S_i = PRF_{k'}(i)$ is generated from the index position using key $k' \neq k$. Then, a hash $Y_i = \langle S_i, F_{k_i}(S_i) \rangle$ of this string of bytes is computed using a keyed hash function F_{k_i} , where the key is the output of a pseudo-random function $k_i = PRF_{k''}(l_i)$ that takes l_i as an input, with key $k'' \neq k'$. Finally, we compute the XOR of X_i with S_i to generate the final ciphertexts $C_i = X_i \oplus Y_i$.

To search, we generate a trapdoor $T = \langle X, K_s \rangle$, where $X = E_k(w) = \langle L, R \rangle$ is the encryption of keyword w and $K_s = PRF_{k''}(L)$. A server searches for the keyword in a document by iterating over all ciphertexts C_i , computing $C_i \oplus X$ and checking if the output is of the form $\langle S, PRF_{K_s}(S) \rangle$ for some S .

A security analysis of the scheme reveals that it leaks keyword position in the document. It is also susceptible to statistical attacks (e.g. after various searches, the words in the document can be inferred).

Curtmola et al. In 2006, two new schemes, for non-adaptive and adaptive adversaries, were proposed [21]. Each one is, respectively, at least **IND-CKA1** and **IND-CKA2** secure. The idea behind these was to use an *inverted index*, which is an index per distinct word in the database, not per document.

Non-Adaptive Scheme Considering a database with n documents, we define two data structures: an array A , where each position of the array contains a node $n_{i,j}$ which belongs to the linked list L_i associated to the keyword w_i ; a look-up table T to identify the first node of a linked list associated with a keyword.

To build A , we start by constructing the nodes. These nodes are encrypted, stored in random order (using permutations) and augmented with information to find and decrypt the next node in the list. The augmentation step happens before the encryption of the node. Each node $n_{i,j} = \langle id_m, k_{next}, p_{next} \parallel \emptyset \rangle$, includes information about the id of document m ($0 \leq m \leq n$) which contains the keyword, a randomly generated key used to encrypt the next node and a pointer to the next node in the list, or $\emptyset$ if it is the last one. The randomized disposition of the nodes, hides the length of the lists.

The look-up table I associates each keyword w_i with the node $n_{i,0} = \langle k_{i,1}, p_{i,1} \rangle \oplus PRF_k(w_i)$ stored in a random position using a **Pseudo-Random Permutation (PRP)**,

⁴Deterministic encryption leaks message equality.

⁵For information on stream ciphers we recommend reading chapter 8.4 of [71].

$PRP_{k'}(w_i)$, where $k \neq k'$. This node contains the encryption key and the pointer to the first node of the linked list L_i .

To search, we generate trapdoor $T_w = \langle PRF_k(w), PRP_{k'}(w) \rangle$. The server uses the trapdoor to find the correct node $n_{i,0} = T[PRP_{k'}(w)]$ and then retrieve all the identifiers of the documents that contain the keyword.

Adaptive Scheme Considering a database with n documents, we now define only a look-up table I , but instead of generating the index for a single keyword w , we generate it for the family of the keyword. The family of a keyword is defined as a set of labels $F_w = \{w|str(j) : 1 \leq j \leq t\}$ where j is an integer converted to string, using str , t is the number of identifiers of documents containing w , and the label is the result of concatenating w with j . The look-up table associates each label with the identifier of a document containing the w , in a random position, using $PRP_{k'}(w_i|str(j))$. To hide differences in the number of labels per keyword, extra entries are padded.

We generate the trapdoor $T_w = (t_1, \dots, t_n) = (PRP_k(w|str(1)), \dots, PRP_k(w|str(n)))$ to search. The server uses the trapdoor values to find valid index positions in I and outputs the set of identifiers.

Generalization to the Dynamic Setting The previous schemes were considered static, as they do not allow modifications on the database after setting up the scheme. In 2012, Kamara et al. [44] proposed a *dynamic* extension to the *non-adaptive* scheme in [21]. They used a *deletion* array to track the positions on the search array that needed to be modified, after an update. They encrypted the array pointers, using homomorphic encryption, to support modifications on their values without decryption.

Generalization to Structured Data In 2010, Chase and Kamara [15] proposed a new adaptively secure scheme based on the *non-adaptive* scheme in [21]. This new scheme supports arbitrarily-structured data. **Structured Symmetric Encryption** is based on generating an inverted index in the form of a padded and permuted dictionary. They associated data items with metadata, allowing information about structure to be preserved if needed.

Generalization to Graph-Structured Data After **STE**, additional research has been conducted to support operations on graph-structured data.

Some authors extended **STE** to support queries on conceptual graphs [74], but they need to compute all possible query answers *a priori*.

Others, proposed a **SSE** scheme that supports *subgraph isomorphism queries* [12]. In their approach, they extract features from the graphs and generate the encrypted-indexes from the extracted features. This scheme was planned for generic graphs.

Finally, in [89] they combine an index-based **SSE** with a bucketization algorithm. By doing so, queries do not leak information about the structure of the graph. Although the authors state that their approach works with any kind of query, they only focused on neighbor and adjacency ones.

2.3 Secure Semantic Data Solutions

One of the simplest works in the field focuses solely on protecting individuals in **OWL2** ontologies [62], uniquely encrypting their values and leaving the rest of the data in plaintext. This approach aimed at retaining functionality but leaks too much information. The following solutions try to achieve more robust security guaranties.

2.3.1 Self-Enforcing Access Control for Encrypted RDF

This work researched the implementation of access control in RDF datasets [28] through the use of functional encryption [9].

Encryption Scheme In the proposed encryption scheme, functions correspond to the computation of inner-products over $\mathbb{Z}_N$. The set of possible ciphertext attributes of length n is $\Sigma = \mathbb{Z}_N^n$ and the class of decryption key functions is $F = \{f_{\vec{x}} \mid \vec{x} \in \mathbb{Z}_N^n\}$. An attribute vector $\vec{y} \in F$ is associated with each ciphertext. A decryption key is a vector $\vec{x}$ such that $f_{\vec{x}}(y) = 1$ iff $\sum_{i=1}^n y_i x_i = 0$

To encrypt an RDF triple $t_i = (s_i, p_i, o_i)$ we functionally encrypt a random seed m_{t_i} , with an attribute vector $\vec{y}_{t_i} = (y_{s_i}, y'_{s_i}, y_{p_i}, y'_{p_i}, y_{o_i}, y'_{o_i})$ such that $\vec{y}_l = -r \cdot \sigma(l)$, $y'_l := r$, $l \in \{s, p, o\}$ with σ being a function that maps a triple's subject, predicate and object to a value in $\mathbb{Z}_N$, and r a random value also in $\mathbb{Z}_N$

Advanced Encryption Standard (AES)⁶ encryption is used to encrypt the plaintext triple using a key derivable from the generated seed m_{t_i} . The final ciphertext $c_{t_i} = \langle AES(t_i), FE(m_{t_i}) \rangle$ corresponds to both the AES ciphertext and the ciphertext of the seed.

The decryption keys are generated taking into account the possible triple patterns $tp = (s, p, o)$ combinations. For each triple pattern we define $\vec{x}_{tp} = (x_s, x'_s, x_p, x'_p, x_o, x'_o)$ such that $x_l := 1$, $x'_l := \sigma(l)$ if l is bound and $x_l := 0$, $x'_l := 0$ otherwise.

Finally, to decrypt a triple $t = (s_t, p_t, o_t)$, given a triple pattern $tp = (s_{tp}, p_{tp}, o_{tp})$ we compute the inner product $\vec{y}_t \cdot \vec{x}_{tp} = y_{s_t} x_{s_{tp}} + y'_{s_t} x'_{s_{tp}} + y_{p_t} x_{p_{tp}} + y'_{p_t} x'_{p_{tp}} + y_{o_t} x_{o_{tp}} + y'_{o_t} x'_{o_{tp}}$

Only if the result of the inner product is 0 can we decrypt the seed, derive the correct AES key and decrypt the triple.

Data Structure The data is stored in a key-value structure where the keys are unique integer IDs and the values correspond to the encrypted triples. The IDs are indexed using two strategies.

The first strategy is a *3-Index* approach. We compute secure hashes $h(s)$, $h(p)$ and $h(o)$ for each s, p, o (unbound values correspond to $h(?)$). Then, we generate three index tables *SPO*, *POS* and *OSP* where we store all triple pattern combinations and associate each value to the corresponding ID.

The second strategy is *vertical partitioning*. Instead of indexing the permutations of subjects, predicates and objects. We create one table per predicate, indexing the subject, object pairs that are related via that predicate. Each ID is indexed by a ("*so*" | "*os*", s, o).

T³

3) TODO: Add image with both index schemes

2.3.2 CryptGraphDB

CryptGraphDB [2] is an adaptation of the *CryptDB* [77, 76, 75] system to graph-structured data. It is capable of answering cypher [63] queries over conceptual graphs. We will start by describing the *CryptDB* architecture, security guaranties, query execution pipeline and encryption schemes. We will then explain the adaptations employed by *CryptGraphDB*.

System Architecture *CryptDB* is a practical solution for processing SQL queries over encrypted relational datasets. This solution leverages onion-encryption to support various

⁶For information on AES we recommend reading chapter 6 of [71].

primitive operations necessary to answer a [SQL](#) query (equality checks, order comparisons, aggregates and joins).

The system is composed by an application server, a proxy and an unmodified [Database Management Systems \(DBMS\)](#). The proxy intercepts all queries, encrypts and decrypts all data, and changes query operators while preserving the semantics of the original query.

T⁴

4) *TODO: Add image of Cryptdb system*

Security Threats *CryptDB* addresses two security threats. The first threat, consists in a passive attacker (e.g. database administrator), which only observes confidential data in the [DBMS](#) server. Against this attacker *CryptDB* provides a few security guaranties: sensitive data is always encrypted; the [DBMS](#) server cannot generate results for queries whose primitives do not belong in the subset of supported operations; information leakage depends on the types of computations that will be executed over the data - if no relational predicate filtering is needed, nothing more than the data's length is leaked; on equality checks, it reveals the histogram of encrypted values in the respective columns; in order checks, the system reveals the order between the encrypted values in the column.

The second threat is an attacker that gains control over the proxy, application server and [DBMS](#) server. Against this adversary, *CryptDB* provides confidentiality of data belonging to users not logged in during the period of the compromise (until it has been fixed). The adversary can, however, delete data stored in the database.

Besides confidentiality, *CryptDB* does not ensure integrity, freshness, or completeness of results.

Query Execution Pipeline For executing a query issued by the application server, a sequence of steps are necessary: i) the proxy intercepts and rewrites the query, anonymizing the table and column names and encrypting each constant present with the adequate encryption scheme for the desired operation, using a master key; ii) the proxy checks whether it should provide keys to the [DBMS](#) server, so that adjustments can be done on the encryption layers of the targeted data, if the check is positive it issues an UPDATE query that triggers a custom function to adjust the encryption of the selected columns; iii) the encrypted query is forwarded to the [DBMS](#) server and is executed using plain [SQL](#); iv) the encrypted results are sent to the proxy, which are then decrypted and forwarded to the application server.

Adjustable Encryption Schemes To avoid unnecessary leakage of information data is encrypted with different layers, like an onion. When operations need encryption schemes with lower security guaranties, the data encryption layers are peeled. For each layer, per onion, per column, the proxy uses the same encryption key, derived from a master key, *MK*. This reduces the leakage of relations between columns and avoids the complexity of having different keys for each row in a column. For table *t*, column *c*, onion *o*, and encryption layer *l*, the proxy would use the encryption key $K_{t,c,o,l} = PRP_{MK}(t, c, o, l)$.

We now present the different encryption schemes supported by *CryptDB*:

- *Random (RND)*: A probabilistic encryption scheme that provides the most security but does not allow any computations on the ciphertexts.
- *Deterministic (DET)*: A deterministic encryption scheme which allows equality checks to be executed over ciphertexts.

- *Order-Preserving (OPE)*: An encryption scheme that preserves order between the ciphertexts in a column. This allows the server to execute range queries.
- *Homomorphic encryption (HOM)*: The Paillier cryptosystem [69] is used to probabilistically encrypt numeric values. This scheme allows additive operations to be executed with ciphertexts.
- *Join (equi-JOIN)*: In *CryptDB* columns are encrypted with a different encryption keys. Essentially, joins can not be performed between columns until the encryption keys of the columns coincide. To support join operations between different columns a new primitive is introduced JOIN-ADJ. It can be interpreted as a *keyed random hash* with the property that the hashes can be adjusted to modify their key without access to the plaintext, using an adjustment key $\Delta K = k_{c_1}/k_{c_2}$. The adjustment of the columns is done only on one of the two (the first column in lexicographic order). The information leakage is similar to that of the DET layer.
- *OPE-JOIN*: *CryptDB* does not provide a construct to dynamically modify the keys of columns encrypted with the *OPE* scheme. To support range joins, these columns have to be specified before, so that they are encrypted under the same key.
- *Word search (SEARCH)*: This encryption scheme allows queries with the LIKE keyword. In *CryptDB* they used the scheme by Song et al. [86]. The information leakage is similar to that of the RND layer.

The following images provide, respectively, an overview of the different onions types in *CryptDB* and an example a query execution where layer adjustments are needed.

T⁵ **T⁶**

CryptGraphDB constructs upon the approach of *CryptDB*, adapting the solution to be able to answer *cypher* queries [63]. The proxy is implemented to handle *cypher* queries instead of *SQL* queries, but keeps the same functionality. Furthermore, as the *cypher* language does not require the join operator, *CryptGraphDB* only maintains the RND, DET, OPE, HOM and SEARCH encryption schemes.

2.3.3 GOOSE

GOOSE [17] is a secure graph outsourcing framework capable of answering *UCRPQ* queries. The solution relies on secure multiparty computation techniques, encryption schemes and a standard *SPARQL* engine.

GOOSE was designed as a distributed system with 5 participants: a *data owner* (DO), who owns the outsourced graph G ; a *user* (U), who submits graph queries against G and receives query answers; a *query translator* (QT); *answer translator* (AT); and a *SPARQL engine* (SE).

T⁷

Preliminaries Let $ENC_{A \rightarrow B}$ or $ENC_{B \leftarrow A}$ denote *AES-CBC* encryption using a shared key between participants A and B.

Let the outsourced *RDF* graph be $G = (N, E)$ where N is a set of nodes and $E \in N \times \Sigma \times N$ is a set of directed edges between the nodes, labeled using the alphabet Σ . Let $\sigma_N : N \rightarrow \{0, \dots, |N| - 1\}$ and $\sigma_\Sigma : \Sigma \rightarrow \{0, \dots, |\Sigma| - 1\}$ be two random bijections, respectively, for the edges and nodes, and let σ^{-1} represent the inverse of σ . The hidden set of edges from E is defined as $\hat{E} = \{(\sigma_N(s), \sigma_\Sigma(p), \sigma_N(o)) \mid (s, p, o) \in E\}$.

5) TODO: Add query flow image

6) TODO: Add onion encryption image

7) TODO: Add protocol image

Consider $\Sigma^+ = \{a, a^- \mid a \in \Sigma\}$, a^- denotes the inverse of the edge label a . Let $V = \{?v_0, \dots, ?v_n\}$, with $n \in \mathbb{Z}^+$, be a set of variables. A *query rule* is an expression of the form $head \leftarrow body$ defined as $(\overline{?v}) \leftarrow (?v_0, r_0, ?v_1), \dots, (?v_{n+1}, r_n, ?v_{n+2})$, where, for every $i \in \mathbb{Z}_0^+$, r_i is a regular expression over Σ^+ using $\{., +, *\}$ and $\overline{?v}$ is a vector composed by a subset of the variables from the *body*. The length of this vector is the *arity* of the rule.

A **UCRPQ** query q is finite and non-empty set of *query rules* with the same *arity*. The answer of a query q evaluated against a graph G , using standard **SPARQL** semantics, is denoted as $ans(G, q)$.

Graph outsourcing To outsource a graph, the *data owner* generates σ_V , σ_Σ and $\hat{E}$. It then sends $ENC_{DO_QT}(\sigma_\Sigma)$ to the *query translator*, $ENC_{DO_AT}(\sigma_V)$ to the *answer translator* and $ENC_{DO_SE}(\hat{E})$ to the **SPARQL** engine.

Query Evaluation In **GOOSE**, query evaluation is a four steps process:

1. First, the *user* generates a query q and sends $ENC_{U_QT}(q)$ to the *query translator*.
2. The *query translator* generates $\hat{q}$ by applying the σ_Σ function to all predicates in q . $ENC_{QT_SE}(\hat{q})$ is sent to the **SPARQL** engine.
3. The **SPARQL** engine evaluates $\hat{q}$ over $\hat{G} = (\bigcup_{(\hat{s}, \hat{p}, \hat{o}) \in \hat{E}} \{\hat{s}, \hat{o}\}, \hat{E})$ obtaining the answer $ans(\hat{G}, \hat{q})$. It then sends $ENC_{SE_AT}(ans(\hat{G}, \hat{q}))$ to the *answer translator*.
4. The *answer translator* deobfuscates $ans(\hat{G}, \hat{q})$ using σ_V^{-1} and generates the final answer $ans(G, q)$. It then sends $ENC_{AT_U}(ans(G, q))$ to the *user*.

Security Model **GOOSE** is designed to be secure against *honest-but-curious* adversaries. It is assumed that data and computations are outsourced to external servers. The solution is expected to have the following security properties, which have been formally proven by the authors:

1. No cloud node can learn G .
2. Let q be a query submitted by a user and $ans(G, q)$ the answer of the query on the graph G . No cloud node can learn simultaneously q and $ans(G, q)$.
3. An external observer cannot learn G , any query q nor answer $ans(G, q)$ by analyzing the network messages.

2.3.4 Discussion

Current solutions limitations In this section we covered a few solutions for secure outsourcing of semantic data.

The *CryptGraphDB* [2] was not designed to support the **SPARQL** query language and suffers from the same security leakage problems inherent to *CryptDB*. Some authors argue that the usage of deterministic encryption to support join operations leaks too much information (e.g. frequency information about the encrypted items) and is vulnerable to *leakage-abuse* attacks [43, 38, 83].

The solution from J. D. Fernández et al. [28] uses hashes as indexes and this leaks the structure of the **RDF** graph. Furthermore, it only supports search operations over simple *triple patterns*.

The [GOOSE](#) [17] solution can answer [SPARQL](#) queries with complex *graph patterns*, but the data obfuscation methods used are deterministic and still leak the topology of the [RDF](#) graphs and frequency information of the encrypted keywords.

Furthermore, none of these solutions offer reasoning capabilities, and it is not clear whether they support write operations besides a single initial upload of the data.

Orthogonal security research Like the previous solutions, the main focus of our thesis will be on data *confidentiality*. Nonetheless, we will briefly mention the research that is being done in other security fields regarding [RDF](#) graphs.

One of those fields is ensuring *integrity* and *authenticity* of [RDF](#) data. Various signage schemes [13, 97, 45, 46, 47, 90] and even one service-oriented solution [53] have been proposed.

The other field is *access control*. Most solutions try to embed the policies within the data, using annotations. This allows for high-granularity in the *access control*, which can be inferred, but requires conflict resolution [42, 1, 30, 70, 18, 48, 19].

2.4 Summary

In Section 2.1, we described the layered architecture of the *Semantic Web* and discussed the standard languages for describing its *data model* and *ontologies*. More specifically, we covered the semantics and syntax of [RDF](#) and explained the primitives of [RDFS](#).

Additionally, we specified the structure of [OWL2](#), focusing on the mechanism it provides to write ontologies. We also discussed the restriction imposed by *OWL DL* to remain decidable.

Furthermore, we explained [SPARQL](#), providing mathematical definitions, defining its semantics, describing the syntax, enumerating a few examples of queries and reasoning preservation techniques.

We finished Section 2.1 by Listing a few of the tools, software, storage solutions and evaluation benchmarks used in this field of research.

In Section 2.2, we discussed various solutions to execute secure computations in untrusted servers. We started by describing the *honest-but-curious*, or *passive*, adversary model and explained why *Encryption at Rest* and [ORAM](#) are not practical approaches.

We proceeded to cover [Homomorphic Encryption](#) schemes, informally mentioning the mathematical basis of [HE](#) and enumerating the type of schemes available, and property-preserving encryption schemes.

In Section 2.2.3, we explained the difference between [SSE](#) and [PEKS](#). We then described the general model for index-based [SE](#), as well as the *Client/Server* architecture of [SE](#) systems. We continued by discussing how to extend [S/S](#) architectures, covered the types of privacy leakage.

In Section 2.2.4 we specified the standard security definitions in use, and gave an in depth explanation of reference [SSE](#) schemes. We finished the [SSE](#) Section by showing recent generalizations, focusing on dynamic solutions and schemes that support structured data.

In Section 2.3, we mentioned related research on the topic of secure semantic data sharing, discussing the trade-off of each approach. Additionally, we briefly covered parallel work in the fields of data *integrity*, *authenticity* and *access control*.

SEnTriS: A SEARCHABLE ENCRYPTED TRIPLESTORE SHARING PROTOCOL

In this chapter we will present the main contribution of our research. In Section 3.1, we introduce SEnTriS and define the system and adversary model. In Section 3.2 we define: 1) two protocols for SPARQL query evaluation, SPARQL-EVAL and SPARQL-EVAL-II; 2) a triple pattern SSE scheme, SSE-TP; 3) a SPARQL query evaluation protocol that functions with encrypted triplestores, SSE-SPARQL-EVAL; a IAM protocol used to adapt the solution to the M/M setting; 4) the SEnTriS protocol.

3.1 Solution Models

3.1.1 Overview

The current *state-of-the-art* solutions solely focus on the problem of secure RDF data outsourcing, leaving the problem of dynamic updates out of scope. Furthermore, these systems do not have reasoning capabilities, suffer from leakage problems and, frequently, solely provide search over simple *triple patterns*.

Considering all these limitations we developed SEnTriS, a solution for secure and cooperative development of semantic knowledge bases.

3.1.2 System Model

SEnTriS possesses RBAC and can execute secure read and write operations. ^{T⁸} It supports most SPARQL operations and retains reasoning capabilities while evaluating queries.

SEnTriS is designed for a multi-user setting. It is purely a software solution, not requiring any special hardware, where storage and heavy computations are outsourced to external servers.

In SEnTriS, authenticated *users* are able to create *encrypted triplestores*. Within a triplestore they are able to upload, modify and query semantic data. Triplestores can be shared between users with varying degrees of access.

SEnTriS follows the usual SE schemes architecture. It is a *client/server* application which outsources data storage and heavy computations to third parties. It follows the M/M model with *secrets sharing* between trusted users.

8) *TODO*: Do not forget to mention verification of authenticity and data integrity as future work

3.1.3 Security Model

We will consider an *honest-but-curious* adversary. This adversary will try to eavesdrop communication between servers. If it has access to the external servers where the semantic data and execution logs are stored, it will also try to infer or decrypt sensitive information. The leakage of the protocols in use will highly influence the effectiveness of the attacks.

We will consider that communication channels are trusted, reliable and secure. Additionally, we will assume that authenticated users are trusted and non-malicious. Such users can read sensitive information and security parameters of the triplestores they have created or have been given access to.

Finally, we will not consider the case where attacks involve *revoked users*. The RBAC prevents these users from interacting (via client) with triplestores they do not have permissions to. Nevertheless, we have not implemented a key rotation, or re-encryption protocol for data. As such, an adversary collaborating with these users, or an attack from such users, will always succeed in the decryption the data of all the triplestores they used to have access to. They can still possess the cryptographic secrets used by the protocols, which makes it trivial to decrypt leaked information.

3.2 Solution Definition

3.2.1 Overview

The functioning of a query evaluation system can be described as a pipeline of sequential operations. Firstly, if the system supports authentication, there is an authentication step. An authenticated user will be able to submit a series of queries while their session is active. A submitted query is parsed, and given the specified heuristics, the system generates a series of query execution plans, which consist in sequences of logical operations (e.g. data retrieval, update or execution of an equijoin). After a query plan is chosen and executed, the final results are sent to the user, with the applied solution modifiers.

SEnTriS follows a similar logic, for query evaluation, with a few modifications to enhance security. The system works with data that is encrypted in a manner that allows for search and set operations to remain executable. Additionally, before each step of the pipeline, access checks are performed against the users' ephemeral access tokens, to ensure that only the right actors are performing actions within the system. ^{T⁹}

SSE techniques are used during query evaluation: we generate trapdoors to search encrypted subjects, predicates and objects present in *triple patterns*. With these, we build *bindings tables*, which are used for the evaluation of complex *graph patterns*

SEnTriS is implemented using two different approaches:

- *SEnTriS* V1: Data encryption retains elements from *CryptDB* [77, 76, 75]. Subjects, predicates and objects are encrypted in two layers. The outer layer is probabilistically encrypted, whereas the second is deterministically encrypted. When equality checks are required, we peel the outer layer and compare the values.
- *SEnTriS* V2: Implemented using an additive homomorphic encryption scheme. The value of a subject, predicate or object is mapped to an *equality tags*. These values are homomorphically encrypted and support re-encryption. Trapdoors index the *equality tags* and equality tags reference the probabilistically encrypted subjects, predicates and objects. The *bindings tables* are built with *re-encrypted equality tags*. When doing equality checks, we decrypt the difference between *equality tags*.

9) TODO: Add pipeline figure to illustrate differences

In the remainder of this section we will: 1) define the core data elements and cryptographic tools; 2) specify the [SPARQL](#) query evaluation protocols; 3) define the triple pattern [SSE](#) schemes, and the adaptations done to the [SPARQL](#) query evaluation protocols; 3) define the [Identity and Access Management \(IAM\)](#) protocol, explaining how triplestores (and their secrets) are shared and how to support concurrent access; 5) define the [SEnTriS](#) protocol.

3.2.2 Core elements

We will denote $E(m)$ as the encryption of m , $O(m)$ as the obfuscation of m , mapping it to a random value.

Triplestores A [RDF](#) dataset within [SEnTriS](#) is represented as a single named triplestore. A triplestore $T = (triplestoreID, nodes, acl)$ is a key value store, where $triplestoreID$ is a unique string that identifies the dataset, $nodes = \{(idx_0, node_0), \dots, (idx_n, node_n)\}$ is a map of predicate, subject or object $nodes$, $node \in I \cup B \cup L$, and each node is uniquely identified by its pattern index, idx , which is a hash of the *triple pattern*. The acl is the triplestore's [Access Control List \(ACL\)](#).

This data structure is a simplified version of the one defined in section 2.1.4, where we only represent the *default* graph.

Triplestore Access Control List A triplestore's [ACL](#) $acl = (owner, read, write, acrl)$, is a tuple where *owner* is the username of the owner of the triplestore, *read* and *write* are both collections of usernames and *acrl* is a collection of *access change requests*. An access change request $accessRequest = (requestID, username, write)$ is a tuple where $requestID$ is a [Universally Unique Identifier \(UUID\)](#) (version 4) [50], *username* identifies the user that this request affects and *write* is a boolean, that specifies if the user desires write access (true) or read access (false).

Encrypted Triplestores An encrypted triplestore $E(T) = (triplestoreID, E(nodes), acl)$ is a key value store, where $triplestoreID$ is a unique string that identifies the dataset, $E(nodes)$ is an encrypted map of predicate, subject or object $nodes$ and the acl is the triplestore's [ACL](#). $E(nodes)$ will be defined in Section 3.2.5

Users A user $u = (username, password, owned, role)$ is a tuple, where *username* and *password* are unique strings, $owned = [triplestoreID_0, \dots, triplestoreID_n]$ is a collection of user's owned triplestores identifiers and *role* is the user's role.

Role Change Request A role change request $roleRequest = (requestID, username, role)$ is a tuple where $requestID$ is a [UUID](#) (version 4), *username* identifies the user that this request affects and *role* specifies the desired new role.

Sessions A user's session $s = (sessionID, username)$ is a tuple, where $sessionID$ is a [UUID](#) (version 4), *username* belongs to a valid registered *user*. Sessions remain active until they expire, or are terminated. The duration of a session will be configurable.

Locks A lock is $l = (lockID, resourceID)$ is a tuple, where $lockID$ is a [UUID](#) (version 4) and *resourceID* is either a valid *username* or *triplestoreID*. Locks remain valid until they expire, or are released. The duration of a lock will be configurable.

Ephemeral Access tokens An access token $at = (tokenID, username, sessionID, triplestoreID, lockID)$ is a tuple, where $tokenID$ is a [UUID](#) (version 4), $username$ is the identifier of the user who holds the token, $sessionID$ is the identifier of the token holder's active session, at the time of its creation, $triplestoreID$ is the identifier of the triplestore whose access is being requested and $lockID$ can either be $\perp$ (for *read* access requests) or the identifier of the current lock, against the targeted triplestore (for *write* access requests). Access tokens are ephemeral and will expire after a certain amount of configurable time.

Secret Maps A secret map $smap = (triplestoreID, secrets)$ is a tuple where $triplestoreID$ is the identifier of an existing triplestore, $secrets = \{(key_0, secret_0), \dots, (key_n, secret_n)\}$ is a map and key_i is a unique string that identifies $secret_i$, the serialized value of a secret parameter.

SPARQL queries The definition of a [SPARQL](#) query within [SEnTriS](#) is equal to that provided in section [2.1.4](#).

Query Execution Plans A query execution plan is a tuple $plan = (jobs, execOrder)$, where $jobs = \{(jobID_0, job_0), \dots, (jobID_n, job_n)\}$ is a map that maps a query job to their respective identifiers, which are [UUIDs](#) (version 4), and $execOrder$ is a sequence of job identifiers.

Randomized Query Execution Plans A randomized query execution plan is a tuple $plan_{rnd} = (jobs, rndExecOrder)$, where $jobs = \{(jobID_0, job_0), \dots, (jobID_n, job_n)\}$ is a map that maps a query job to their respective identifiers, which are [UUIDs](#) (version 4), and $rndExecOrder$ is a randomly chosen sequence of job identifiers, from the subset of valid execution sequences.

Randomized Prepared Query Execution Plans After the query preparation phase, the *searches* maps, in the secure search jobs, have been updated to map obfuscated variables to search identifiers. A plan where all search jobs have been updated like this is said to be *prepared*.

Query Jobs A query job is an internal operation that is executed during the evaluation of a [SPARQL](#) query. Query jobs can be of three types: *basic*, *unary* and *binary* jobs. Unary and binary jobs represent operations that are executed against the *bindings* of, respectively, one and two other jobs, whereas basic jobs are independent of other jobs' results. The execution of a job will always generate *bindings*, even if it is an empty collection. These *bindings* will be associated to the respective *jobID*.

There exist two basic jobs, job_{SEARCH} and job_{EMPTY} .

The search job is a tuple $job_{SEARCH} = (jobID, varPattern, s, p, o)$, where $varPattern \in \{S, P, O, SP, SO, PO, SPO\}$ and s, p, o are $\perp$ or have the value of, respectively a subject, predicate or object of the *triple pattern* being searched.

The empty result job is a tuple $job_{EMPTY} = (jobID, vars)$, where $vars = [var_0, \dots, var_n]$ is a collection of variables.

There exist four unary jobs, which represent the operations that apply result modifiers. These jobs are tuples of type $job_I = (jobID, jobID_{prev})$, where $I \in \{PROJECT, ORDERBY, SLICE, DISTINCT\}$ and dictates which solution modifier will be applied to the *bindings* associated with $jobID_{prev}$.

Finally, there are four binary jobs, which represent set operations between *binding*. These are tuples of type $job_{II} = (jobID, jobID_l, jobID_r)$, where $II \in \{JOIN, UNION, OPTIONAL, MINUS\}$ and specifies which set operation will be executed between the *bindings* of $jobID_l$ and $jobID_r$.

Secure Query Jobs These jobs have their variables obfuscated and output encrypted *bindings* tables.

Only the *secure search job* is structurally different. It is a tuple $job_{SEC_SEARCH} = (jobID, vars, searches)$, where $vars = [O(var_0), \dots, O(var_n)]$ is a collection of obfuscated variables. During the query preparation phase, $searches = \{(O(var_0), keyword_0), \dots, (O(var_n), keyword_n)\}$ is a map that maps variables to the keywords to be searched, whereas, in the query evaluation phase, $searches = \{(O(var_0), searchID_0), \dots, (O(var_n), searchID_n)\}$ is a map that maps variables to the identifiers of search results. These identifiers are **UUIDs** (version 4).

Bindings tables The output of a query job is a *bindings* table. This data structure is a tuple $bindingsTable = (bindings, patterns)$ where $bindings = \{(var_0, bindingsMap_0), \dots, (var_n, bindingsMap_n)\}$ and $patterns = \{(var_0, patternsMap_0), \dots, (var_n, patternsMap_n)\}$ are maps that map query variables, respectively, to *bindings* and *triple patterns* maps.

A *bindings* map, $bindingsMap = \{(binding_0, [patternID_{0,0}, \dots, patternID_{0,k}]), \dots, (binding_n, [patternID_{n,0}, \dots, patternID_{n,k}])\}$, is a map where *bindings* are mapped to a list of *triple pattern* identifiers, which are **UUIDs** (version 4).

A *triple patterns* map, $patternsMap = \{(patternID_0, binding_0), \dots, (patternID_n, binding_n)\}$, is a map where *triple pattern* identifiers are mapped to *bindings*.

Encrypted Bindings tables Similar to a *bindings* table, an encrypted *bindings* table is a tuple $E(bindingsTable) = (E(bindings), E(patterns))$ where $E(bindings) = \{(O(var_0), E(bindingsMap_0)), \dots, (O(var_n), E(bindingsMap_n))\}$ and $E(patterns) = \{(O(var_0), E(patternsMap_0)), \dots, (O(var_n), E(patternsMap_n))\}$ are maps that map obfuscated query variables, respectively, to encrypted *bindings* and *triple patterns* maps.

An encrypted *bindings* map, $E(bindingsMap) = \{(E(binding_0), [patternID_{0,0}, \dots, patternID_{0,k}]), \dots, (E(binding_n), [patternID_{n,0}, \dots, patternID_{n,k}])\}$, is a map where encrypted *bindings* are mapped to a list of *triple pattern* identifiers, which are **UUIDs** (version 4).

An encrypted *triple patterns* map, $E(patternsMap) = \{(patternID_0, E(binding_0)), \dots, (patternID_n, E(binding_n))\}$, is a map where *triple pattern* identifiers are mapped to encrypted *bindings*.

We will define $E(binding)$ in Section 3.2.5 for both **SEnTriS V1** and **SEnTriS V2**.

SPARQL Query Result A **SPARQL** query result is a tuple $sparqlResult = (bindings, distinct, ordered, sliced, sortConditions, length, offset)$ where $bindings = [pattern_0, \dots, pattern_n]$ is a collection of *triple patterns*, *ordered*, *distinct* and *sliced* are boolean values, *length* and *offset* are integers and $sortConditions = [(var_0, dir_0), \dots, (var_n, dir_n)]$ is a collection of variables and their respective sorting direction, which can either be 1 or -1. A positive *dir* represents an *ascending* sorting whereas a negative *dir* represents a *descending* sorting condition.

Let $n \in \mathbb{Z}^+$ be the number of variables in a **SPARQL** query, a *triple pattern* $pattern_i = \{(var_0, binding_{i,0}), \dots, (var_n, binding_{i,n})\}$ is a map where each variable is mapped to a *binding*. A *binding* can be a plaintext value or $\perp$ for the particular case of **OPTIONAL** queries.

Encrypted SPARQL Query Result An encrypted [SPARQL](#) query result is a tuple $E(sparqlResult) = (E(bindings), distinct, ordered, sliced, sortConditions, length, offset)$ where $E(bindings) = [E(pattern_0), \dots, E(pattern_n)]$ is a collection of encrypted *triple patterns*, *ordered*, *distinct* and *sliced* are boolean values, *length* and *offset* are integers and $sortConditions = [(var_0, dir_0), \dots, (var_n, dir_n)]$ is a collection of variables and their respective sorting direction, which can either be 1 or -1 . A positive *dir* represents an *ascending* sorting whereas a negative *dir* represents a *descending* sorting condition.

Let $n \in \mathbb{Z}^+$ be the number of variables (obfuscated) in a [SPARQL](#) query, an encrypted triple pattern $E(pattern_i) = \{(O(var_0), E(binding_{i,0})), \dots, (O(var_n), E(binding_{i,n}))\}$ is a map where each obfuscated variable is mapped to an encrypted *binding*. An encrypted *binding* can either have a value or be $\perp$ for the particular case of OPTIONAL queries.

SPARQL Final Query Result A [SPARQL](#) query result *sparqlResult* is the result of evaluating a *WHERE* clause. The final output of a query evaluation must respect the standards present in [115, 116]. This final query result, *sparqlFinalResult*, will vary depending on the query type:

- *SELECT*: it will be a query result in JSON format;
- *ASK*: it will be a boolean;
- *CONSTRUCT*: it will be a set of triples;
- *DESCRIBE*: it will be a frequency table of the query variables;
- *SPARQL UPDATE*: it will be either SUCCESS, or $\perp$ in case of failure;

3.2.3 Cryptographic Tools

3.2.3.1 Notation

For the remainder of the document, the function $E(K, m, iv)$ denotes the encryption of message m , using the initialization vector iv and key K . The notation $E(data)$ will represent uniquely that *data* is encrypted. The function $D(K, ct)$ denotes the decryption of ciphertext ct , using key K . The keys used as input will identify the encryption scheme (e.g. K_A implies the usage of the symmetric encryption scheme A and $K_{B_{priv}}$ implies the usage of public key encryption scheme B).

3.2.3.2 DGK Encryption

[SEnTriS](#) V2 can be implemented with any additive homomorphic encryption scheme, but we have chosen to use the [DGK](#) encryption scheme [23, 22, 32, 52]. Compared to other schemes (e.g. Paillier [69]) it produces smaller ciphertexts, and the reduced plaintext and ciphertext spaces result in large performance gains. It also permits us to employ more efficient equality checks, which is a very common operation when evaluating [SPARQL](#) queries.

Definition 7 (DGK Encryption). Let k , t , and l be public parameters. Parameter k determines the size of a [Rivest-Shamir-Adleman \(RSA\)](#) modulus, l is the bit length of the messages to encrypt and t should be large enough so that exhaustive search is infeasible.

Key Generation Let p, q be two primes, generated such that $n = pq$ is a k -bit RSA modulus and that the u, v_p and v_q also exist and possess the following properties:

1. u is the smallest prime greater than $l + 2$
2. u is a divisor of both $p - 1$ and $q - 1$
3. v_p is t -bit prime numbers that divides $p - 1$
4. v_q is t -bit prime numbers that divides $q - 1$

Choose random values $g, h \in \mathbb{Z}_n^*$, so that:

1. the multiplicative order of h is $v_p v_q$
2. g is an element of order $uv_p v_q \in \mathbb{Z}_n^*$

The message space is $M = \mathbb{Z}_u$ and the ciphertext space is $C = \{g^m h^r \bmod n \mid m \in \mathbb{Z}_u, r \in \mathbb{Z}_{v_p v_q}\} \subseteq \mathbb{Z}_n^*$. The public and private keys are, respectively, $K_{DGK_{pub}} = (n, g, h, u)$ and $K_{DGK_{priv}} = (p, q, v_p, v_q)$.

Probabilistic Encryption An encryption ct of m is calculated as:

$$ct = g^m h^r \bmod n \quad (3.1)$$

The value $r \in \mathbb{Z}_{v_p v_q}$ is only possible to calculate if we have access to private key. Knowing the public key, we instead determine the value of r as a random $2.5t$ -bit integer, so that h^r is indistinguishable from a uniformly random element from the subgroup generated by h .

Deterministic Encryption To compute a deterministic encryption of m we remove the randomization part, $h^r \bmod n$, from the calculations:

$$ct = g^m \bmod n \quad (3.2)$$

Re-encryption To re-encrypt a probabilistic encryption of m we sample a new random $2.5t$ -bit integer r_2 and compute:

$$\begin{aligned} ct_2 &= ct \cdot h^{r_2} \bmod n \\ &= g^m h^r h^{r_2} \bmod n \\ &= g^m h^{r+r_2} \bmod n \end{aligned} \quad (3.3)$$

Decryption The decryption of a ciphertext ct is done by calculating:

$$\begin{aligned} c^{v_p v_q} \bmod n &= (g^m h^r)^{v_p v_q} \bmod n \\ &= (g^m)^{v_p v_q} (h^r)^{v_p v_q} \bmod n \\ &= (g^m)^{v_p v_q} * 1 \bmod n && \text{because order of } h \text{ is } v_p v_q \\ &= (g^{v_p v_q})^m \bmod n \end{aligned} \quad (3.4)$$

Due to the value of u being small, it is feasible to search for the calculated value in a look-up table with tuples of type $((g^{v_p v_q})^i \bmod n, i)$, for $0 < i < u$. Even without the look-up table, one can quickly determine if ct is an encryption of 0, because $D(E(0)) = 1$:

$$c^{v_p v_q} \bmod n = (g^{v_p v_q})^m \bmod n \stackrel{m \equiv 0}{=} (g^{v_p v_q})^0 \bmod n = 1 \bmod n = 1 \quad (3.5)$$

So far we have presented the decryption equations using all p, q, v_p, v_q values, but we can utilize a faster decryption method. This method only requires the computation of modulo p or q , because m is uniquely determined from $c^{v_p} \bmod p$ or $c^{v_q} \bmod q$. For the remaining of the document we will only be considering this decryption method.

Homomorphic property The DKG cryptosystem is homomorphic with respect to the addition operation (+). For $ct_1, ct_2 \in C$, there exists $m_1, m_2 \in M$ and $r_1, r_2 \in \mathbb{Z}_{v_p v_q}$, where $ct_1 = g^{m_1} h^{r_1} \bmod n$ and $ct_2 = g^{m_2} h^{r_2} \bmod n$ and Equation 2.5 holds:

$$\begin{aligned} E(m_1) * E(m_2) &= (g^{m_1} h^{r_1} \bmod n) * (g^{m_2} h^{r_2} \bmod n) \\ &= (g^{m_1} h^{r_1} * g^{m_2} h^{r_2}) \bmod n \\ &= g^{(m_1+m_2)} h^{(r_1+r_2)} \bmod n \\ &= E(m_1 + m_2) \end{aligned} \quad (3.6)$$

Additional operations Multiplication of an encrypted value by a positive constant k is supported, just like in Equation 2.6, by raising it to the power of k .

Multiplication by a negative constant c is done similarly to what is specified in Equation 2.7, but instead of considering a plaintext space $\mathbb{Z}_n$, we consider $M = \mathbb{Z}_u$. As such, this operation requires raising the encrypted value to the power of $u - c$:

$$E(m)^{(u-c)} = E(-c * m) \quad (3.7)$$

Given Equations 2.8 and 3.7, subtraction of two ciphertexts is achieved by computing:

$$E(m_1 - m_2) = E(m_1 + (-1) * m_2) = E(m_1) * E(m_2)^{(u-1)} \quad (3.8)$$

To perform an equality check, between two plaintext messages, m_1, m_2 , we compute the difference of their respective ciphertexts, ct_1, ct_2 , and check if the result is 1 ($m_1 = m_2 \Leftrightarrow m_1 - m_2 = 0$):

$$D(E(m_1) - E(m_2)) = D(E(0)) = 1 \quad (3.9)$$

3.2.4 Evaluation of SPARQL queries

Now that we have defined all the core data structures and the cryptographic tools, we will proceed to explain how we evaluate SPARQL queries. We will start by explaining the protocol for doing so with plaintext data. For simplicity, we will not take into consideration role based access control and permissions, as it will not alter in any way the protocol logic.

3.2.4.1 Storage of triple patterns

The first step when designing a solution that supports SPARQL, is to define how we will store *triple patterns*. The storage strategy must preserve semantics and ensure fast access during query evaluation.

To achieve this, when we index data, our goal is to preserve the relationships between subjects, predicates and objects (which might appear in queries' graph patterns). In the

remainder of the document we will use *node* as a common term to denominate *subjects*, *predicates* and *objects*.

We will index a triple pattern, $t = (s, p, o)$ based on all of its possible variations within a SPARQL query. For example, to answer a query where we need to search for $(?v, p, o)$ (the subject is a variable) we need an index $p, o \leftarrow [s_0, \dots, s_n]$. With this index, all subjects are referenced by the predicate and object pairs in the triplestore. Applying this strategy to other possible patterns, we can index all triple pattern variations with the following 6 indexes:

$$\begin{array}{ll}
 (p, o) \leftarrow [s_0, \dots, s_n] & \text{answers } (?v, p, o) \\
 (s, o) \leftarrow [p_0, \dots, p_n] & \text{answers } (s, ?v, o) \\
 (s, p) \leftarrow [o_0, \dots, o_n] & \text{answers } (?v, p, o) \\
 s \leftarrow [(p, o)_0, \dots, (p, o)_n] & \text{answers } (s, ?v_1, ?v_2) \\
 p \leftarrow [(s, o)_0, \dots, (s, o)_n] & \text{answers } (?v_1, p, ?v_2) \\
 o \leftarrow [(p, s)_0, \dots, (p, s)_n] & \text{answers } (?v_1, ?v_2, o)
 \end{array} \tag{3.10}$$

We do not index patterns with 3 variables, as a join of two complementary indexes (e.g. s with (p, o)) can answer this kind of queries. Nonetheless, because the encryption schemes of SEnTriS V1 and V2 do not support this kind of triple pattern, we will not support this type of queries.

For patterns with 2 variables, the index consists in the concatenation of the pattern type (e.g. "PO" for indexes referencing predicate and object pairs) with the known value (e.g. a subject). This index will reference a set of node pairs, whose values we concatenate together.

For patterns with a single variable, the index is the concatenation of the pattern type (e.g. "S" for indexes referencing subjects) with the concatenation of the other two known values (e.g. predicate and object). This index references a set of individual nodes (e.g. subjects).

In Listing A.1, we define the logic for triple pattern storage engine. We will assume that this engine executes write operations *atomically*: it either executes all operations, or none.

A search operation always produces a *bindings table*. This data structure was designed to support the operations that need to be executed between intermediate collections of *bindings*, during query evaluation. In Listing A.2, we show the underlying logic of a *bindings table*.

3.2.4.2 SPARQL Query Evaluation Protocol

Having defined how data is stored, we will now explain the protocol for evaluating SPARQL queries.

Definition 8 (SPARQL Query Evaluation Protocol). Let SPARQL-EVAL be a SPARQL query evaluation protocol. It is a tuple $\text{SPARQL-EVAL} = (\text{PlanQuery}, \text{EvaluateQuery}, \text{GenerateFinalResult})$ of 3 algorithms where:

- $\{(plan, construct, upload, delete), \perp\} \leftarrow \text{PlanQuery}(\text{query})$: Takes as input a SPARQL query. If the query is well-formed, and all required operations are supported, outputs a query execution plan and three collections of triple patterns used while constructing the final query results: the *construct*, *upload* and *delete* collections. In case of any error, it outputs $\perp$.

- $\{sparqlResult, \perp\} \leftarrow \text{EvaluateQuery}(\text{plan}, \text{triplestoreID})$: Takes as input a query execution plan and the identifier of the triplestore to be queried. In case of success, outputs the query results, $\perp$ otherwise.
- $\{sparqlFinalResult, \perp\} \leftarrow \text{GenerateFinalResult}(\text{triplestoreID}, \text{queryType}, \text{variables}, \text{bindings}, \text{construct}, \text{upload}, \text{delete})$: Takes as input a triplestore identifier, the query type, a collection of variables, the *bindings* obtained from the evaluation of a WHERE clause, and three collections of triple patterns. The *construct* template will be used for generating the final result of a CONSTRUCT query, whereas the *upload* and *delete* templates will be used to compute the triples to, respectively, upload and delete in the specified triplestore. In case of success outputs a *sparqlFinalResult*, $\perp$ otherwise.

SPARQL-EVAL is a distributed, two-party protocol: the query plan generation and final result generation are executed by a *client* component; whereas the query evaluation and data updates are executed by a *triplestore* component.

10) TODO: Add SPARQL-Eval figure

T¹⁰

Query Plan Generation We have created two abstractions to separate the logic for parsing a query from the logic for generating a query execution plan. Respectively, we denominate them as the *SPARQL query engine* and *query planner*. The query engine will parse the query into processable operations. These operations will be consumed by the query planner, which will create a query execution plan.

The parsing of a **SPARQL** query is out of the scope of this thesis. We’ve based our solution in *Apache Jena* [92]. We will utilize the solution provided by the API to parse a query and decompose it into a collection of *operations*. We will not formally define these data structures, but we will assume that their implementation is correct and sufficient to generate valid query jobs. With the collection of *operations*, we will construct a plan composed by our *query jobs*. All of this logic is scoped to the query engine, and it is specified in Listing A.3.

One can argue that the collection of *Apache Jena* operations, created after parsing and compiling a query, are equivalent to a query execution plan: executing them by their precedence order, will generate a valid query evaluation. In Chapter 4, we will explain why these are not compatible with our solution, and why we have decided to implement our custom query planner. The functioning of a simplified query planner (without any reasoning capabilities) is shown in Listing A.4.

The query planner generates a query execution plan where equi-joins are randomly ordered. The basis for this decision is to not incur in any bias, when evaluating performance, by always generating a deterministic execution order. Like the parsing of a *SPARQL* query, optimization of its execution plan is out of the scope of this thesis, hence this approach.

We generate a random ordering of the equi-joins, by randomly sampling one of the valid execution orders. **BGP**s are the only elements from a query that generate equi-joins. When we have visited all **BGP** operations, we create an auxiliary graph where compatible equi-joins are neighbors. Two joins are *compatible* if they have, at least, 1 common variable. We generate the execution order of the equi-joins, by exploring the graph and computing a path that contains each job_{JOIN} once. To generate the random ordering, we randomly sample the origin of the path. After generating the execution order we add the query jobs to the execution plan.

Query Evaluation Following the dispatcher pattern, we have created two other abstractions: the *query execution*, and the *query worker*. A query execution dispatches work, and a query worker processes pieces of work.

In the specific case of our query evaluation protocol, the query execution will process the query execution plan and dispatch the next available query job, to the worker. The worker has access to the stored data and will execute the assigned query job. Currently, [SEnTriS](#) only supports sequential dispatching of jobs to a single worker.

The dispatching of jobs continuous until there are no more jobs pending execution. After the last job is executed, the query result is generated and stays available to be retrieved as output.

Listings [A.5](#) and [A.6](#) specify, respectively, the functioning of the query execution and query worker.

Generating Final Result & Applying modifications After the evaluation of a query's WHERE clause, or the generation of the *delete* and *update* templates, there is still one final phase to execute. During this phase we generate the correct final query answer, which matches the query type, and update the triplestore data, upon the presence of an update query. Listings [A.7](#) and [A.8](#) present, respectively, the logic of the SPARQL-EVAL *client* and *triplestore* components.

3.2.4.3 SPARQL Query Evaluation Protocol with Reasoning

We will now explain how we modify the SPARQL-EVAL protocol, by applying *query rewriting* techniques, during the generation of query execution plans, to retain reasoning capabilities.

Definition 9 (SPARQL Query Evaluation Protocol with Reasoning). Let SPARQL-EVAL-II be a SPARQL query evaluation protocol that supports reasoning, via the application of query rewriting techniques.

It is a tuple $\text{SPARQL-EVAL-II} = (\text{PlanQuery}, \text{SPARQL-EVAL.EvaluateQuery}, \text{SPARQL-EVAL.GenerateFinalResult})$ of 3 algorithms where:

- $\{(plan, construct, upload, delete), \perp\} \leftarrow \text{PlanQuery}(\text{query}, \text{triplestoreID}, \text{inference}, \text{expansionDepth}, \text{transitivityDepth})$: Takes as input a [SPARQL](#) query, a triplestore identifier, a boolean which indicates if the query needs to be rewritten and two integer parameters that control the rewriting depth. The *expansionDepth* specifies the number of rounds and the *transitivityDepth* specifies the number of transitivity relationships to considered. If the query is well-formed, and all required operations are supported, outputs a query execution plan and three collections of triple patterns used while constructing the final query results: the *construct*, *upload* and *delete* collections. In case of any error, it outputs $\perp$.

SPARQL-EVAL-II is a distributed, two-party protocol: the query plan generation and final result generation are executed by a *client* component; whereas the query evaluation and data updates are executed by a *triplestore* component.

Storage of Schema triples We will refer to the subset of triples containing nodes that belong to the Tbox as *schema triples*. Storage of such triples will not require indexation by the pattern type. Furthermore, in [SEnTriS](#) we will only support a single set of schema triples per triplestore, which can either be uploaded or deleted, but not modified. In

Listings A.9 and A.10 we specify, respectively, the modifications to the triplestore storage engine and the extra operations in the SPARQL-EVAL-II triplestore component, required to support the storage and retrieval of such triples.

Query Rewriting To rewrite the queries we first fetch the schema triples, and then instantiate the triplestore’s ontology. An ontology will be an abstraction where we map class and property axioms to nodes.

Using the *Apache Jena API* [92] we create an ontological model. After the generation of the model, we load it with the schema triples. Once the triples are loaded, we iterate over all the classes and map each one to equivalent classes, subclasses, restrictions and intersections they are part of.

In the next step, we iterate over all properties and map each one to their subproperties, equivalent and inverse properties. We also map the properties to two boolean parameters which are true if, respectively, the property is symmetric and transitive.

All of this information is stored in adequate auxiliary data structures which will be accessible by the query planner. Listing A.11 provides the specification of a triplestore’s ontology.

During the creation of the query execution plan, when visiting a **BGP**, the query planner consults the ontology and searches for information related with the nodes present in the triple patterns. This information is used to expand the **BGP** with new patterns. This expansion is done using a leveled approach and is uniquely scoped to the query planner. All the modifications required to the query planner and client are depicted in Listings A.12 and A.13 respectively.

Optimization and development of a fully functional query rewriting technique is out of the scope of this thesis. The system can not evaluate property paths, functions and nested queries, which limits our available solutions. We focused on trying to support a few expansions related with: subclasses, equivalent classes, restrictions, intersections, sub-properties, equivalent and inverse properties, property transitivity and symmetry. Furthermore, we assume that the data which will be queried is consistent with the ontology provided.

We decided to implement the query rewriting with parametrized expansion and transitivity depth because it offers control over the expansion of the query, which impacts performance as it introduces additional operations.

To conclude, this approach may also induce duplicates in the results of a query, so we are required to use the **DISTINCT** modifier.

3.2.4.4 Supported Queries

The two protocols defined for query evaluation do not offer full **SPARQL** support. **SEnTriS** will only be capable of evaluating a subset of queries with a few restrictions. Queries will solely be evaluated against the default graph. Only the **DISTINCT**, **ORDERBY** and **LIMIT** modifiers are allowed. Finally, graph patterns must not contain property paths, nested subqueries queries, graph selection, nor triple patterns composed solely by variables. Following the grammar used in Definition 3, a graph pattern will be defined as:

$$P ::= tp \mid P_1 . P_2 \mid P_1 \text{ UNION } P_2 \mid P_1 \text{ OPTIONAL } P_2 \mid P_1 \text{ MINUS } P_2 \quad (3.11)$$

All query types will be supported apart from the **LOAD** and **CLEAR**, which will be implemented with dedicated operations to create, upload and delete a triplestore.

3.2.5 Encryption of triple patterns

So far we have only explained how to evaluate a [SPARQL](#) query that is executed over plaintext data, but [SEnTriS](#) operates with encrypted data. As such, the previous protocols will need to be adapted to work with encrypted data and the encryption schemes used must encrypt data in such a way that it retains its functionality. For the subset of supported queries, data must be searchable via triple patterns and allow the equality operation, essential for the execution of equi-joins.

3.2.5.1 Triple Pattern SSE scheme

Definition 10 (Triple Pattern SSE Scheme). Let SSE-TP be triple pattern [SSE](#) scheme. It is defined as a tuple $\text{SSE-TP} = (\text{CreateTriplestore}, \text{Upload}, \text{Delete}, \text{UpdateTriplestore}, \text{DeleteTriplestore}, \text{SearchTriplestore}, \text{SearchSchema}, \text{Setup}, \text{ReInitialize}, \text{Encrypt}, \text{Decrypt}, \text{Trapdoor}, \text{Search}, \text{PrepareUpload}, \text{PrepareDelete}, \text{Update}, \text{DeleteAll})$ of 17 algorithms where:

- $\{\text{SUCCESS}, \perp\} \leftarrow \text{CreateTriplestore}(\text{triplestoreID}, [t_0, t_1, \dots, t_n])$: Takes as input a triplestore identifier and a set of *ABox* triples. Outputs SUCCESS if the triplestore is successfully created and encrypted, $\perp$ otherwise.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{Upload}(\text{triplestoreID}, [t_0, t_1, \dots, t_n], \text{schema})$: Takes as input a triplestore identifier, a set of triples and a boolean parameter, which is set to true if the triples are schema triples. If triples are successfully encrypted and uploaded outputs SUCCESS, $\perp$ otherwise.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{Delete}(\text{triplestoreID}, [t_0, t_1, \dots, t_n])$: Takes as input a triplestore identifier, a set of triples and a boolean parameter, which is set to true if the triples are schema triples. If triples are successfully deleted from the triplestore outputs SUCCESS, $\perp$ otherwise.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{UpdateTriplestore}(\text{triplestoreID}, \text{triples}_{\text{delete}}, \text{triples}_{\text{upload}})$: Takes as input a triplestore identifier, a set of triples to upload and a set of triples to delete. Outputs SUCCESS if the triples to delete are successfully deleted and the triples to upload are successfully encrypted and uploaded, $\perp$ otherwise.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{DeleteTriplestore}(\text{triplestoreID}, \text{schema})$: Takes as input the triplestore identifier and a boolean parameter, which is set to true if only the schema triples should be deleted. Outputs SUCCESS if the requested data is deleted, $\perp$ otherwise.
- $\{[node_0, \dots, node_n], \perp\} \leftarrow \text{SearchTriplestore}(\text{triplestoreID}, \text{triplePatterns})$: Takes as input the triplestore identifier and a collection of triples patterns. Searches the triplestore for matches against the triple patterns. Outputs a collection of nodes, or $\perp$ if an error occurred.
- $\{[t_0, t_1, \dots, t_n], \perp\} \leftarrow \text{SearchSchema}(\text{triplestoreID})$: Takes as input the triplestore identifier. Searches the encrypted triplestore for its schema triples. Outputs a collection of triples, or $\perp$ if an error occurred.
- $\lambda \leftarrow \text{Setup}()$: Takes no input. Outputs a set of newly generated security parameters.

- $\{\text{SUCCESS}, \perp\} \leftarrow \mathbf{ReInitialize}(\lambda)$: Takes as input a set of security parameters. In case of success, sets the security parameters, updates the state of the scheme and outputs $\text{SUCCESS}, \perp$ otherwise.
- $E(\text{nodes}) \leftarrow \mathbf{Encrypt}([t_0, \dots, t_n], \text{schema})$: Takes as input a set of triples and a boolean parameter, which is set to true if the triples are schema triples. Parses the data received and outputs a map of encrypted nodes, indexed by trapdoors.
- $\text{node} \leftarrow \mathbf{Decrypt}(E(\text{node}), \text{schema})$: Takes as input an encrypted node and a boolean parameter, which is set to true if the encrypted node belongs to the schema triples. Decrypts the node and outputs the plaintext value.
- $\text{trapdoor}_{n,i} \leftarrow \mathbf{Trapdoor}(\text{keyword}_n, i)$: Takes as input a keyword, and an integer. Generates and outputs the trapdoor that corresponds to the i^{th} trapdoor of the keyword.
- $[E(\text{node}_0), \dots, E(\text{node}_n)] \leftarrow \mathbf{Search}(\text{triplestoreID}, [\text{trapdoor}_0, \dots, \text{trapdoor}_n])$: Takes as input the triplestore identifier and a collection of trapdoors. Searches the encrypted triplestore for matches against the received trapdoors. Outputs a collection of encrypted nodes, where there was a match.
- $\{\text{uploadID}, \perp\} \leftarrow \mathbf{PrepareUpload}(\text{triplestoreID}, E(\text{nodes}))$: Takes as input the triplestore identifier and a map of encrypted nodes, indexed by trapdoors. Temporarily saves the operation. Outputs the identifiers of the uploads to be applied during the execution of a SSE-TP.Update operation, $\perp$ in case of an error.
- $\{\text{deletionID}, \perp\} \leftarrow \mathbf{PrepareDelete}(\text{triplestoreID}, [\text{trapdoor}_0, \dots, \text{trapdoor}_n])$: Takes as input the triplestore identifier and a collection of trapdoors, whose entries should be deleted. Temporarily saves the operation. Outputs the identifier of the deletions to be applied during the execution of a SSE-TP.Update operation, $\perp$ in case of an error.
- $\{\text{SUCCESS}, \perp\} \leftarrow \mathbf{Update}(\text{triplestoreID}, [\text{uploadID}_0, \dots, \text{uploadID}_n], [\text{deletionID}_0, \dots, \text{deletionID}_n])$: Takes as input the triplestore identifier, a collection of upload identifiers, and a collection of deletion identifiers. Outputs SUCCESS if the uploads and deletions are successfully executed, $\perp$ otherwise.
- $\{\text{SUCCESS}, \perp\} \leftarrow \mathbf{DeleteAll}(\text{triplestoreID})$: Takes as input the triplestore identifier. Outputs SUCCESS if the triplestore data is deleted, $\perp$ otherwise.

SSE-TP is a distributed, two-party protocol: search and storage of encrypted data is scoped to the *triplestore* component; whereas triplestore creation, and data encryption and decryption is scoped to the *client* component.

3.2.5.2 Trapdoor generation

Following the [SSE](#) approach, in [SEnTriS](#) we do not reference nodes using plaintext indexes. Instead, we generate *trapdoors* by deterministically encrypting the value of the indexes that would reference the node. To enforce the property that every trapdoor must uniquely reference a single node, we deterministically encrypt the index value using an incremental *initialization vector*. This initialization vector is equal to the observed frequency of the index keyword.

Trapdoors generated with $iv = 0$ always reference the latest value of the index keyword frequency.

For each index, the respective trapdoors are encrypted using a unique key. This key is derived from a master key, for the context of the index keyword.

Definition 11 (Trapdoor Generation). Let $keyword_n$ be the keyword that references node $n \in ABox$. Let K_{MASTER} be a master key for a deterministic encryption scheme and $K_{keyword_n}$ be the derived key for the context of $keyword_n$. Finally, let i be the observed frequency for $keyword_n$, and $trapdoor_{n,i}$ be the trapdoor that is built from $keyword_n$, after $i - 1$ observations of the term:

$$trapdoor_{n,i} = E(K_{keyword_n}, keyword_n, i) \quad (3.12)$$

In the case of the trapdoors that reference schema nodes, we will also base the solution in the previous indexing scheme, but instead of using a fixed schema keyword, we will randomly generate it for each triplestore.

Definition 12 (Schema Node Trapdoor Generation). Let $keyword_{schema}$ be a randomly generated keyword, specific to a triplestore t . Let $n \in TBox$ be a schema node, K_{MASTER} be a master key for a deterministic encryption scheme and K_{schema} be the derived key for the context of $keyword_{schema}$. Finally, let i be the observed frequency for $keyword_{schema}$, and $trapdoor_{schema,i}$ be the trapdoor that is built from $keyword_{schema}$, after $i - 1$ observations of the term:

$$trapdoor_{schema,i} = E(K_{schema}, keyword_{schema}, i) \quad (3.13)$$

We also reference with trapdoors the frequencies of a triple pattern, which are necessary for computing deletions.

Definition 13 (Triple Frequencies Trapdoor Generation). Let $keyword_s$, $keyword_p$ and $keyword_o$ be the keywords that reference the nodes belonging to a triple pattern $t = (s, p, o)$. Let $keyword_t = SPO \parallel keyword_s \parallel keyword_p \parallel keyword_o$ be the keyword that references t , K_{MASTER} be a master key for a deterministic encryption scheme and $K_{keyword_t}$ be the derived key for the context of $keyword_t$. Finally, let i be the observed frequency for $keyword_t$, and $trapdoor_{t,i}$ be the trapdoor that is built from $keyword_t$, after $i - 1$ observations of the term:

$$trapdoor_{t,i} = E(K_{keyword_t}, keyword_t, i) \quad (3.14)$$

Finally, uniquely to [SEnTriS V2](#), trapdoors generated with $iv = 0$ for the context of the parsed value of a node, reference its *equality tag*. We will define equality tags in [Section 3.2.5.4](#)

3.2.5.3 Initial setup & re-initialization of scheme

During initialization, we generate all the encryption keys and security parameters of the scheme. These vary between versions the two versions. Re-initialization is required because the scheme is stateful, the state exclusive to each triplestore, and we might need to execute multiple operations against different triplestores.

SEnTriS V1 The security parameters correspond to $\lambda = (K_{MASTER}, K_{RND}, K_{DET}, iv_{DET}, keyword_{schema})$ where, K_{MASTER} , K_{RND} and K_{DET} are secret keys, iv_{DET} is a randomly generated initialization vector and $keyword_{schema}$ is a freshly generated a [UUID](#) (version 4).

SEnTriS V2 The security parameters correspond to $\lambda = (K_{MASTER}, K_{RND}, K_{DGK_{priv}}, K_{DGK_{pub}}, iv_{DET}, keyword_{schema})$ where, K_{MASTER} and K_{RND} are secret keys, $K_{DGK_{priv}}$ and $K_{DGK_{pub}}$ are, respectively, a **DGK** private and public key, iv_{DET} is a randomly generated initialization vector and $keyword_{schema}$ is a freshly generated a **UUID** (version 4).

3.2.5.4 Triple Pattern encryption and decryption

The implementation of the scheme in both versions is very similar: the generation of trapdoors does not vary; the encryption of schema nodes and frequencies belonging to the triple patterns are probabilistically encrypted. They only differ on the encryption of nodes belonging to the *ABox*. This common logic is specified in Listing B.1.

SEnTriS V1 We follow the *CryptDB* [77, 76, 75] approach and encrypt *ABox* nodes in two layered onions $E(node) = E(K_{RND}, E(K_{DET}, node, iv_{DET}))$. The outer layer is probabilistically encrypted, whereas the second is deterministically encrypted. When equality checks are required, we peel the outer layer and compare the values. The version specific functions for **SEnTriS V1** are shown in Listing B.2.

SEnTriS V2 *ABox* nodes are mapped to *equality tags*. We generate equality tags by using an injective function $eqTag : (I \cup L \cup V) \rightarrow \mathbb{Z}^+$ which maps the value of a node to a unique positive integer.

Trapdoors index encrypted equality tags and encrypted equality tags index encrypted *ABox* nodes. Respectively, nodes are probabilistically encrypted $E(node) = E(K_{RND}, node)$ and equality tags are homomorphically encrypted using the **DGK** cryptosystem [23, 22, 32, 52] $E(eq_{tag}) = E(K_{DGK_{pub}}, eqTag(node))$. Equality checks are done by computing the difference between two equality tags, and checking if the value is 1. The version specific functions for **SEnTriS V2** are shown in Listing B.3.

3.2.5.5 Storage, searches & updates

When we create a triplestore, we initialize its encryption protocol by generating the security parameters. Upon successfully setting up the encryption scheme, we parse its *ABox* triple patterns, encrypt and upload them. The schema triples are uploaded in separate upload operations, if required. The triplestore creation and encryption operations are handled by the client.

Due to the high volumes of data that can be processed during write operations, these algorithms were designed to be executed in batches. Consequently, upon receiving the encrypted data pertaining to a triplestore, the triplestore component will only store it in temporary entries. These are fully persisted when a *SSE-TP.Update* operation is executed. We will continue to assume that the storage engine operations are atomic. In Listings B.4 and B.5 we specify, respectively, the storage engine and the triplestore component logic.

The search operation is a two-step process in **SEnTriS V1**. Firstly, we parse the triple patterns and generate their keyword frequency trapdoors (where $i = 0$) and search for these values. Upon receiving the encrypted values, we decrypt the frequencies, generate trapdoors accordingly and execute the final search operation.

In **SEnTriS V2**, the results of the second search yield a collection of encrypted equality tags. We need an extra communication round if we want to retrieve the encrypted nodes, referenced by the encrypted equality tags. This extra communication round is only required for search over the *ABox*.

Listings B.6 and B.7 show the search logic in the client component for versions V1 and V2, respectively.

In **SEnTriS** V1, to execute deletions, we fetch the frequencies of the triple patterns, generate their trapdoors and send a deletion request for data referenced by those trapdoors. In the case of **SEnTriS** V2, we are required to fetch the equality tags, because data referenced by them needs to be deleted as well. The deletion requests are sent to the triplestore, where they are temporarily saved.

To upload new data, firstly, we fetch the frequencies of the triple patterns keywords. Additionally, in **SEnTriS** V2, we will need to fetch the equality tags. Upon receiving all the contextual data, we decrypt it, and set the encryption scheme to the right context. We then issue deletion requests for any duplicated data, encrypt the new triple patterns and send the upload requests to the triplestore component.

Update operations are the only way to actually persist changes in the triplestores, so they are always executed at the end of an upload or deletion request. To ensure that triplestores always stay in a consistent state after subsequent modifications, deletions are processed first, and then the updated values are stored.

Listings B.8 and B.9 show the underlying logic of write operations in the client component. For version V2, we only highlight the algorithms that differ between versions.

3.2.6 Evaluation of SPARQL queries over encrypted data

We will now explain how we adapt the SPARQL-EVAL-II to work with the SSE-TP encrypted data, and create a protocol capable of evaluating **SPARQL** queries over encrypted triplestores.

Definition 14 (Secure SPARQL Query Evaluation Protocol). Let SSE-SPARQL-EVAL be a **SPARQL** query evaluation protocol, that works with SSE-TP encrypted triplestores. It is a tuple $\text{SSE-SPARQL-EVAL} = (\text{SSE-TP.ReInitialize}, \text{SSE-TP.Encrypt}, \text{SSE-TP.Decrypt}, \text{SSE-TP.PrepareUpload}, \text{SSE-TP.PrepareDelete}, \text{SSE-TP.Update}, \text{SSE-TP.Trapdoor}, \text{SSE-TP.Search}, \text{PlanQuery}, \text{PrepareQuery}, \text{PrepareSearch}, \text{PrepareBindings}, \text{EvaluateQuery}, \text{GenerateFinalResult})$ of 8 algorithms where:

- $\{(plan_{rnd}, construct, upload, delete), \perp\} \leftarrow \text{PlanQuery}(\text{query}, \text{triplestoreID}, \text{inference}, \text{expansionDepth}, \text{transitivityDepth})$: Takes as input a **SPARQL** query, a triplestore identifier, a boolean which indicates if the query needs to be rewritten and two integer parameters that control the rewriting depth. The *expansionDepth* specifies the number of rounds and the *transitivityDepth* specifies the number of transitivity relationships to considered. If the query is well-formed, and all required operations are supported, outputs a randomized query execution plan and three collections of triple patterns used while constructing the final query results: the *construct*, *upload* and *delete* collections. In case of any error, it outputs $\perp$.
- $\{preparedPlan_{rnd}, \perp\} \leftarrow \text{PrepareQuery}(\text{triplestoreID}, plan_{rnd})$: Takes as input a triplestore identifier and a randomized query execution plan. For each distinct job_{SEC_SEARCH} in the plan a $\text{SSE-SPARQL-EVAL.PrepareSearch}$ request will be issued. Upon receiving the *searchID*, the respective job_{SEC_SEARCH} keywords will be updated accordingly. Outputs a randomized prepared query execution plan. In case of any error, outputs $\perp$.

- $\{searchID, \perp\} \leftarrow \text{PrepareSearch}(\text{triplestoreID}, \text{mask}, [\text{trapdoor}_0, \dots, \text{trapdoor}_n])$: Takes as input a triplestore identifier, a masking value and a collection of trapdoors. Upon finding the values, these will be re-encrypted, if $\text{mask} \neq \perp$. Finally, issues a SSE-SPARQL-EVAL.*PrepareBindings* request and outputs the result of the algorithm.
- $\{searchID, \perp\} \leftarrow \text{PrepareBindings}([E(\text{binding}_0), \dots, E(\text{binding}_n)])$: Takes as input a collection of encrypted bindings. Generates a fresh **UUID** (version 4) and persists the received values in memory, associating them with the generated identifier. Outputs the identifier, or $\perp$ if an error occurs.
- $\{E(\text{sparqlResult}), \perp\} \leftarrow \text{EvaluateQuery}(\text{preparedPlan}_{\text{rnd}}, \mathbf{K})$: Takes as input a randomized prepared query execution plan and a secret key. Executes the query plan and outputs the encrypted **SPARQL** results. If there is an error, outputs $\perp$.
- $\{\text{sparqlFinalResult}, \perp\} \leftarrow \text{GenerateFinalResult}(\text{triplestoreID}, \text{queryType}, \text{variables}, \text{bindings}, \text{construct}, \text{upload}, \text{delete})$: Takes as input a triplestore identifier, the query type, a variables map, the *bindings* obtained from the evaluation of a WHERE clause, and three collections of triple patterns. The *construct* template will be used for generating the final result of a CONSTRUCT query, whereas the *upload* and *delete* templates will be used to compute the triples to, respectively, upload and delete in the specified triplestore. In case of success outputs a *sparqlFinalResult*, $\perp$ otherwise.

SSE-SPARQL-EVAL is a distributed, three-party protocol: search and storage of encrypted data is scoped to the *triplestore* component; encryption and decryption of data, query plan generation, preparation and final result generation is scoped to the *client* component; the evaluation of the query execution plans are scoped to a secure computation *proxy* component.

3.2.6.1 Query Preparation Phase

In this phase, depicted in Figure ... **T**¹¹, we generate the query execution plan. We adapt the SPARQL-EVAL-II protocol in two ways. Firstly, we obfuscate the variables. Secondly, to introduce extra randomness in the query plan structure, the ontology and the expansion algorithm are randomized. Listing C.1 and C.2 specify the modifications to the **SPARQL** planner and ontology logic.

Upon the generation of a randomized query execution plan, the client component processes the secure search jobs and requests the necessary SSE-SPARQL-EVAL.*PrepareSearch* operations. In **SEnTriS** V2, a masking value is also computed and sent alongside each of the previous requests.

After receiving the collections of trapdoors, the triplestore component matches them against the triplestore's data. The bindings found are sent to the computation proxy, to be temporarily persisted and used during the query evaluation phase. In **SEnTriS** V2, before sending the data to the proxy, the bindings are re-encrypted, using the received *mask* value.

After receiving the search identifiers, the client substitutes the search jobs keywords with the respective values. After all search jobs have been successfully updated, the plan is said to be prepared and ready to be sent to the proxy component. Listings C.3 and C.4 show the logic of the client for **SEnTriS** V1 and V2 respectively.

11) *TODO*: Create figure for query preparation phase

Listings C.5 and C.6 show the logic of the triplestore component and how search values are blinded in the storage engine.

Finally, listings C.7 and C.8 provide the logic for the proxy component and how the bindings are temporarily stored. We have only highlighted the algorithms that differ from the previous protocols.

3.2.6.2 Query Evaluation Phase

Query evaluation, illustrated in Figure ... T^{12} , is done in the computation proxy and retains most of the logic present in the previous SPARQL query evaluation protocols. The only difference is in the query worker. This worker will have access to a secret key and will operate with encrypted bindings tables. Listing C.9 shows the query evaluation logic in the proxy component. Listing C.10 specifies the modifications to the query worker.

12) TODO: Create figure for query evaluation phase

In SEnTriS V1 the secret key corresponds to K_{RND} , and it will be used to peel the outer layer of the encrypted nodes. In SEnTriS V2, this key corresponds to an equality key $K_{DGK_{eq}} = (p, vp, n, u)$, which is sufficient to compute the difference between two encrypted equality tags, but not enough to perform a decryption operation. To execute a decryption, in the DGK cryptosystem, we require the g parameter, and it is not derivable from the values that compose the $K_{DGK_{eq}}$. The only exception happens with the encryption of 0, which is always 1.

As seen in Listing C.11, in SEnTriS V1, before constructing the bindings tables, the probabilistic encryption layer of the bindings will be decrypted. The deterministic values allow the bindings tables to operate as normal.

In SEnTriS V2, the values are always probabilistic. Consequently, as shown in Listing C.12, equality checks, and set operations that require this operation, will have to be keyed and use the provided equality key.

Before constructing the encryption tables we will need to group the bindings, keeping track of pattern indexes but retaining one unique value, per distinct group of equal equality tags. This will enable us to construct the encryption tables without altering their data structures. The non-deterministic nature of the data will result in an increase of complexity for internal searches, as they will have to be done using sequential scans. In Chapter 4, we will explain how we mitigate this impact using parallelism.

Listings C.13 and C.14 provide the logic for the construction of bindings tables, in SEnTriS V2, and their respective modifications.

3.2.6.3 Decryption of SPARQL query results

Upon receiving the results from the query evaluation, the client component will decrypt them and construct the final results. This final step will require an extra communication round for SEnTriS V2, because the returned values are encrypted equality tags. These will need to be unmasked, and used to search for the encrypted nodes.

After the decryption of the results, any modifications required are executed. Listings C.15 and C.16 illustrate this logic for SEnTriS V1 and V2, respectively.

3.2.7 The M/M setting

So far we have specified the building blocks for the secure evaluation of SPARQL queries, but SEnTriS is not uniquely a query evaluation protocol. It is also a triplestore management and sharing protocol: users can create, manage and share encrypted and non-encrypted triplestores with varying degrees of access.

To adapt the previous protocols to the *M/M* setting we will require an identity and access management protocol, with secret sharing between authorized users.

Authentication will be validated with *sessions*, management of access to triplestores will be achieved with the definition of triplestore *Access Control Lists* and strict enforcement of permissions. Furthermore, we will rely on a lock-based access strategy, and *ephemeral access tokens*, to ensure that concurrent access is done in a controlled and safe manner.

Finally, roles will grant different execution rights to users, which restrict the type of operations that they can execute within the system. Users can have the *ADMIN*, *PRIVILEGED* and *BASIC* roles. For ease of communication from now onwards, we will refer to a user with each role, respectively, as *admin*, *privileged* or *basic* user.

Basic users will only be able to query existing triplestores, if they have been given access to do so. They can request *read* or *write* access to existing triplestores.

Privileged users will have all the permissions of a *basic* user. Additionally, they will be able to create new triplestores, grant or revoke read/write access to their triplestores, transfer ownership to other users and process access requests targeting triplestores they own.

Admin users will possess all the permissions of a *privileged* user, but they will also be able to manage any triplestore and process role change requests.

Definition 15 (Identity and Access Management Protocol with Secret Sharing). Let IAM be an identity and access management protocol between authenticated users, where temporary access to triplestores is granted using ephemeral access tokens. It is a tuple of $IAM = (\text{Init}, \text{CreateSession}, \text{SaveUser}, \text{DeleteUser}, \text{IssueGrantRoleRequest}, \text{GetPendingRoleRequests}, \text{ProcessRoleRequest}, \text{CreateTriplestoreAccessList}, \text{GetTriplestores}, \text{GetUsersWithAccess}, \text{ChangeTriplestoreOwner}, \text{DeleteTriplestoreAccessList}, \text{GrantTriplestoreAccess}, \text{RevokeTriplestoreAccess}, \text{IssueAccessRequest}, \text{GetPendingAccessRequests}, \text{ProcessAccessRequest}, \text{CreateAccessToken}, \text{DeleteAccessToken}, \text{CheckReadAccess}, \text{CheckWriteAccess}, \text{CheckOwnerAccess}, \text{CheckIfActive}, \text{AcquireTriplestoreLock}, \text{ReleaseTriplestoreLock}, \text{Auth}, \text{RegisterUser}, \text{UnregisterUser}, \text{IssueUpgradeRequest}, \text{IssueDowngradeRequest}, \text{ListPendingRoleRequests}, \text{IssueRoleDecision}, \text{ListTriplestores}, \text{UpdateTriplestoreOwner}, \text{GrantAccess}, \text{RevokeAccess}, \text{RequestAccess}, \text{ListPendingAccessRequests}, \text{IssueAccessDecision}, \text{ListUsersWithAccess}, \text{CreateTriplestore}, \text{DeleteTriplestore}, \text{SaveSecrets}, \text{DeleteSecrets}, \text{GetSecrets}, \text{SSE-TP.Setup})$ of 46 algorithms where:

- $\{\text{SUCCESS}, \perp\} \leftarrow \text{Init}()$: initializes the protocol by registering a default admin user. Outputs SUCCESS if the user is created and registered, $\perp$ otherwise.
- $\{\text{cookie}, \perp\} \leftarrow \text{CreateSession}(\text{username}, \text{password})$: takes as input a pair of user credentials. In case of valid credentials, creates a new session and outputs a *cookie* holding its data, $\perp$ otherwise.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{SaveUser}(\text{username}, \text{password})$: takes as input a pair of user credentials. Outputs SUCCESS if a new user is registered, $\perp$ otherwise.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{DeleteUser}(\text{cookie}, \text{username})$: takes as input a cookie and the username of the user to delete. Outputs SUCCESS if the user is successfully deleted, $\perp$ otherwise.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{IssueGrantRoleRequest}(\text{cookie}, \text{username}, \text{issuer}, \text{role})$: takes as input a cookie, the usernames of the target and issuer of the role changing request

and the desired role. Outputs SUCCESS if the request is successfully persisted, $\perp$ otherwise.

- $\{[roleRequest_0, \dots, roleRequest_n], \perp\} \leftarrow \text{GetPendingRoleRequests}(\text{cookie}, \text{username})$: takes as input a cookie, and the request issuer username. Outputs the collection of pending role changing requests, $\perp$ in case of error.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{ProcessRoleRequest}(\text{cookie}, \text{username}, \text{target}, \text{requestID}, \text{decision})$: takes as input a cookie, the username of the processing request's issuer, the username of the role change target, the request identifier and the issuer's decision. Outputs SUCCESS if the role changing request is successfully processed, according to the provided decision, $\perp$ otherwise.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{CreateTriplestoreAccessList}(\text{cookie}, \text{issuer}, \text{triplestoreID})$: takes as input a cookie, the username of the request's issuer and a triplestore identifier. Outputs SUCCESS if the triplestore access policy is successfully generated, $\perp$ otherwise.
- $[triplestoreID_0, \dots, triplestoreID_n] \leftarrow \text{GetTriplestores}(\text{cookie}, \text{issuer}, \text{owns}, \text{read}, \text{write})$: takes as input a cookie, the username of the request's issuer and three boolean filter parameters. Outputs a collection of triplestore identifiers: if *owns* is *true*, returns identifiers of *owned* triplestores; if *read* is *true*, returns identifiers of triplestores that the user has *read* access; if *write* is *true*, returns identifiers of triplestores that the user has *write* access;
- $\{msg, \perp\} \leftarrow \text{GetUsersWithAccess}(\text{cookie}, \text{triplestoreID}, \text{write}, \text{at})$: takes as input a cookie, a triplestore identifier, a boolean filter parameter and an access token. If the access token has sufficient permissions, outputs a message with the specified triplestore's *owner* username and the collection of *read* and *write* access (if the *write* parameter is *true*) users' usernames, $\perp$ otherwise.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{ChangeTriplestoreOwner}(\text{cookie}, \text{triplestoreID}, \text{target}, \text{at})$: takes as input a cookie, a triplestore identifier, a username and an access token. Outputs SUCCESS if the ownership of the specified triplestore is successfully transferred to the *target* user, $\perp$ otherwise.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{DeleteTriplestoreAccessList}(\text{cookie}, \text{triplestoreID}, \text{at})$: takes as input a cookie a triplestore identifier and an access token. Outputs SUCCESS if the specified triplestore access policy is successfully deleted, $\perp$ otherwise.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{GrantTriplestoreAccess}(\text{cookie}, \text{triplestoreID}, \text{target}, \text{write}, \text{at})$: takes as input a cookie, a triplestore identifier, a username, a boolean parameter and an access token. Outputs SUCCESS if the *target* user is successfully granted *read*, or *write* (if *write* parameter is *true*), access against the specified triplestore, $\perp$ otherwise.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{RevokeTriplestoreAccess}(\text{cookie}, \text{triplestoreID}, \text{target}, \text{write}, \text{at})$: takes as input a cookie, a triplestore identifier, a username, a boolean parameter and an access token. Outputs SUCCESS if the *target* user *read*, or *write* (if *write* parameter is *true*), access against the specified triplestore is successfully revoked, $\perp$ otherwise.

- $\{\text{SUCCESS}, \perp\} \leftarrow \text{IssueAccessRequest}(\text{cookie}, \text{triplestoreID}, \text{target}, \text{write})$: takes as input a cookie, a triplestore identifier, a username, a boolean parameter and an access token. Outputs SUCCESS if an access change request, for the *target* user, granting *read* or *write* (if the *write* parameter is *true*) permissions against the specified triplestore is successfully issued, $\perp$ otherwise.
- $\{\text{accessRequest}_0, \dots, \text{accessRequest}_n, \perp\} \leftarrow \text{GetPendingAccessRequests}(\text{cookie}, \text{triplestoreID}, \text{at})$: takes as input a cookie, a triplestore identifier and an access token. Outputs the collection of the triplestore's pending access change requests, if the specified access token has sufficient permissions, $\perp$ otherwise.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{ProcessAccessRequest}(\text{cookie}, \text{triplestoreID}, \text{requestID}, \text{decision}, \text{at})$: takes as input a cookie, a triplestore identifier, an access change request identifier, a boolean parameter and an access token. Outputs SUCCESS if the specified request is successfully processed according to the provided *decision* parameter, $\perp$ otherwise.
- $\{\text{tokenID}, \perp\} \leftarrow \text{CreateAccessToken}(\text{cookie}, \text{triplestoreID}, \text{target})$: takes as input a cookie, a triplestore identifier, and a username. Outputs the newly generated token identifier, if the *target* user has sufficient permissions against the specified triplestore, $\perp$ otherwise.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{DeleteAccessToken}(\text{cookie}, \text{triplestoreID}, \text{at})$: takes as input a cookie, a triplestore identifier, and an access token. Outputs SUCCESS if the token is successfully deleted, $\perp$ otherwise.
- $\{\text{OK}, \perp\} \leftarrow \text{CheckReadAccess}(\text{triplestoreID}, \text{at})$: takes as input a triplestore identifier and an access token. Outputs OK if the token grants *read* access to the specified triplestore, $\perp$ otherwise.
- $\{\text{OK}, \perp\} \leftarrow \text{CheckWriteAccess}(\text{triplestoreID}, \text{at})$: takes as input a triplestore identifier and an access token. Outputs OK if the token grants *write* access to the specified triplestore, $\perp$ otherwise.
- $\{\text{OK}, \perp\} \leftarrow \text{CheckOwnerAccess}(\text{triplestoreID}, \text{at})$: takes as input a triplestore identifier and an access token. Outputs OK if the token grants *owner* access to the specified triplestore, $\perp$ otherwise.
- $\{\text{OK}, \perp\} \leftarrow \text{CheckIfActive}(\text{at})$: takes as input a triplestore identifier and an access token. Outputs OK if the token is active (the associated session or the token has not expired), $\perp$ otherwise.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{AcquireTriplestoreLock}(\text{cookie}, \text{at})$: takes as input a cookie and an access token. Outputs SUCCESS if the token is successfully associated with a newly generated lock against the triplestore affected by the token, $\perp$ otherwise.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{ReleaseTriplestoreLock}(\text{cookie}, \text{at})$: takes as input a cookie and an access token. Outputs SUCCESS if the lock associated with the token is successfully unlocked, $\perp$ otherwise.
- $\{\text{cookie}, \perp\} \leftarrow \text{Auth}(\text{username}, \text{password})$: takes as input a pair of user credentials, issues a *IAM.auth* request and outputs the result of the algorithm.

- $\{\text{SUCCESS}, \perp\} \leftarrow \text{RegisterUser}(\text{username}, \text{password})$: takes as input a pair of user credentials, issues a *IAM.registerUser* request and outputs the result of the algorithm.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{UnregisterUser}(\text{cookie}, \text{username})$: takes as input a cookie and the username of the user to delete, issues a *IAM.deleteUser* request and outputs the result of the algorithm.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{IssueUpgradeRequest}(\text{cookie}, \text{username})$: takes as input a cookie, and a username. Issues a *IAM.issueGrantRoleRequest* request, to grant the PRIVILEGED role, and outputs the result of the algorithm.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{IssueDowngradeRequest}(\text{cookie}, \text{username})$: takes as input a cookie, and a username. Issues a *IAM.issueGrantRoleRequest* request, to grant the BASIC role, and outputs the result of the algorithm.
- $\{[\text{roleRequest}_0, \dots, \text{roleRequest}_n], \perp\} \leftarrow \text{ListPendingRoleRequests}(\text{cookie}, \text{username})$: takes as input a cookie, and the request issuer username. Issues a *IAM.getPendingRoleRequests* request and outputs the result of the algorithm.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{IssueRoleDecision}(\text{cookie}, \text{username}, \text{target}, \text{requestID}, \text{decision})$: takes as input a cookie, the username of the processing request's issuer, the username of the role change target, the request identifier and the issuer's decision. Issues a *IAM.processRoleRequest* request and outputs the result of the algorithm.
- $\{[\text{triplestoreID}_0, \dots, \text{triplestoreID}_n], \perp\} \leftarrow \text{ListTriplestores}(\text{cookie}, \text{issuer}, \text{write}, \text{read}, \text{owns})$: takes as input a cookie, the username of the request's issuer and three boolean filter parameters. Issues a *IAM.getTriplestores* request and outputs the result of the algorithm, or $\perp$ in case of error.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{UpdateTriplestoreOwner}(\text{cookie}, \text{triplestoreID}, \text{issuer}, \text{target})$: takes as input a cookie, a triplestore identifier, and two usernames, respectively, of the request issuer and the new triplestore owner. Tries to create an access token and lock the triplestore. Issues a *IAM.changeTriplestoreOwner* request and outputs the retrieved response.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{GrantAccess}(\text{cookie}, \text{triplestoreID}, \text{issuer}, \text{target}, \text{write})$: takes as input a cookie, a triplestore identifier, the usernames of the request issuer and the user to grant access, and a boolean parameter to specify the access type. Tries to create an access token and lock the triplestore. Issues a *IAM.grantTriplestoreAccess* request and outputs the retrieved response.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{RevokeAccess}(\text{cookie}, \text{triplestoreID}, \text{issuer}, \text{target}, \text{write})$: takes as input a cookie, a triplestore identifier, the usernames of the request issuer and the user to revoke access, and a boolean parameter to specify the access type. Tries to create an access token and lock the triplestore. Issues a *IAM.revokeTriplestoreAccess* request and outputs the retrieved response.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{RequestAccess}(\text{cookie}, \text{triplestoreID}, \text{issuer}, \text{write})$: takes as input a cookie, a triplestore identifier, a username and a boolean parameter to specify the access type. Issues a *IAM.issueAccessRequest* request and outputs the retrieved response.

- $\{\text{SUCCESS}, \perp\} \leftarrow \text{ListPendingAccessRequests}(\text{cookie}, \text{triplestoreID}, \text{issuer})$: takes as input a cookie, a triplestore identifier and a request issuer username. Issues a `IAM.getPendingAccessRequests` request and outputs the retrieved response.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{IssueAccessDecision}(\text{cookie}, \text{triplestoreID}, \text{issuer}, \text{requestID}, \text{decision})$: takes as input a cookie, a triplestore identifier, the issuer username, an access request identifier and the final decision. Tries to create an access token and lock the triplestore. Issues a `IAM.processAccessRequest` request and outputs the retrieved response.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{ListUsersWithAccess}(\text{cookie}, \text{triplestoreID}, \text{issuer}, \text{write})$: takes as input a cookie, a triplestore identifier, the request issuer username and a boolean parameter specifying the type of access. Tries to create an access token. Issues a `IAM.getUsersWithAccess` request and outputs the retrieved response.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{CreateTriplestore}(\text{cookie}, \text{triplestoreID})$: takes as input a cookie and a new triplestore identifier. Issues a `IAM.createTriplestoreAccessList` request and outputs the retrieved response. In case of SUCCESS, will also try to create an access token and lock the newly created triplestore. Executes `SSE-TP.Setup` and tries to save the generated secrets by issuing a `IAM.saveSecrets` request. In case of error, deletes the newly generated triplestore and outputs $\perp$. Outputs SUCCESS if the triplestore `ACL` is successfully created and the secrets saved correctly.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{DeleteTriplestore}(\text{cookie}, \text{issuer}, \text{triplestoreID})$: takes as input a cookie, a username and a triplestore identifier. Tries to create an access token and lock the triplestore. It issues a `IAM.DeleteSecrets` request and a `IAM.DeleteTriplestoreAccessList` request. Outputs SUCCESS if the triplestore is deleted, $\perp$ otherwise.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{SaveSecrets}(\text{triplestoreID}, \text{secrets}, \text{at})$: takes as input a triplestore identifier, a collection of secrets and an access token. If the token as sufficient permissions, the secrets are saved and outputs SUCCESS, $\perp$ otherwise.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{DeleteSecrets}(\text{triplestoreID}, \text{at})$: takes as input a triplestore identifier and an access token. If the token as sufficient permissions, the triplestore secrets are deleted and outputs SUCCESS, $\perp$ otherwise.
- $\{\text{secrets}, \perp\} \leftarrow \text{GetSecrets}(\text{triplestoreID}, \text{at})$: takes as input a triplestore identifier and an access token. If the token as sufficient permissions, outputs the triplestore secrets, $\perp$ otherwise.

IAM is a distributed, multi-party protocol between a *client*, a *IAM provider* and a *vault* component. The client is the public facing component, that users interact with. The IAM provider validates sessions, manages user data and triplestore access policies, stores role and permission change requests, validates permission rules, emits session cookies and ephemeral access tokens, validates access requests and controls the locking and unlocking of resources. The vault component is responsible for storage of secrets.

3.2.7.1 Authentication Procedure & Session Validation

Authentication is the process where we verify the identity of a user. We do so, by checking that the provided credentials, during login, match the ones saved within our system. To avoid having to verify this condition in every request sent by the user, we use *sessions*.

The session is generated upon a successful login request. A valid session will prove the identity of the user sending the requests. In [SEnTriS](#), an authenticated user needs to be valid (has been registered within the system) and hold an active session.

The login request is initiated by the client component, which will redirect the provided credentials to the IAM provider, responsible for the emission of a cookie, holding the session data. In the IAM provider we store the user credentials, following [National Institute of Standards and Technology \(NIST\)](#) best practices [64]. The passwords are stored using a suitable one-way key derivation function (e.g. [Password-based Key Derivation Function 2 \(PBKDF2\)](#)). Listings [D.1](#) and [D.2](#) provide the authentication logic for the client and IAM provider components respectively. Listing [D.3](#) shows the IAM provider storage engine logic for storing user data. We will assume that the storage engine is atomic.

3.2.7.2 Lock-based Concurrent Access & Permission Enforcement

Concurrent operations can target users or triplestores. The correctness of concurrent write and read operations is achieved using a *lock-based* approach. Resources (users and triplestores) can be in a *locked* or *unlocked* state. Write operations can only be requested against *unlocked* resources, which remain locked while the operation is being executed. Consequently, read operations can occur simultaneously, but write operations are sequential, against the same resource.

Locks are enough for requests targeting users, because the execution context is always the same. Operations can only be executed by users, against themselves, or by *admin* users, and the origin of the requests is always scoped to a *Client* component.

Operations against triplestores are more complex. The query evaluation pipeline is distributed between different components. As such, whenever the execution context changes we must verify if the author of the request has sufficient permissions to execute the operations being requested.

To fulfill such requirements, every request will have an authorization header, following the bearer authentication scheme, containing an *access token*. Permission checks are executed, by issuing a verification request, with the token, to a verifier component. In the case of [SEnTriS](#) it will be the [IAM](#) provider. The [IAM](#) provider extracts the token from the authorization header, verifies the validity of the token, and if its associated permissions are correct.

Queries that contain write predicates, will require that their target remains *locked*, during their evaluation. As such, an *access token* can be associated with a *lock*. The targeted triplestore will remain in a *locked* state until its *access token* is deleted or the *lock* expires.

The emission of locks and ephemeral access tokens is also scoped to the IAM provider. The creation and management of locks is scoped to a *locks client*. This abstraction functions like an atomic storage engine as shown in Listing [D.4](#). Listing [D.5](#) provides the logic for the emission of access tokens, triplestore locking and permission verification. Finally, Listing [D.6](#) shows the underlying storage engine for storing and modifying [ACLs](#) and access tokens.

3.2.7.3 Creation and Sharing of Triplestores

Any privileged user can create new triplestores. Access to triplestores can either be given, requested or revoked. An access request will have to be processed manually by an admin user or the triplestore owner. Listings [D.7](#), [D.8](#) and [D.9](#) show the logic for triplestore access management.

3.2.7.4 Sharing of Secrets

Access to a triplestore will grant permission to consult the secrets of the triplestore. Permission to modify a triplestore's secrets will only be granted to its owner. Revoking access to a triplestore, will also revoke access to its secrets. Listing D.10 shows the vault component logic.

3.2.7.5 Management of Roles

To manage their roles, users can issue role upgrade and downgrade requests via the client. This will trigger the issuing of a role change request, which can either be processed manually, by an admin user, or automatically in case of a downgrade request. Listings D.11, D.12, and D.13 specify this logic.

3.2.8 Definition of SEnTriS

Now that we have explained how we manage access in the M/M setting, we are ready to fully define SEnTriS.

Definition 16 (SEnTriS). Let SEnTriS be a protocol that supports the collaborative development and sharing of SSE-TP encrypted triplestores and is capable of answering SPARQL queries using the SSE-SPARQL-EVAL protocol, adapted to the M/M setting. It is a tuple $\text{SEnTriS} = (\text{IAM.Init}, \text{IAM.CreateSession}, \text{IAM.SaveUser}, \text{IAM.DeleteUser}, \text{IAM.IssueGrantRoleRequest}, \text{IAM.GetPendingRoleRequests}, \text{IAM.ProcessRoleRequest}, \text{IAM.CreateTriplestoreAccessList}, \text{IAM.GetTriplestores}, \text{IAM.GetUsersWithAccess}, \text{IAM.ChangeTriplestoreOwner}, \text{IAM.DeleteTriplestoreAccessList}, \text{IAM.GrantTriplestoreAccess}, \text{IAM.RevokeTriplestoreAccess}, \text{IAM.IssueAccessRequest}, \text{IAM.GetPendingAccessRequests}, \text{IAM.ProcessAccessRequest}, \text{IAM.CreateAccessToken}, \text{IAM.DeleteAccessToken}, \text{IAM.CheckReadAccess}, \text{IAM.CheckWriteAccess}, \text{IAM.CheckOwnerAccess}, \text{IAM.CheckIfActive}, \text{IAM.AcquireTriplestoreLock}, \text{IAM.ReleaseTriplestoreLock}, \text{IAM.Auth}, \text{IAM.RegisterUser}, \text{IAM.UnregisterUser}, \text{IAM.IssueUpgradeRequest}, \text{IAM.issueDowngradeRequest}, \text{IAM.ListPendingRoleRequests}, \text{IAM.IssueRoleDecision}, \text{IAM.ListTriplestores}, \text{IAM.UpdateTriplestoreOwner}, \text{IAM.GrantAccess}, \text{IAM.RevokeAccess}, \text{IAM.RequestAccess}, \text{IAM.ListPendingAccessRequests}, \text{IAM.IssueAccessDecision}, \text{IAM.ListUsersWithAccess}, \text{IAM.SaveSecrets}, \text{IAM.DeleteSecrets}, \text{IAM.GetSecrets}, \text{SSE-TP.Setup}, \text{SSE-TP.ReInitialize}, \text{SSE-TP.Encrypt}, \text{SSE-TP.Decrypt}, \text{SSE-TP.Trapdoor}, \text{SSE-TP.Upload}, \text{IAM.CreateTriplestore}, \text{UploadTriples}, \text{DeleteTriplestore}, \text{SearchSchema}, \text{ExecuteSPARQLQuery}, \text{PrepareSearch}, \text{Search}, \text{PrepareUpload}, \text{PrepareDelete}, \text{Update}, \text{DeleteAll}, \text{PrepareBindings}, \text{EvaluateQuery})$ of 62 algorithms where:

- $\{\text{SUCCESS}, \perp\} \leftarrow \text{UploadTriples}(\text{cookie}, \text{triplestoreID}, [t_0, t_1, \dots, t_n], \text{schema})$: Takes as input a cookie, a triplestore identifier, a set of triples and a boolean parameter, which is set to true if the triples are schema triples. If *schema* is false generates a SPARQL Update query with the provided triples, and executes SEnTriS.ExecuteSPARQLQuery. If *schema* is true, tries to create an access token and lock the triplestore, by issuing a IAM.CreateAccessToken followed by a IAM.AcquireTriplestoreLock request. Upon successfully generating the token, and locking the triplestore, issues a IAM.GetSecrets request and executes SSE-TP.ReInitialize. With the encryption scheme fully initialized, executes a modified implementation of SSE-TP.Upload where the issued requests follow the definitions of SEnTriS.PrepareDelete, SEnTriS-

PrepareUpload and *SEnTriS.Update*. In case of any error, issues a *IAM.DeleteAccessToken* and outputs $\perp$.

- $\{\text{SUCCESS}, \perp\} \leftarrow \text{DeleteTriplestore}(\text{cookie}, \text{triplestoreID})$: Takes as input the triplestore identifier. Tries to create an access token and lock the triplestore, by issuing a *IAM.CreateAccessToken* followed by a *IAM.AcquireTriplestoreLock* request. Upon successful generation of the access token, and locking of the triplestore, issues a *SEnTriS.DeleteAll*, followed by the issuing of a *IAM.DeleteSecrets* request and a *IAM.DeleteTriplestoreAccessList* request. Outputs *SUCCESS* if the triplestore is deleted, $\perp$ otherwise. In case of any error, issues a *IAM.DeleteAccessToken* and outputs $\perp$.
- $\{[t_0, t_1, \dots, t_n], \perp\} \leftarrow \text{SearchSchema}(\text{cookie}, \text{triplestoreID})$: Takes as input a cookie, a triplestore identifier. Tries to create an access token, by issuing a *IAM.CreateAccessToken* request. Upon successful generation of the token, executes a modified implementation of *SSE-TP.SearchSchema*, where the search requests issued follow the *SEnTriS.Search* definition. Outputs a collection of triples. In case of any error, issues a *IAM.DeleteAccessToken* and outputs $\perp$.
- $\{\text{sparqlFinalResult}, \perp\} \leftarrow \text{ExecuteSPARQLQuery}(\text{cookie}, \text{query}, \text{triplestoreID}, \text{inference}, \text{expansionDepth}, \text{transitivityDepth})$: Takes as input a cookie, a *SPARQL* query, a triplestore identifier, a boolean which indicates if the query needs to be rewritten and two integer parameters that control the rewriting depth. The *expansionDepth* specifies the number of rounds and the *transitivityDepth* specifies the number of transitivity relationships to considered. Firstly, the client will try to generate an access token, by issuing a *IAM.CreateAccessToken* request, if it is a *SPARQL* update query it will also issue a *IAM.AcquireTriplestoreLock* request to lock the triplestore. Upon the successful generation of the access token, it will issue a *IAM.GetSecrets* request to fetch the secrets and execute *SSE-TP.ReInitialize*. If no error occurs, it starts executing the *SSE-SPARQL-EVAL.GeneratePlan*, followed by a modified implementation of *SSE-SPARQL-EVAL.PreparePlan*, where requests issued follow the definition of *SEnTriS.PrepareSearch*, and carry the access token. Issues a *SEnTriS.EvaluateQuery* request and, upon receiving the encrypted bindings, executes *SSE-SPARQL-EVAL.GenerateFinalResult*. In case of success, outputs a *sparqlFinalResult*. If any error occurs, issues a *IAM.DeleteAccessToken* and outputs $\perp$.
- $\{\text{searchID}, \perp\} \leftarrow \text{PrepareSearch}(\text{triplestoreID}, \text{mask}, [\text{trapdoor}_0, \dots, \text{trapdoor}_n], \text{at})$: Takes as input a triplestore identifier, a masking value, a collection of trapdoors and an access token. If the provided access token has sufficient permissions, executes a modified implementation of *SSE-SPARQL-EVAL.PrepareSearch* where the requests issued follow the definition of *SEnTriS.PrepareBindings*. Outputs the received result.
- $\{[E(\text{node}_0), \dots, E(\text{node}_n)], \perp\} \leftarrow \text{Search}(\text{triplestoreID}, [\text{trapdoor}_0, \dots, \text{trapdoor}_n], \text{at})$: Takes as input the triplestore identifier, a collection of trapdoors and an access token. Issues a *IAM.CheckReadAccess* request to validate if the access token has sufficient permission against the specified triplestore. In case of a positive result, executes *SSE-TP.Search* and outputs the result.
- $\{\text{uploadID}, \perp\} \leftarrow \text{PrepareUpload}(\text{triplestoreID}, E(\text{nodes}), \text{at})$: Takes as input the triplestore identifier, a map of encrypted nodes, indexed by trapdoors, and an access

token. Issues a *IAM.CheckWriteAccess* request to validate if the access token has sufficient permission against the specified triplestore. In case of a positive result, starts execution and executes *SSE-TP.PrepareUpload*. Outputs the identifiers of the uploads to be applied during the execution of a *SEnTriS.Update* operation, $\perp$ in case of an error.

- $\{deletionID, \perp\} \leftarrow \text{PrepareDelete}(\text{triplestoreID}, [\text{trapdoor}_0, \dots, \text{trapdoor}_n], \text{at})$: Takes as input the triplestore identifier, a collection of trapdoors, whose entries should be deleted, and an access token. Issues a *IAM.CheckWriteAccess* request to validate if the access token has sufficient permission against the specified triplestore. In case of a positive result, executes *SSE-TP.PrepareDelete* and outputs the identifier of the deletions to be applied during the execution of a *SEnTriS.Update* operation, $\perp$ in case of an error.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{Update}(\text{triplestoreID}, [\text{uploadID}_0, \dots, \text{uploadID}_n], [\text{deletionID}_0, \dots, \text{deletionID}_n], \text{at})$: Takes as input the triplestore identifier, a collection of upload identifiers, a collection of deletion identifiers and an access token. Issues a *IAM.CheckWriteAccess* request to validate if the access token has sufficient permission against the specified triplestore. In case of a positive result, executes *SSE-TP.Update* and outputs the result.
- $\{\text{SUCCESS}, \perp\} \leftarrow \text{DeleteAll}(\text{triplestoreID}, \text{at})$: Takes as input the triplestore identifier and an access token. Issues a *IAM.CheckOwnerAccess* request to validate if the access token has sufficient permission against the specified triplestore. In case of a positive result, executes *SSE-TP.DeleteAll* and outputs the result.
- $\{searchID, \perp\} \leftarrow \text{PrepareBindings}([E(\text{binding}_0), \dots, E(\text{binding}_n)], \text{at})$: Takes as input a collection of encrypted bindings and an access token. Issues a *IAM.CheckIfActive* request to validate if the access token is still active. In case of a positive result, executes *SSE-SPARQL-EVAL.PrepareBindings* and outputs the results.
- $\{E(\text{sparqlResult}), \perp\} \leftarrow \text{EvaluateQuery}(\text{preparedPlan}_{\text{rnd}}, \text{K}, \text{at})$: Takes as input a randomized prepared query execution plan, a secret key and an access token. Issues a *IAM.CheckIfActive* request to validate if the access token is still active. In case of a positive result, executes *SSE-SPARQL-EVAL.EvaluateQuery* and outputs the results.

SEnTriS is distributed, multi-party protocols between a *client*, a *IAM provider*, a *vault*, a *triplestore* and a secure computation *proxy* component. The client is the public facing component, that users interact with. It provides a unified view of the system (hiding the internal communication logic and interaction between the other components) and abstracts the logic for generating access tokens and locking resources. It is also responsible for the encryption and decryption of data, as well as, generation of trapdoors. During query evaluation, it constructs, and prepares, the query execution plan and computes the final results (applying solution modifiers if necessary). The IAM provider validates sessions, manages user data and triplestore access policies, stores role and permission change requests, validates permission rules, emits session cookies and ephemeral access tokens, validates access requests and controls the locking and unlocking of resources. The vault component stores and distributes triplestore secrets. Search and storage of encrypted data is scoped to the triplestore component. Finally, the secure computation proxy evaluates query execution plans, over encrypted bindings.

T¹³13) *TODO*: Add figure for SEnTriS

3.3 Summary

In this Chapter we have explained all the building blocks that compose [SEnTriS](#).

We gave an overview of the system model, and the security model. We specified the core elements that are part of the various protocols and explained the cryptographic tools used.

Furthermore, we defined two [SPARQL](#) query evaluation protocols, SPARQL-EVAL and SPARQL-EVAL-II. In the latter we apply query rewriting techniques to retain reasoning capabilities.

Additionally, we defined a [SSE](#) scheme for encrypting triple patterns, SSE-TP, and explained how to adapt the SPARQL-EVAL-II protocol to work with SSE-TP encrypted triplestores, defining the SSE-SPARQL-EVAL protocol.

Finally, we specified a [IAM](#) protocol, necessary for the functioning of the SSE-SPARQL-EVAL protocol in a [M/M](#) setting.

We concluded with the definition of the [SEnTriS](#) protocol.

IMPLEMENTATION

4.1 General Considerations

4.1.1 Apache JENA Limitations

4.1.2 Memory Constraints

4.1.3 Design & Architecture Choices

Each individual component is a web server that exposes a [REST API](#). All communication channels between components are secure and reliable. They follow the [HTTPS](#) [40] protocol and [TLS](#) [94] data encryption with mutual authentication.

4.2 Component Specific Implementation Details

4.2.1 Client

4.2.1.1 Exposed API

4.2.1.2 Architecture

4.2.1.3 Implementation Details

4.2.2 Triplestore

4.2.2.1 Exposed API

4.2.2.2 Architecture

4.2.2.3 Implementation Details

Redis (Key Structure & Lua Scripts)

Atomicity

4.2.3 IAM Provider

4.2.3.1 Exposed API

4.2.3.2 Architecture

4.2.3.3 Implementation Details

Redis (Key Structure & Lua Scripts)

4.2.4 Proxy

4.2.4.1 Exposed API

4.2.4.2 Architecture

4.2.4.3 Implementation Details

Redis (Key Structure)

Parallelization (SEnTriS V2)

4.2.5 Vault

4.2.5.1 Exposed API

4.2.5.2 Architecture

4.2.5.3 Implementation Details

Redis (Key Structure)

4.2.6 Discussion

4.3 Summary

SYSTEM EVALUATION

5.1 Experimental Evaluation

5.1.1 Setup & Methodology

5.1.1.1 LUBM Benchmark adaptation

5.1.1.2 Experiments

5.1.2 Experimental Results

5.1.2.1 Repository Size & (Up)load Time

5.1.2.2 Query Response Time

5.1.2.3 Query Completeness & Soundness

5.1.2.4 Combined Metric

5.1.2.5 Result Analysis

5.2 Security Analysis

In this Section we will discuss the security of [SEnTriS](#), taking into consideration its security model, defined in Section 3.1.3. We will start by enumerating the security properties of the different versions.

5.2.1 Security Properties

We expect all implementations of [SEnTriS](#) to possess the following *security properties*:

1. A user will be able to create triplestores and access tokens, if and only if they have an active session.
2. A user will be able to interact with a triplestore, if and only if they have a valid access token.
3. A user will not be able to query triplestores which they have not been given access to.
4. A user will only be able to execute operations according to its role.

5. An external observer will not be able to learn any useful information from eavesdropping messages sent between components.
6. An attacker will not be able to impersonate any component.
7. An attacker will not be able to tamper with messages sent between components.
8. An attacker with access to the *Triplestore* component will learn nothing more than the size of a triplestore and its *search* and *access* patterns.
9. Regarding *term equality* during intermediate set operations between *bindings*, an attacker with access to the *Proxy* component will learn nothing more than:
 - a) *SEnTriS V1*: the equality of all terms, belonging to a series of queries, targeting the same triplestore.
 - b) *SEnTriS V2* the equality of all terms present in the graph patterns of individual queries. The bindings of different queries are indistinguishable.

5.2.2 Leakage Analysis

5.2.2.1 SEnTriS V1

5.2.2.2 SEnTriS V2

5.2.3 Security Properties Proofs

5.3 Complexity Analysis

5.3.1 Storage Complexity

5.3.1.1 SEnTriS V1

5.3.1.2 SEnTriS V2

5.3.2 Query Evaluation Complexity

5.3.2.1 SEnTriS V1

5.3.2.2 SEnTriS V2

5.4 Final Discussion

Both implementations come with a trade-off between query evaluation leakage and performance.

SEnTriS V1 provides both data confidentiality and access control, but suffers from the same leakage problems of *CryptDB*, where an adversary can associate the leakage produced by multiple queries targeting the same triplestore. This facilitates the computation of the frequency distribution of all observed terms, which simplifies the execution of *leakage-abuse* attacks [43, 38, 83]. This can be especially problematic during *SPARQL* queries evaluation, as they tend to require a high volume of equality checks. A situation further aggravated when we apply query rewriting techniques: the modified queries usually require an even higher number of equality joins.

In *SEnTriS V2* provides both data confidentiality, access control and reduced leakage. This is achieved at the cost of: 1) an increase in complexity, during set operations between binding tables: 2) an extra communication round during query evaluation, where we

have to reverse the masking applied to the *equality tags* and need to fetch, and decrypt, the node values to generate the final results.

CONCLUSION & FUTURE WORK

BIBLIOGRAPHY

- [1] F. Abel et al. “Enabling advanced and context-dependent access control in RDF stores”. In: *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)*. Vol. 4825 LNCS. 2007, pp. 1–14. ISBN: 9783540762973. DOI: [10.1007/978-3-540-76298-0_1](https://doi.org/10.1007/978-3-540-76298-0_1) (cit. on p. 30).
- [2] N. Aburawi, A. Lisitsa, and F. Coenen. “Querying Encrypted Graph Databases”. In: *ICISSP 2018 - Proceedings of the 4th International Conference on Information Systems Security and Privacy*. Vol. 2018-January. 2018. DOI: [10.5220/0006660004470451](https://doi.org/10.5220/0006660004470451) (cit. on pp. 4, 26, 29).
- [3] A. Acar et al. “A survey on homomorphic encryption schemes: Theory and implementation”. In: vol. 51. 2018. DOI: [10.1145/3214303](https://doi.org/10.1145/3214303) (cit. on pp. 3, 20).
- [4] Amazon Web Services, Inc. or its affiliates. *Amazon Neptune*. <https://docs.aws.amazon.com/neptune/index.html>. Accessed: 2022-2-17 (cit. on p. 19).
- [5] G. Antoniou et al. “A Semantic Web Primer (Third Edition)”. In: *The MIT Press* (2012), p. 287 (cit. on pp. 7, 8, 11, 12, 14).
- [6] N. Ashish. “Semantic-Web Technology: Applications at NASA”. In: *Dagstuhl Seminar Proceedings*. 2005 (cit. on p. 2).
- [7] T. Berners-Lee, J. Hendler, and O. Lassila. *The semantic web*. Vol. 284. 2001. DOI: [10.1038/scientificamerican0501-34](https://doi.org/10.1038/scientificamerican0501-34) (cit. on p. 1).
- [8] M. Blaze, G. Bleumer, and M. Strauss. “Divertible protocols and atomic proxy cryptography”. In: *Advances in Cryptology — EUROCRYPT’98*. Springer Berlin Heidelberg, 1998, pp. 127–144. DOI: [10.1007/BFb0054122](https://doi.org/10.1007/BFb0054122) (cit. on p. 23).
- [9] D. Boneh, A. Sahai, and B. Waters. “Functional Encryption: Definitions and Challenges”. In: *Cryptology ePrint Archive* (2010) (cit. on p. 26).
- [10] C. Bösch et al. “A Survey of Provably Secure Searchable Encryption”. In: *ACM Comput. Surv.* 47.2 (Aug. 2014), pp. 1–51. ISSN: 0360-0300. DOI: [10.1145/2636328](https://doi.org/10.1145/2636328) (cit. on pp. 2, 3, 20, 22, 23).
- [11] Z. Brakerski, C. Gentry, and V. Vaikuntanathan. “(Leveled) Fully Homomorphic Encryption without Bootstrapping”. In: *ACM Trans. Comput. Theory* 6.3 (July 2014), pp. 1–36. ISSN: 1942-3454. DOI: [10.1145/2633600](https://doi.org/10.1145/2633600) (cit. on pp. 3, 20).
- [12] N. Cao et al. “Privacy-preserving query over encrypted graph-structured data in cloud computing”. In: *Proceedings - International Conference on Distributed Computing Systems*. 2011. DOI: [10.1109/ICDCS.2011.84](https://doi.org/10.1109/ICDCS.2011.84) (cit. on p. 25).

- [13] J. J. Carroll. “Signing RDF graphs”. In: *Lect. Notes Comput. Sci.* 2870 (2003), pp. 369–384. issn: 0302-9743, 1611-3349. doi: [10.1007/978-3-540-39718-2_24](https://doi.org/10.1007/978-3-540-39718-2_24) (cit. on p. 30).
- [14] M. Carroll, A. Van Der Merwe, and P. Kotzé. “Secure cloud computing: Benefits, risks and controls”. In: *2011 Information Security for South Africa - Proceedings of the ISSA 2011 Conference*. 2011. doi: [10.1109/ISSA.2011.6027519](https://doi.org/10.1109/ISSA.2011.6027519) (cit. on p. 2).
- [15] M. Chase and S. Kamara. “Structured encryption and controlled disclosure”. In: *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)*. Vol. 6477 LNCS. 2010. doi: [10.1007/978-3-642-17373-8_33](https://doi.org/10.1007/978-3-642-17373-8_33) (cit. on pp. 3, 20, 25).
- [16] B. Chris and S. Andreas. *Berlin SPARQL Benchmark*. <http://wbsg.informatik.uni-mannheim.de/bizer/berlinsparqlbenchmark/spec/index.html>. Accessed: 2022-2-17 (cit. on p. 19).
- [17] R. Ciucanu and P. Lafourcade. “GOOSE: A Secure Framework for Graph Outsourcing and SPARQL Evaluation”. In: *Data and Applications Security and Privacy XXXIV*. Springer International Publishing, 2020, pp. 347–366. doi: [10.1007/978-3-030-49669-2_20](https://doi.org/10.1007/978-3-030-49669-2_20) (cit. on pp. 4, 28, 30).
- [18] L. Costabello, S. Villata, and F. Gandon. “Context-aware access control for RDF graph stores”. In: *Frontiers in Artificial Intelligence and Applications*. Vol. 242. IOS Press, 2012, pp. 282–287. isbn: 9781614990970. doi: [10.3233/978-1-61499-098-7-282](https://doi.org/10.3233/978-1-61499-098-7-282) (cit. on p. 30).
- [19] L. Costabello et al. “Access control for HTTP operations on linked data”. In: *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)*. Vol. 7882 LNCS. Springer Verlag, 2013, pp. 185–199. isbn: 9783642382871. doi: [10.1007/978-3-642-38288-8_13](https://doi.org/10.1007/978-3-642-38288-8_13) (cit. on p. 30).
- [20] O. Curé and G. Blin. *RDF Database Systems: Triples Storage and SPARQL Query Processing*. en. Morgan Kaufmann, Nov. 2014. isbn: 9780128004708 (cit. on pp. 15, 16).
- [21] R. Curtmola et al. “Searchable symmetric encryption: Improved definitions and efficient constructions”. In: *J. Comput. Secur.* 19.5 (Nov. 2011), pp. 895–934. issn: 0926-227X, 1875-8924. doi: [10.3233/jcs-2011-0426](https://doi.org/10.3233/jcs-2011-0426) (cit. on pp. 23–25).
- [22] I. Damgård, M. Geisler, and M. Kroigard. *A correction to ‘efficient and secure comparison for on-line auctions’*. 2009. doi: [10.1504/ijact.2009.028031](https://doi.org/10.1504/ijact.2009.028031) (cit. on pp. 5, 20, 36, 46).
- [23] I. Damgård, M. Geisler, and M. Krøigaard. “Efficient and Secure Comparison for On-Line Auctions”. In: *Information Security and Privacy*. Springer Berlin Heidelberg, 2007, pp. 416–430. doi: [10.1007/978-3-540-73458-1_30](https://doi.org/10.1007/978-3-540-73458-1_30) (cit. on pp. 5, 20, 36, 46).
- [24] C. Dong, G. Russello, and N. Dulay. *Shared and Searchable Encrypted Data for Untrusted Servers*. 2008. doi: [10.1007/978-3-540-70567-3_10](https://doi.org/10.1007/978-3-540-70567-3_10) (cit. on p. 23).
- [25] M. Dürst and M. Suignard. *Internationalized Resource Identifiers (IRIs)*. en. Tech. rep. 2005. doi: [10.17487/RFC3987](https://doi.org/10.17487/RFC3987) (cit. on p. 6).
- [26] T. Elgamal. “A public key cryptosystem and a signature scheme based on discrete logarithms”. In: *IEEE Trans. Inf. Theory* 31.4 (July 1985), pp. 469–472. issn: 0018-9448, 1557-9654. doi: [10.1109/TIT.1985.1057074](https://doi.org/10.1109/TIT.1985.1057074) (cit. on p. 20).

- [27] D. Evans, V. Kolesnikov, and M. Rosulek. “A Pragmatic Introduction to Secure Multi-Party Computation”. In: *Foundations and Trends® in Privacy and Security* 2.2-3 (2018), pp. 70–246. ISSN: 2474-1558. DOI: [10.1561/33000000019](https://doi.org/10.1561/33000000019) (cit. on pp. 2, 19).
- [28] J. D. Fernández et al. “Self-Enforcing Access Control for Encrypted RDF”. In: *The Semantic Web*. Springer International Publishing, 2017, pp. 607–622. DOI: [10.1007/978-3-319-58068-5_37](https://doi.org/10.1007/978-3-319-58068-5_37) (cit. on pp. 4, 26, 29).
- [29] B. Ferreira et al. “Boolean Searchable Symmetric Encryption with Filters on Trusted Hardware”. In: *IEEE Trans. Dependable Secure Comput.* (2020), pp. 1–1. ISSN: 1941-0018. DOI: [10.1109/TDSC.2020.3012100](https://doi.org/10.1109/TDSC.2020.3012100) (cit. on p. 3).
- [30] G. Flouris et al. “Controlling access to RDF graphs”. In: *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)*. Vol. 6369 LNCS. 2010, pp. 107–117. ISBN: 9783642158766. DOI: [10.1007/978-3-642-15877-3_12](https://doi.org/10.1007/978-3-642-15877-3_12) (cit. on p. 30).
- [31] C. Fontaine and F. Galand. “A survey of homomorphic encryption for nonspecialists”. In: *Eurasip Journal on Information Security* 2007 (2007). ISSN: 2510-523X. DOI: [10.1155/2007/13801](https://doi.org/10.1155/2007/13801) (cit. on pp. 3, 20).
- [32] M. Franz. *Secure Computations on Non-Integer Values*. 2011 (cit. on pp. 5, 20, 21, 36, 46).
- [33] *Gartner forecasts worldwide public cloud end-user spending to reach nearly \$600 billion in 2023*. en. <https://www.gartner.com/en/newsroom/press-releases/2022-10-31-gartner-forecasts-worldwide-public-cloud-end-user-spending-to-reach-nearly-600-billion-in-2023>. Accessed: 2023-8-21 (cit. on p. 2).
- [34] C. Gentry. *A Fully Homomorphic Encryption Scheme*. 2009 (cit. on p. 20).
- [35] O. Goldreich. *Foundations of Cryptography: Volume 2, Basic Applications*. en. Cambridge University Press, 2001. ISBN: 9780521830843 (cit. on pp. 2, 19).
- [36] O. Goldreich and R. Ostrovsky. “Software protection and simulation on oblivious RAMs”. In: *J. ACM* 43.3 (May 1996), pp. 431–473. ISSN: 0004-5411. DOI: [10.1145/233551.233553](https://doi.org/10.1145/233551.233553) (cit. on pp. 3, 20).
- [37] Y. Guo, Z. Pan, and J. Heflin. “LUBM: A benchmark for OWL knowledge base systems”. In: *Journal of Web Semantics* 3.2 (Oct. 2005), pp. 158–182. ISSN: 1570-8268. DOI: [10.1016/j.websem.2005.06.005](https://doi.org/10.1016/j.websem.2005.06.005) (cit. on p. 19).
- [38] F. Hahn, N. Loza, and F. Kerschbaum. “Joins Over Encrypted Data with Fine Granular Security”. In: *2019 IEEE 35th International Conference on Data Engineering (ICDE)*. Apr. 2019, pp. 674–685. DOI: [10.1109/ICDE.2019.00066](https://doi.org/10.1109/ICDE.2019.00066) (cit. on pp. 29, 63).
- [39] M. A. Harris et al. “The Gene Ontology (GO) database and informatics resource”. en. In: *Nucleic Acids Res.* 32.Database issue (Jan. 2004), pp. D258–61. ISSN: 0305-1048, 1362-4962. DOI: [10.1093/nar/gkh036](https://doi.org/10.1093/nar/gkh036) (cit. on p. 2).
- [40] *HTTP Over TLS*. en. <https://datatracker.ietf.org/doc/html/rfc2818>. Accessed: 2022-2-18 (cit. on p. 60).
- [41] A.-A. Ivan and Y. Dodis. “Proxy Cryptography Revisited”. In: *NDSS*. 2003 (cit. on p. 23).

- [42] A. Jain and C. Farkas. “Secure resource description framework: An access control model”. In: *Proceedings of ACM Symposium on Access Control Models and Technologies, SACMAT*. Vol. 2006. 2006, pp. 121–129. ISBN: 9781595933546 (cit. on p. 30).
- [43] C. Jutla and S. Patranabis. “Efficient Searchable Symmetric Encryption for Join Queries”. In: *Cryptology ePrint Archive* (2021) (cit. on pp. 29, 63).
- [44] S. Kamara, C. Papamanthou, and T. Roeder. “Dynamic searchable symmetric encryption”. In: *Proceedings of the 2012 ACM conference on Computer and communications security*. CCS ’12. Raleigh, North Carolina, USA: Association for Computing Machinery, Oct. 2012, pp. 965–976. ISBN: 9781450316514. DOI: [10.1145/2382196.2382298](https://doi.org/10.1145/2382196.2382298) (cit. on p. 25).
- [45] A. Kasten and A. Scherp. *Iterative Signing of RDF(S) Graphs, Named Graphs, and OWL Graphs: Formalization and Application*. Tech. rep. Mainz: Universität Koblenz-Landau, Mar. 2013 (cit. on p. 30).
- [46] A. Kasten and A. Scherp. “Towards a configurable framework for iterative signing of distributed graph data”. In: *CEUR Workshop Proceedings*. Vol. 1121. 2014 (cit. on p. 30).
- [47] A. Kasten, A. Scherp, and P. Schauß. “A framework for iterative signing of graph data on the web”. In: *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)*. Vol. 8465 LNCS. Springer Verlag, 2014, pp. 146–160. ISBN: 9783319074429. DOI: [10.1007/978-3-319-07443-6_11](https://doi.org/10.1007/978-3-319-07443-6_11) (cit. on p. 30).
- [48] S. Kirrane et al. “Secure manipulation of linked data”. In: *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)*. Vol. 8218 LNCS. 2013. DOI: [10.1007/978-3-642-41335-3_16](https://doi.org/10.1007/978-3-642-41335-3_16) (cit. on p. 30).
- [49] G. Ladwig. *Efficient Optimization and Processing of Queries Over Text-rich Graph-structured Data*. en. KIT Scientific Publishing, 2013. ISBN: 9783731500155. DOI: [10.5445/KSP/1000034423](https://doi.org/10.5445/KSP/1000034423) (cit. on p. 16).
- [50] P. Leach, M. Mealling, and R. Salz. *A Universally Unique Identifier (UUID) URN Namespace*. Tech. rep. July 2005. DOI: [10.17487/rfc4122](https://doi.org/10.17487/rfc4122) (cit. on p. 33).
- [51] A. Letelier et al. “Static analysis and optimization of semantic web queries”. In: *ACM Trans. Database Syst.* 38.4 (Dec. 2013), pp. 1–45. ISSN: 0362-5915. DOI: [10.1145/2500130](https://doi.org/10.1145/2500130) (cit. on p. 16).
- [52] M. X. Makkes. *Efficient implementation of homomorphic cryptosystems*. 2010 (cit. on pp. 5, 36, 46).
- [53] A. Maña, G. Pujol, and C. Pandolfo. “A distributed secure ontology for certified services oriented applications”. In: *Proceedings - 5th IEEE International Conference on Semantic Computing, ICSC 2011*. 2011. DOI: [10.1109/ICSC.2011.105](https://doi.org/10.1109/ICSC.2011.105) (cit. on p. 30).
- [54] MarkLogic. *MarkLogic server - multi-model database*. en. <https://www.marklogic.com/product/marklogic-database-overview/>. Accessed: 2022-2-17. Feb. 2020 (cit. on p. 19).
- [55] J. P. McCrae. *The linked open data cloud*. en. <https://lod-cloud.net/>. Accessed: 2022-2-8 (cit. on p. 2).

- [56] B. D. McKay and A. Piperno. “Practical graph isomorphism, II”. In: *J. Symbolic Comput.* 60 (Jan. 2014), pp. 94–112. ISSN: 0747-7171. DOI: [10.1016/j.jsc.2013.09.003](https://doi.org/10.1016/j.jsc.2013.09.003) (cit. on p. 16).
- [57] G. A. Miller. “WordNet: a lexical database for English”. In: *Commun. ACM* 38.11 (Nov. 1995), pp. 39–41. ISSN: 0001-0782. DOI: [10.1145/219717.219748](https://doi.org/10.1145/219717.219748) (cit. on p. 2).
- [58] G. E. Modoni, M. Sacco, and W. Terkaj. “A survey of RDF store solutions”. In: *2014 International Conference on Engineering, Technology and Innovation: Engineering Responsible Innovation in Products and Services, ICE 2014*. 2014. DOI: [10.1109/ICE.2014.6871541](https://doi.org/10.1109/ICE.2014.6871541) (cit. on p. 15).
- [59] *MongoDB Manual: Queryable encryption*. en. <https://www.mongodb.com/docs/v7.0/core/queryable-encryption/>. Accessed: 2023-8-21 (cit. on p. 3).
- [60] M. A. Musen and Protégé Team. “The Protégé Project: A Look Back and a Look Forward”. en. In: *AI Matters* 1.4 (June 2015), pp. 4–12. ISSN: 2372-3483. DOI: [10.1145/2757001.2757003](https://doi.org/10.1145/2757001.2757003) (cit. on pp. 2, 19).
- [61] R. Neches et al. “Enabling technology for knowledge sharing”. In: *AI Magazine* 12.3 (1991). ISSN: 0738-4602 (cit. on p. 1).
- [62] R. Nejahi, A. Marzouk, and N. Gherabi. “An Enhanced Method for Web Ontologies Encryption”. In: *European Journal of Electrical Engineering and Computer Science* 3.4 (July 2019). ISSN: 2506-9853. DOI: [10.24018/ejece.2019.3.4.99](https://doi.org/10.24018/ejece.2019.3.4.99) (cit. on p. 25).
- [63] Neo4j, Inc. *Cypher Query Language*. en. <https://neo4j.com/developer/cypher/>. Accessed: 2023-8-24 (cit. on pp. 4, 26, 28).
- [64] *NIST Special Publication 800-63B*. <https://pages.nist.gov/800-63-3/sp800-63b.html>. Accessed: 2024-9-18 (cit. on p. 55).
- [65] N. Noy et al. “Industry-scale knowledge graphs: Lessons and challenges”. In: *Commun. ACM* 62.8 (2019). ISSN: 0001-0782, 1557-7317. DOI: [10.1145/3331166](https://doi.org/10.1145/3331166) (cit. on p. 1).
- [66] Ontotext. *GraphDB*. en. <https://graphdb.ontotext.com/documentation/>. Accessed: 2022-2-17 (cit. on p. 19).
- [67] OpenLink Software. *Virtuoso Universal Server*. <https://virtuoso.openlinksw.com/>. Accessed: 2022-2-17 (cit. on p. 19).
- [68] owl.cs. *OWL API repository*. en. <https://github.com/owlcs/owlapi>. Accessed: 2022-2-17 (cit. on p. 19).
- [69] P. Paillier. “Public-Key Cryptosystems Based on Composite Degree Residuosity Classes”. In: *Advances in Cryptology — EUROCRYPT ’99*. Springer Berlin Heidelberg, 1999, pp. 223–238. DOI: [10.1007/3-540-48910-X_16](https://doi.org/10.1007/3-540-48910-X_16) (cit. on pp. 20, 28, 36).
- [70] V. Papakonstantinou et al. “Access control for RDF graphs using abstract models”. In: *Proceedings of ACM Symposium on Access Control Models and Technologies, SACMAT*. 2012. DOI: [10.1145/2295136.2295155](https://doi.org/10.1145/2295136.2295155) (cit. on p. 30).
- [71] W. Pearson. *Cryptography and Network Security: Principles and Practice, Global Edition*. 7th ed. Pearson Education, 2017. ISBN: 9781292158587 (cit. on pp. 3, 24, 26).

- [72] Z. Peng, S. Tang, and L. Jiang. *A Symmetric Authenticated Proxy Re-encryption Scheme with Provable Security*. 2017. DOI: [10.1007/978-3-319-68542-7_8](https://doi.org/10.1007/978-3-319-68542-7_8) (cit. on p. 23).
- [73] F. Picalausa and S. Vansummeren. “What are real SPARQL queries like?” In: *Proceedings of the International Workshop on Semantic Web Information Management*. SWIM ’11 Article 7. Athens, Greece: Association for Computing Machinery, June 2011, pp. 1–6. ISBN: 9781450306515. DOI: [10.1145/1999299.1999306](https://doi.org/10.1145/1999299.1999306) (cit. on pp. 15–17).
- [74] G. S. Poh, M. S. Mohamad, and M. R. Z’Aba. “Structured encryption for conceptual graphs”. In: *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)*. Vol. 7631 LNCS. Springer Verlag, 2012, pp. 105–122. ISBN: 9783642341168. DOI: [10.1007/978-3-642-34117-5_7](https://doi.org/10.1007/978-3-642-34117-5_7) (cit. on pp. 4, 25).
- [75] R. A. Popa, N. Zeldovich, and H. Balakrishnan. “CryptDB : A Practical Encrypted Relational DBMS”. In: *Designfax* (2011). ISSN: 0163-6669 (cit. on pp. 4, 5, 26, 32, 46).
- [76] R. A. Popa et al. “CryptDB: Processing queries on an encrypted database”. In: *Communications of the ACM*. Vol. 55. 2012. DOI: [10.1145/2330667.2330691](https://doi.org/10.1145/2330667.2330691) (cit. on pp. 4, 5, 21, 26, 32, 46).
- [77] R. A. Popa et al. “CryptDB: Protecting confidentiality with encrypted query processing”. In: *SOSP’11 - Proceedings of the 23rd ACM Symposium on Operating Systems Principles*. 2011. DOI: [10.1145/2043556.2043566](https://doi.org/10.1145/2043556.2043566) (cit. on pp. 4, 5, 26, 32, 46).
- [78] *RDF 1.1: On semantics of RDF datasets*. en. <https://www.w3.org/TR/rdf11-datasets/>. Accessed: 2023-10-14 (cit. on p. 15).
- [79] R. L. Rivest, L. Adleman, and M. L. Dertouzos. “On data banks and privacy homomorphisms”. In: *Foundations of secure* (1978) (cit. on p. 20).
- [80] R. L. Rivest, A. Shamir, and L. Adleman. “A method for obtaining digital signatures and public-key cryptosystems”. en. In: *Commun. ACM* 21.2 (Feb. 1978), pp. 120–126. ISSN: 0001-0782, 1557-7317. DOI: [10.1145/359340.359342](https://doi.org/10.1145/359340.359342) (cit. on p. 20).
- [81] K. Sakurai, T. Nishide, and A. Syalim. *Improved proxy re-encryption scheme for symmetric key cryptography*. 2017. DOI: [10.1109/iwbis.2017.8275110](https://doi.org/10.1109/iwbis.2017.8275110) (cit. on p. 23).
- [82] Security Compass. *The 2021 State of Cloud Adoption*. en. <https://resources.securitycompass.com/reports/2021-state-of-cloud-adoption>. Accessed: 2022-2-9. Sept. 2021 (cit. on p. 2).
- [83] M. Shafieinejad et al. “Equi-Joins Over Encrypted Data for Series of Queries”. In: (Mar. 2021). arXiv: [2103.05792 \[cs.CR\]](https://arxiv.org/abs/2103.05792) (cit. on pp. 29, 63).
- [84] N. Sioutos et al. “NCI Thesaurus: a semantic model integrating cancer-related clinical and molecular information”. en. In: *J. Biomed. Inform.* 40.1 (Feb. 2007), pp. 30–43. ISSN: 1532-0464, 1532-0480. DOI: [10.1016/j.jbi.2006.02.013](https://doi.org/10.1016/j.jbi.2006.02.013) (cit. on p. 2).
- [85] E. Sirin et al. “Pellet: A practical OWL-DL reasoner”. In: *Journal of Web Semantics* 5.2 (June 2007), pp. 51–53. ISSN: 1570-8268. DOI: [10.1016/j.websem.2007.03.004](https://doi.org/10.1016/j.websem.2007.03.004) (cit. on p. 19).

- [86] D. X. Song, D. Wagner, and A. Perrig. “Practical techniques for searches on encrypted data”. In: *Proceeding 2000 IEEE Symposium on Security and Privacy. S P 2000*. May 2000, pp. 44–55. DOI: [10.1109/SECPRI.2000.848445](https://doi.org/10.1109/SECPRI.2000.848445) (cit. on pp. 3, 22, 28).
- [87] D. X. Song, D. Wagner, and A. Perrig. “Practical techniques for searches on encrypted data”. In: *Proceeding 2000 IEEE Symposium on Security and Privacy. S P 2000*. May 2000, pp. 44–55. DOI: [10.1109/SECPRI.2000.848445](https://doi.org/10.1109/SECPRI.2000.848445) (cit. on p. 24).
- [88] R. Studer, V. R. Benjamins, and D. Fensel. “Knowledge engineering: Principles and methods”. In: *Data Knowl. Eng.* 25.1 (Mar. 1998), pp. 161–197. ISSN: 0169-023X. DOI: [10.1016/S0169-023X\(97\)00056-6](https://doi.org/10.1016/S0169-023X(97)00056-6) (cit. on p. 10).
- [89] G. Suntaxi, A. A. El Ghazi, and K. Böhm. “Secrecy and performance models for query processing on outsourced graph data”. In: *Distributed and Parallel Databases* 39.1 (Mar. 2021), pp. 35–77. ISSN: 1573-7578. DOI: [10.1007/s10619-020-07284-0](https://doi.org/10.1007/s10619-020-07284-0) (cit. on p. 25).
- [90] A. Sutton and R. Samavi. “Timestamp-based Integrity Proofs for Linked Data”. In: *Proceedings of the International Workshop on Semantic Big Data, SBD 2018 - In conjunction with the 2018 ACM SIGMOD/PODS Conference*. Association for Computing Machinery, Inc, June 2018. ISBN: 9781450357791. DOI: [10.1145/3208352.3208353](https://doi.org/10.1145/3208352.3208353) (cit. on p. 30).
- [91] A. Syalim, T. Nishide, and K. Sakurai. “Realizing Proxy Re-encryption in the Symmetric World”. In: *Informatics Engineering and Information Science*. Springer Berlin Heidelberg, 2011, pp. 259–274. DOI: [10.1007/978-3-642-25327-0_23](https://doi.org/10.1007/978-3-642-25327-0_23) (cit. on p. 23).
- [92] The Apache Software Foundation. *Apache Jena*. en. <https://jena.apache.org/>. Accessed: 2022-2-17 (cit. on pp. 19, 40, 42).
- [93] The Apache Software Foundation. *Apache Jena - TDB*. en. <https://jena.apache.org/documentation/tdb/index.html>. Accessed: 2022-2-17 (cit. on p. 19).
- [94] *The Transport Layer Security (TLS) Protocol Version 1.3*. en. <https://datatracker.ietf.org/doc/html/rfc8446>. Accessed: 2022-2-18 (cit. on p. 60).
- [95] T. Tudorache, J. Vendetti, and N. F. Noy. “Web-Protégé: A lightweight OWL ontology editor for the web”. In: *CEUR Workshop Proceedings*. Vol. 432. 2009 (cit. on pp. 2, 19).
- [96] T. Tudorache et al. “Supporting collaborative ontology development in Protégé”. In: *Lect. Notes Comput. Sci.* 5318 LNCS (2008), pp. 17–32. ISSN: 0302-9743, 1611-3349. DOI: [10.1007/978-3-540-88564-1_2](https://doi.org/10.1007/978-3-540-88564-1_2) (cit. on pp. 2, 19).
- [97] G. Tummarello et al. “Signing individual fragments of an RDF graph”. In: *14th International World Wide Web Conference, WWW2005*. 2005, pp. 1020–1021. ISBN: 9781595930514. DOI: [10.1145/1062745.1062848](https://doi.org/10.1145/1062745.1062848) (cit. on p. 30).
- [98] W3C. *Extensible markup language (XML) 1.0 (fifth edition)*. en. <https://www.w3.org/TR/xml1/>. Accessed: 2022-2-11 (cit. on p. 7).
- [99] W3C. *OWL 2 web ontology language direct semantics (second edition)*. en. <https://www.w3.org/TR/2012/REC-owl2-direct-semantics-20121211/>. Accessed: 2022-2-14 (cit. on p. 11).

-
- [100] W3C. *OWL 2 web ontology language document overview (second edition)*. en. <https://www.w3.org/TR/owl2-overview/>. Accessed: 2022-2-11 (cit. on pp. viii, 1, 7, 11).
 - [101] W3C. *OWL 2 web ontology language Manchester syntax (second edition)*. en. <https://www.w3.org/TR/2012/NOTE-owl2-manchester-syntax-20121211/>. Accessed: 2022-2-14 (cit. on p. 11).
 - [102] W3C. *OWL 2 web ontology language profiles (second edition)*. en. <https://www.w3.org/TR/2012/REC-owl2-profiles-20121211/>. Accessed: 2022-2-14 (cit. on p. 11).
 - [103] W3C. *OWL 2 web ontology language quick reference guide (second edition)*. en. <https://www.w3.org/TR/owl2-quick-reference/>. Accessed: 2022-2-14 (cit. on p. 12).
 - [104] W3C. *OWL 2 web ontology language RDF-based semantics (second edition)*. en. <https://www.w3.org/TR/2012/REC-owl2-rdf-based-semantics-20121211/>. Accessed: 2022-2-14 (cit. on p. 11).
 - [105] W3C. *OWL 2 web ontology language structural specification and functional-style syntax (second edition)*. en. <https://www.w3.org/TR/2012/REC-owl2-syntax-20121211/>. Accessed: 2022-2-14 (cit. on p. 11).
 - [106] W3C. *OWL 2 web ontology language XML serialization (second edition)*. en. <https://www.w3.org/TR/2012/REC-owl2-xml-serialization-20121211/>. Accessed: 2022-2-14 (cit. on p. 11).
 - [107] W3C. *RDF 1.1 concepts and abstract syntax*. en. <https://www.w3.org/TR/rdf11-concepts/>. Accessed: 2022-2-11 (cit. on pp. 1, 7).
 - [108] W3C. *RDF 1.1 TriG*. en. <https://www.w3.org/TR/trig/>. Accessed: 2022-2-13 (cit. on p. 8).
 - [109] W3C. *RDF 1.1 turtle*. en. <https://www.w3.org/TR/turtle/>. Accessed: 2022-2-13 (cit. on p. 8).
 - [110] W3C. *RDF 1.1 XML syntax*. en. <https://www.w3.org/TR/rdf-syntax-grammar/>. Accessed: 2022-2-13 (cit. on p. 8).
 - [111] W3C. *RDF Schema 1.1*. en. <https://www.w3.org/TR/rdf-schema/>. Accessed: 2022-2-11 (cit. on pp. 1, 7, 9).
 - [112] W3C. *RDFa core 1.1 - third edition*. en. <https://www.w3.org/TR/rdfa-core/>. Accessed: 2022-2-13 (cit. on p. 8).
 - [113] W3C. *Semantic Web Case Studies and Use Cases*. en. <https://www.w3.org/2001/sw/sweo/public/UseCases/>. Accessed: 2022-2-8 (cit. on p. 2).
 - [114] W3C. *SPARQL 1.1 graph store HTTP protocol*. en. <https://www.w3.org/TR/sparql11-http-rdf-update/>. Accessed: 2022-2-4. Mar. 2013 (cit. on p. 15).
 - [115] W3C. *SPARQL 1.1 Overview*. <https://www.w3.org/TR/sparql11-overview/>. Mar. 2013 (cit. on pp. 1, 15, 36).
 - [116] W3C. *SPARQL 1.1 update*. en. <https://www.w3.org/TR/sparql11-update/>. Accessed: 2022-2-15 (cit. on pp. 1, 15, 18, 36).
 - [117] W3C. *SPARQL Query Language for RDF*. <https://www.w3.org/TR/rdf-sparql-query/>. Jan. 2008 (cit. on pp. 1, 15, 17).

BIBLIOGRAPHY

- [118] W3C. *SWRL: A Semantic Web Rule Language Combining OWL and RuleML*. <https://www.w3.org/Submission/SWRL/>. May 2004 (cit. on p. 7).
- [119] *What is SNOMED CT*. en. <https://www.snomed.org/what-is-snomed-ct>. Accessed: 2024-9-24 (cit. on p. 2).
- [120] Wikipedia contributors. *Semantic web stack*. https://en.wikipedia.org/wiki/File:Semantic_web_stack.svg. Accessed: 2022-2-15 (cit. on pp. viii, 6).

SPARQL-EVAL LISTINGS

Listing A.1: Triple Pattern Storage Engine

```

1  State:
2    triplestores ← {}
3    SEPARATORI ← fetchConfig(SEPARATORI) // A configurable parameter, must provide
        ↳ a value that does not belong to the charset used for nodes
4    SEPARATORII ← fetchConfig(SEPARATORII) // Must be different than SEPARATORI,
        ↳ and follow the same restrictions
5    SEPARATORIII ← fetchConfig(SEPARATORIII) // Must be different than SEPARATORI
        ↳ and SEPARATORII, and follow the same restrictions
6
7  function save(triplestoreID, triples):
8    foreach t in triples:
9      saveTriple(t)
10   return "SUCCESS"
11
12  function delete(triplestoreID):
13    triplestores ← triplestores / { triplestores[triplestoreID] }
14    return "SUCCESS"
15
16  function delete(triplestoreID, triples):
17    foreach t in triples:
18      deleteTriple(t)
19    return "SUCCESS"
20
21  function saveTriple(triple=(s,p,o)):
22    parseds ← parseNode(s)
23    parsedp ← parseNode(p)
24    parsedo ← parseNode(o)
25    keywords ← parseKeyword(s)
26    keywordp ← parseKeyword(p)
27    keywordo ← parseKeyword(o)
28    save(triplestoreID, "P0", keywords, parsedp, parsedo)
29    save(triplestoreID, "S0", keywordp, parseds, parsedo)
30    save(triplestoreID, "SP", keywordo, parseds, parsedp)
31    save(triplestoreID, "S", keywordp, keywordo, parseds)
32    save(triplestoreID, "P", keywords, keywordo, parsedp)
33    save(triplestoreID, "O", keywords, keywordp, parsedo)
34
35  function deleteTriple(triple=(s,p,o)):
36    parseds ← parseNode(s)
37    parsedp ← parseNode(p)

```

```

38   parsedo ← parseNode(o)
39   keywords ← parseKeyword(s)
40   keywordp ← parseKeyword(p)
41   keywordo ← parseKeyword(o)
42   delete(triplestoreID, "P0", keywords, parsedp, parsedo)
43   delete(triplestoreID, "S0", keywordp, parseds, parsedo)
44   delete(triplestoreID, "SP", keywordo, parseds, parsedp)
45   delete(triplestoreID, "S", keywordp, keywordo, parseds)
46   delete(triplestoreID, "P", keywords, keywordo, parsedp)
47   delete(triplestoreID, "O", keywords, keywordp, parsedo)
48
49   function save(triplestoreID, varPattern, keyword, node1, node2):
50     index ← generateIndex(varPattern, keyword)
51     triplestores[triplestoreID][index] ← triplestores[triplestoreID][index] ∪ { node1 ||
        ↪ SEPARATORIII || node2 }
52
53   function save(triplestoreID, varPattern, keyword1, keyword2, node):
54     index ← generateIndex(varPattern, keyword1, keyword2)
55     triplestores[triplestoreID][index] ← triplestores[triplestoreID][index] ∪ { node }
56
57   function delete(triplestoreID, varPattern, keyword, node1, node2):
58     index ← generateIndex(varPattern, keyword)
59     triplestores[triplestoreID][index] ← triplestores[triplestoreID][index] / { node1 ||
        ↪ SEPARATORIII || node2 }
60
61   function delete(triplestoreID, varPattern, keyword1, keyword2, node):
62     index ← generateIndex(varPattern, keyword1, keyword2)
63     triplestores[triplestoreID][index] ← triplestores[index] / { node }
64
65   function generateIndex(varPattern, keyword):
66     return SEPARATORII || varPattern || SEPARATORII || keyword
67
68   function generateIndex(varPattern, keyword1, keyword2):
69     return SEPARATORII || varPattern || SEPARATORII || node1 || SEPARATORII || node2
70
71   function search(triplestoreID, varPattern, s, p, o):
72     switch(varPattern):
73       case "S": return fetchBindings(triplestoreID, generateIndex("S", p, o), s)
74       case "P": return fetchBindings(triplestoreID, generateIndex("P", s, o), p)
75       case "O": return fetchBindings(triplestoreID, generateIndex("P", s, p), o)
76       case "P0": return fetchBindings(triplestoreID, generateIndex("P0", s), p, o)
77       case "S0": return fetchBindings(triplestoreID, generateIndex("S0", p), s, o)
78       case "SP": return fetchBindings(triplestoreID, generateIndex("SP", o), s, p)
79       default: throw NotImplementedException()
80
81   function fetchBindings(triplestoreID, index, var):
82     table ← createBindingsTable()
83     table.setVariables({var})
84     foreach node in triplestores[triplestoreID][index]:
85       table.add(generateID(), var, node)
86     return table
87
88   function fetchBindings(triplestoreID, index, var1, var2):
89     table ← createBindingsTable()
90     table.setVariables({var1, var2})
91     foreach nodePair in triplestores[triplestoreID][index]:
92       patternID ← generateID()
93       nodes = nodePair.split(SEPARATORIII)
94       table.add(patternID, var1, nodes[0])
95       table.add(patternID, var2, nodes[1])

```

```

96     return table
97
98 function parseNode(n):
99     if n ∈ I:
100         return "S" || SEPARATORI || n.value // IRI
101     elif n ∈ L:
102         return "L" || SEPARATORI || n.value || SEPARATORI || n.datatype
103     elif n ∈ B:
104         return "B" || SEPARATORI || n.value // identifier of the blank node
105     return ⊥
106
107 function parseKeyword(n):
108     if n ∈ I:
109         return "S" || SEPARATORI || n.value
110     elif n ∈ L:
111         return "L" || SEPARATORI || n.value || SEPARATORI || n.datatype
112     elif n ∈ B:
113         return "BLANK"
114     return ⊥
115
116 function generateNode(parsedValue):
117     (prefix, value, datatype) ← parsedValue.split(SEPARATORI);
118     if prefix == "BLANK":
119         return createBlankNode(value)
120     elif prefix == "S":
121         return createIRINode(value)
122     else
123         return createLiteralNode(value, datatype)

```

```

1 State:
2     bindings ← {}
3     patterns ← {}
4
5 function setVariables(vars):
6     foreach var in vars:
7         bindings[var] ← {}
8         patterns[var] ← {}
9
10 function add(patternID, var, binding):
11     addBinding(binding, var, patternID)
12     addPattern(binding, var, patternID)
13
14 function addPattern(binding, var, patternID):
15     varPatternIDs ← patterns[var]
16     if varPatternIDs == ⊥:
17         varPatternIDs ← {}
18         varPatternIDs[patternID] ← binding
19         patterns[var] ← varPatternIDs
20     else:
21         varPatternIDs[patternID] ← binding
22
23 function addBinding(binding, var, patternID):
24     varBindings = binding[var]
25     if varBindings == ⊥:
26         varBindings ← {}
27         savePatternID(varBindings, binding, patternID)
28         binding[var] ← varBindings
29     else
30         varBindingPatternIDs = varBindings[binding]

```

```
31     if varBindingPatternIDs ==  $\perp$ :
32         savePatternID(varBindings, binding, patternID)
33     else
34         varBindingPatternIDs  $\leftarrow$  varBindingPatternIDs  $\cup$  { patternID }
35
36 function savePatternID(varBindings, binding, patternID):
37     varBindingPatternIDs  $\leftarrow$  {}
38     varBindingPatternIDs  $\leftarrow$  varBindingPatternIDs  $\cup$  { patternID }
39     varBindings[binding]  $\leftarrow$  varBindingPatternIDs
40
41 function getVariables():
42     return bindings.getKeys()
43
44 function getBindings(var):
45     return bindings[var]
46
47 function getBindings():
48     return bindings
49
50 function getPatternIDs(var):
51     return patterns[var]
52
53 function getPatternIDs():
54     return patterns
55
56 function getPatterns():
57     patterns  $\leftarrow$  createLinkedList()
58     vars  $\leftarrow$  getVariables()
59     patternIDs  $\leftarrow$  {}
60     foreach var in vars:
61         patternIDs  $\leftarrow$  patternIDs  $\cup$  patterns[var].getKeys()
62     foreach patternID in patternIDs:
63         pattern  $\leftarrow$  [#vars]
64         i  $\leftarrow$  0
65         foreach var in vars:
66             binding  $\leftarrow$  patterns[var][patternID]
67             pattern[i]  $\leftarrow$  binding
68             if binding ==  $\perp$ :
69                 i++
70         if i < #vars:
71             patterns  $\leftarrow$  patterns  $\cup$  { pattern }
72     return patterns
73
74 function project(vars):
75     bindings.retain(vars)
76     patterns.retain(vars)
77
78 function join(other):
79     varsmutual  $\leftarrow$  copy(getVariables())
80     varsmutual.retain(other.getVariables())
81     return join(varsmutual, this, other, "INNER")
82
83 function leftOuterJoin(other):
84     varsmutual  $\leftarrow$  copy(getVariables())
85     varsmutual.retain(other.getVariables())
86     return join(varsmutual, this, other, "LEFT")
87
88 function union(other):
89     varsl  $\leftarrow$  getVariables()
90     varsr  $\leftarrow$  other.getVariables()
```

```

91  vars ← copy(varsl)
92  vars ← vars ∪ varsr
93  tableunion ← createBindingsTable()
94  tableunion.setVariables(vars)
95  copyAllBindings(varsl, this, tableunion)
96  copyAllBindings(varsr, other, tableunion)
97  return tableunion
98
99  function minus(other):
100  varsmutual ← copy(getVariables())
101  varsmutual.retain(other.getVariables())
102  diff ← {}
103  tableminus ← createBindingsTable()
104  tableminus.setVariables(vars)
105  foreach var : varsmutual:
106    bindingsl ← getBindings(var)
107    bindingsr ← other.getBindings(var)
108    foreach (binding, patternIDs) in bindingsl:
109      if bindingsr[binding] == ⊥:
110        diff ← diff ∪ patternIDs
111  varBindingsRND ← getBindings(sample(varsmutual))
112  foreach (binding, patternIDs) in varBindingsRND:
113    foreach patternID in patternIDs:
114      if patternID ∈ diff:
115        copyBindings(patternID, patternID, varsmutual, this, tableminus)
116  return tableminus
117
118  function copyBindings(newPatternID, oldPatternID, vars, source, target)
119    foreach var in vars
120      binding ← source.getPatternIDs(var)[oldPatternID]
121      if binding ≠ ⊥:
122        target.add(newPatternID, var, binding)
123
124  function copyAllBindings(vars, source, target)
125    foreach var in vars:
126      varBindings ← source.getBindings(var)
127      foreach (binding, patternIDs) in varBindings:
128        foreach patternID in patternIDs:
129          target.add(patternID, var, binding)
130
131  function join(varsmutual, tablel, tabler, type)
132  varsl ← copy(tablel.getVariables())
133  varsr ← copy(tabler.getVariables())
134  vars ← varsl ∪ varsr
135  varsl ← varsl / varsr
136  varsr ← varsr / varsl
137  tablejoin ← createBindingsTable()
138  tablejoin.setVariables(vars)
139  var ← sample(varsmutual)
140  bindingsl = tablel.getBindings(var)
141  bindingsr = tabler.getBindings(var)
142  foreach (binding, patternIDsl) in bindingsl:
143    patternIDsr = bindingsr[binding]
144    if patternIDsr ≠ ⊥ or type == "LEFT":
145      joinPatterns(varsmutual, tablel, varsl, patternIDsl, tabler, varsr, patternIDsr,
146        ↪ tablejoin, type)
147  return tablejoin
148
149  function joinPatterns(varsmutual, tablel, varsl, patternIDsl, tabler, varsr, patternIDsr,
150    ↪ tablejoin, type)

```

```

149  foreach patternIDl in patternIDsl:
150      match ← false
151      if patternIDsr ≠ ⊥:
152          foreach patternIDr in patternIDsr:
153              if equalPatterns(varsmutual, tablel, patternIDl, tabler, patternIDr):
154                  match ← true
155                  patternID = generateID()
156                  copyBindings(patternID, patternIDl, varsmutual, tablel, tablejoin)
157                  copyBindings(patternID, patternIDl, varsl, tablel, tablejoin)
158                  copyBindings(patternID, patternIDr, varsr, tabler, tablejoin)
159              if match == false and type == "LEFT":
160                  copyBindings(patternIDl, patternIDl, varsmutual, tablel, tablejoin)
161                  copyBindings(patternIDl, patternIDl, varsl, tablel, tablejoin)
162
163  function equalPatterns(varsmutual, tablel, patternIDl, tabler, patternIDr)
164      foreach var in varsmutual:
165          bindingl = tablel.getPatternIDs(var)[patternIDl]
166          bindingr = tabler.getPatternIDs(var)[patternIDr]
167          if bindingl == ⊥ and bindingr ≠ ⊥
168              return false
169          elif bindingl ≠ ⊥ and bindingr == ⊥
170              return false
171          elif bindingl ≠ ⊥ and bindingl ≠ bindingr
172              return false
173      return true

```

Listing A.2: Bindings Table

```

1  State:
2      JENAalgebraGenerator ← JENA.createAlgebraGenerator(JENAARQ.getContext()) // the algebra
   ↪ generator compiles a query into operations
3
4  function getQueryPlan(planner, queryString):
5      try:
6          query = JENAQueryFactory.create(queryString) // the query factory parses a string and
   ↪ generates a query
7          planner.setQueryType(query.getType())
8          if planner.getQueryType() == "CONSTRUCT":
9              planner.setConstructTemplate(query.getConstructTemplate().getTriples())
10             return planner.generatePlan(JENAalgebraGenerator.compile(query))
11         catch QueryParseException: // the query factory throws an exception when parsing a
   ↪ string for an update query
12             updateQuery = JENAUpdateFactory.create(queryString) // the update factory parses a
   ↪ string and generates an update query
13             return planner.generatePlan(updateQuery.getOperations().get(0), JENAUpdateFactory)

```

Listing A.3: SPARQL Query Engine

```

1  State:
2      operationsparsed ← {}
3      uploadtemplate ← {}
4      deletetemplate ← {}
5      constructtemplate ← {}
6      queryType ← :bot
7      plan ← (jobs = {}, execOrder = createLinkedList(), vars = {})
8      variables ← {}
9

```

```

10 function getQueryType():
11     return queryType
12
13 function setQueryType(type):
14     queryType ← type
15
16 function setConstructTemplate(triples)
17     constructtemplate ← triples
18
19 function getConstructTemplate:
20     return constructtemplate
21
22 function getUploadTemplate:
23     return uploadtemplate
24
25 function getDeleteTemplate:
26     return deletetemplate
27
28 function pushJob(job):
29     jobID ← job.getID()
30     plan.jobs[jobID] ← job
31     plan.execOrder.addLast(jobID)
32
33 function generatePlan(op)
34     JENA.OpWalker.walk(op, this) // the operation walker will call visit while there are
35     ↪ operations which have not been visited, using a visitor (the query planner is
36     ↪ a JENAVisitor)
37     if #plan.vars == 0:
38         plan.vars ← variables
39     return plan
40
41 function generatePlan(op, JENAalgebraGenerator):
42     if op is UpdateDataInsert:
43         visitUpdate(op)
44     elif op is UpdateDataDelete:
45         visitUpdate(op)
46     elif op is UpdateModify:
47         visitUpdate(op, JENAalgebraGenerator)
48     elif op is UpdateDeleteWhere:
49         visitUpdate(op)
50     else:
51         throw NotImplementedException()
52     if #plan.vars == 0:
53         plan.vars ← variables
54     return plan
55
56 function visit(op)
57     if op is OpBGP:
58         generateSearchJobs(op)
59     elif op is OpTriple:
60         generateSearchJobs(op.asBGP())
61     elif op is OpJoin:
62         generateJoinJob(op)
63     elif op is OpLeftJoin:
64         generateOptionalJob(op)
65     elif op is OpMinus:
66         generateMinusJob(op)
67     elif op is OpUnion:
68         generateUnionJob(op)
69     elif op is OpModifier:

```

```

68     if op is OpDistinctReduced:
69         generateDistinctJob(op)
70     elif op is OpOrder:
71         generateOrderByJob(op)
72     elif op is OpProject:
73         generateProjectJob(op)
74     elif op is OpSlice:
75         generateSliceJob(op)
76     else
77         throw NotImplementedException()
78 else:
79     throw NotImplementedException()
80
81 function visitUpdate(opUpdateDataInsert)
82     setQueryType("INSERT_DATA")
83     foreach quad = (graph, triple) in opUpdateDataInsert.getQuads(): // we only support a single
84         ↪ default graph, hence quads can be converted to triples
85         uploadtemplate ← uploadtemplate ∪ { triple }
86
87 function visitUpdate(opUpdateDataDelete)
88     setQueryType("DELETE_DATA")
89     foreach quad = (graph, triple) in opUpdateDataDelete.getQuads():
90         deletetemplate ← deletetemplate ∪ { triple }
91
92 function visitUpdate(opUpdateDeleteWhere):
93     setQueryType("MODIFY")
94     bgp ← {}
95     foreach quad = (graph, triple) in opUpdateDeleteWhere.getQuads():
96         bgp ← bgp ∪ {triple}
97         deletetemplate ← deletetemplate ∪ { triple }
98     JENA.OpWalker.walk(JENA.createOpBGP(bgp), this)
99
100 function visitUpdate(opUpdateModify, JENAalgebraGenerator)
101     setQueryType("MODIFY")
102     if opUpdateModify.hasInsertClause():
103         foreach quad = (graph, triple) in opUpdateDataInsert.getInsertQuads():
104             uploadtemplate ← uploadtemplate ∪ { triple }
105     if opUpdateModify.hasDeleteClause():
106         foreach quad = (graph, triple) in opUpdateDeleteWhere.getDeleteQuads():
107             deletetemplate ← deletetemplate ∪ { triple }
108     JENA.OpWalker.walk(JENAalgebraGenerator.compile(opUpdateModify.getWherePattern()), this)
109
110 function generateSearchJobs(opOpBGP):
111     patterns = opOpBGP.getPattern()
112     jobIDs = {}
113     jobs = {}
114     jobVars = {}
115     foreach t = (s,p,o) in patterns:
116         jobID = pushSearch(s, p, o, jobs, jobIDs)
117         job = jobs[jobID]
118         varsjob ← extractVars(job)
119         variables ← variables ∪ varsjob
120         jobVars[jobID] ← varsjob
121     if #patterns > 1:
122         generateRandomJoinPipeline(opOpBGP, jobVars)
123     else:
124         operationsparsed.put(opOpBGP, jobID)
125
126 function pushSearch(s, p, o, jobs, jobIDs):

```

```

126  searchPattern ← parsePattern(s, p, o)
127  jobID = jobIDs.get(searchPattern)
128  if jobID == :bot:
129      jobID ← generateID()
130      jobSEARCH ← createSearchJob(jobID, extractVariablesPattern(s, p, o), s, p, o)
131      jobs[jobID] ← jobSEARCH
132      jobIDs[searchPattern] ← jobID
133      pushJob(jobSEARCH)
134  return jobID
135
136  function generateRandomJoinPipeline(opOpBGP, jobVars)
137      (graph, vars, jobs) ← generateAdjacencyMatrix(jobVars)
138      path ← createArrayList(numJobs)
139      sampled ← {}
140      stop ← false
141      numJobs ← #jobVars.getKeys()
142      while #sampled < numJobs and stop == false:
143          next ← sample(jobs \ sampled)
144          sampled ← sampled ∪ { next }
145          bfsSearch(next, graph, path, numJobs)
146          if #path == numJobs:
147              stop ← true
148      if #path == numJobs:
149          joins = createLinkedList()
150          for i in range(0, #path):
151              jobID1 = path.get(i)
152              jobID2 = path.get(i + 1)
153              joinID = generateID()
154              path.get(i + 1) ← joinID
155              joins.add(createJoinJob(joinID, jobID1, jobID2))
156          foreach jobJOIN in join:
157              pushJob(jobJOIN)
158          operationsparsed[opBGP] ← joins.getLast().getID()
159      else :
160          jobID = generateID()
161          pushJob(createEmptyResultJob(jobID, vars))
162          operationsparsed[opBGP] ← jobID
163
164  function generateAdjacencyMatrix(jobVars)
165      graph ← {}
166      jobs ← {}
167      vars ← {}
168      foreach jobID1 in jobVars.getKeys():
169          jobs ← jobs ∪ { jobID1 }
170          v1 ← jobVars[jobID1]
171          vars ← vars ∪ v1
172          edges ← {}
173          foreach jobID2 in jobVars.getKeys():
174              v2 ← jobVars[jobID2]
175              if jobID1 ≠ jobID2
176                  foreach var \; in \; v1:
177                      if v ∈ v2:
178                          edges ← edges ∪ { JobID2 }
179                      break
180          graph[jobID1] ← edges
181      return (graph, jobs, vars)
182
183
184  function bfsSearch(root, graph, path, depth)
185      queue = createQueue()

```

```

186   path.add(root)
187   queue.push(root)
188   while #path < depth:
189       next ← queue.poll()
190       foreach neighbour in graph[next]:
191           if neighbour ∉ path:
192               path.add(neighbour)
193               queue.push(neighbour)
194       if #queue == 0:
195           return
196
197   function getJobID(op)
198       jobID = operationsparsed[op]
199       if jobID == ⊥:
200           throw QueryBuildException()
201       return jobID
202
203   function generateProjectJob(opOpProject)
204       jobID ← generateID()
205       variables ← op.getVars()
206       plan.vars ← variables
207       pushJob(createProjectJob(jobID, getJobID(op.getSubOp()), variables))
208       operationsparsed[op] ← jobID
209
210   function generateDistinctJob(opOpDistinctReduced)
211       jobID ← generateID()
212       pushJob(createDistinctJob(jobID, getJobID(op.getSubOp())))
213       operationsparsed[op] ← jobID
214
215   function generateOrderByJob(opOpOrder)
216       jobID ← generateID()
217       pushJob(createOrderByJob(jobID, getJobID(op.getSubOp()), op.getConditions()))
218       operationsparsed[op] ← jobID
219
220   function generateSliceJob(opOpSlice)
221       jobID ← generateID()
222       pushJob(createSliceJob(jobID, getJobID(op.getSubOp()), op.getStart(), op.getLength()))
223       operationsparsed[op] ← jobID
224
225   function generateJoinJob(opOpJoin)
226       jobID ← generateID()
227       pushJob(createJoinJob(jobID, getJobID(op.getLeft()), getJobID(op.getRight())))
228       operationsparsed[op] ← jobID
229
230   function generateOptionalJob(opOpLeftJoin)
231       jobID ← generateID()
232       pushJob(new OptionalJob(jobID, getJobID(op.getLeft()), getJobID(op.getRight())))
233       operationsparsed[op] ← jobID
234
235   function generateMinusJob(opOpMinus)
236       jobID ← generateID()
237       pushJob(createMinusJob(jobID, getJobID(op.getLeft()), getJobID(op.getRight())))
238       operationsparsed[op] ← jobID
239
240   function generateUnionJob(opOpUnion)
241       jobID ← generateID()
242       pushJob(createUnionJob(jobID, getJobID(op.getLeft()), getJobID(op.getRight())))
243       operationsparsed[op] ← jobID

```

Listing A.4: SPARQL Query Planner (without reasoning capabilities)

```

1 State:
2   pending ← {}
3   finished ← {}
4   results ← {}
5   current ← ⊥
6   sparqlResult ← ⊥
7
8 function exec(plan = (jobs, execOrder), worker):
9   pending ← execOrder
10  while not isFinished():
11    current = execOrder.peek()
12    results[current] ← delegateJob(worker, jobs[current])
13    finished ← finished ∪ { pending.poll() };
14    sparqlResult ← worker.generateSPARQLResult(results[current]);
15
16 function delegateJob(worker, job):
17   res ← results[job.jobID]
18   if res == ⊥:
19     if job is UNARY:
20       return worker.exec(job, results[job.jobIDprev])
21     elif job is BINARY:
22       return worker.exec(job, results[job.jobIDl], results[job.jobIDr])
23     else:
24       return worker.exec(job);
25   else
26     return res;
27
28 function getFinishedJobs():
29   return finished
30
31 function getPendingJobs():
32   return pending
33
34 function getCurrentJob():
35   return current
36
37 function isFinished():
38   return #pending > 0
39
40 function isFinished(jobID):
41   return results[jobID] ≠ ⊥
42
43 function getSPARQLResult():
44   return sparqlResult

```

Listing A.5: Query Execution

```

1 State:
2   triplestoreID ← ⊥
3   storageEngine ← ⊥
4   sparqlResult ← ({}, false, false, false, ⊥, ⊥)
5
6 function setContext(target, storage):
7   triplestoreID ← target
8   storageEngine ← storage
9
10 function exec(job):
11   if job is SEARCH:
12     return execSearch(job)

```

```
13     elif job is EMPTY:
14         return createBindingsTable(job.vars)
15     throw JobInstanceException()
16
17 function exec(job, prevJobResults):
18     if job is PROJECT:
19         return execProject(job, prevJobResults)
20     elif job is ORDERBY:
21         return execOrderBy(job, prevJobResults)
22     elif job is DISTINCT:
23         return execDistinct(prevJobResults)
24     elif job is SLICE:
25         return execSlice(job, prevJobResults)
26     throw JobInstanceException()
27
28 function exec(job, left, right):
29     if job is JOIN:
30         return execJoin(left, right)
31     elif job is UNION:
32         return execUnion(left, right)
33     elif job is OPTIONAL:
34         return execOptional(left, right)
35     elif job is MINUS:
36         return execMinus(left, right)
37     throw JobInstanceException()
38
39 private execSearch(jobSEARCH = (jobID, varPattern, s, p, o)):
40     if varPattern ≠ "SP0"
41         return storageEngine.search(triplestoreID, varPattern, s, p, o)
42     else
43         throw NotImplementedException()
44
45 function execProject(jobPROJECT = (jobID, vars), prevJobResults):
46     prevJobResults.project(vars)
47     return prevJobResults
48
49 function execOrderBy(jobORDERBY = (jobID, jobIDprev, sortCondititions), prevJobResults):
50     sparqlResult.ordered ← true
51     sparqlResult.sortCondititions ← sortCondititions
52     return prevJobResults
53
54 function execSlice(jobSLICE = (jobID, jobIDprev, offset, length), prevJobResults):
55     sparqlResult.sliced ← true
56     sparqlResult.length ← length
57     sparqlResult.offset ← offset
58     return prevJobResults
59
60 function execDistinct(prevJobResults):
61     sparqlResult.distinct ← true
62     return prevJobResults
63
64 function execJoin(left, right):
65     return left.join(right)
66
67 function execUnion(left, right):
68     return left.union(right)
69
70 function execOptional(left, right):
71     return left.leftOuterJoin(right)
72
```

```

73 function execMinus(left, right):
74     return left.minus(right)
75
76 function generateResults(jobResults):
77     bindings ← ⊥
78     if sparqlResult.ordered == true:
79         comparator ← JENA.BindingComparator.create(sparqlResult.sortConditions)
80         if sparqlResult.distinct == true:
81             bindings ← generateBindings(createTreeSet(comparator), jobResults)
82         else:
83             bindings ← generateBindings(createLinkedList(), jobResults)
84             sort(bindings, comparator)
85     else:
86         if sparqlResult.distinct == true:
87             bindings ← generateBindings(createHashSet(), jobResults)
88         else
89             bindings ← generateBindings(createLinkedList(), jobResults)
90     if sparqlResult.sliced == true:
91         bindings ← slice(bindings, sparqlResult.offset, sparqlResult.length)
92     sparqlResult.bindings ← bindings
93     return sparqlResult.bindings
94
95 function generateBindings(collector, jobResults):
96     vars = jobResults.getVariables()
97     values ← {}
98     foreach patterns in jobResults.getPatterns():
99         binding ← {}
100         i ← 0
101         foreach parsedValue in patterns:
102             if parsedValue ≠ ⊥:
103                 binding[vars[i]] ← storageEngine.generateNode(parsedValue)
104                 i++
105         collector.add(binding)
106     return collector

```

Listing A.6: Query Worker

```

1 State:
2     queryEngine ← createSPARQLQueryEngine()
3     planner ← ⊥
4
5 function executeSPARQLQuery(triplestoreID, query):
6     try:
7         (plan, construct, upload, delete) ← planQuery(query)
8         queryType ← planner.getQueryType()
9         if queryType ∈ { SELECT, CONSTRUCT, ASK, DESCRIBE, MODIFY,
10             ↪ DELETE_WHERE }:
11             sparqlResult ← Triplestore.evaluateQuery(triplestoreID, plan)
12             return generateFinalResult(triplestoreID, queryType, plan.vars, sparqlResult.bindings,
13                 ↪ construct, upload, delete)
14         catch Exception:
15             return ⊥
16
17 // {(plan, construct, upload, delete), ⊥} ← SPARQL-EVAL.planQuery(query)
18 function planQuery(query):
19     planner ← createSPARQLQueryPlanner()
20     return (queryEngine.getQueryPlan(planner, query), planner.getConstructTemplate(),
21         ↪ planner.getUploadTemplate(), planner.getDeleteTemplate())

```

```

20 // {sparqlFinalResult,  $\perp$ }  $\leftarrow$  SPARQL-EVAL.generateFinalResult(triplestoreID, variables,
     $\hookrightarrow$  queryType, bindings, construct, upload, delete):
21 function generateFinalResult(triplestoreID, queryType, variables, sparqlResult,
     $\hookrightarrow$  constructtemplate, uploadtemplate, deletetemplate)
22   if queryType == "SELECT":
23     return generateSELECTResults(variables, bindings)
24   elif queryType == "CONSTRUCT":
25     return generateCONSTRUCTResults(constructtemplate, bindings)
26   elif queryType == "ASK":
27     return generateASKResults(bindings)
28   elif queryType == "DESCRIBE":
29     return generateDESCRIBEResults(variables, bindings)
30   else:
31     return applyModifications(triplestoreID, bindings, uploadtemplate, deletetemplate)
32
33 function applyModifications(triplestoreID, bindings, uploadtemplate, deletetemplate):
34   triplesupload  $\leftarrow$  {}
35   triplesdelete  $\leftarrow$  {}
36   if queryType == "MODIFY":
37     triplesupload  $\leftarrow$  generateTriplesFromBindings(uploadtemplate, bindings) //foreach
     $\hookrightarrow$  binding, map values to variables and output the set of generated triples
38   triplesdelete  $\leftarrow$  generateTriplesFromBindings(deletetemplate, bindings) //foreach binding,
     $\hookrightarrow$  map values to variables and output the set of generated triples
39   else if (queryType == INSERT_DATA)
40     triplesupload  $\leftarrow$  uploadtemplate
41   else if (queryType == DELETE_DATA)
42     triplesdelete  $\leftarrow$  deletetemplate
43   try:
44     Triplestore.deleteTriples(triplestoreID, triplesdelete)
45     Triplestore.saveTriples(triplestoreID, triplesupload)
46   catch Exception:
47     return  $\perp$ 
48   return "SUCCESS"
49
50 function generateASKResults(bindings):
51   return #bindings > 0
52
53 function generateCONSTRUCTResults(constructtemplate, bindings):
54   return generateGraphFromBindings(constructtemplate, bindings) //foreach binding, map
     $\hookrightarrow$  values to variables and output a graph with all the generated triples
55
56 function generateSELECTResults(variables, bindings):
57   return generateJSONQueryResult(variables, bindings) // map bindings collection to a
     $\hookrightarrow$  well-formed JSON, as per the official specification
58
59 private Response generateDESCRIBEResults(variables, bindings) :
60   frequencies  $\leftarrow$  {}
61   foreach var in variables:
62     frequencies[var]  $\leftarrow$  0
63   foreach binding in bindings:
64     foreach (var, node) in binding.getKeys():
65       frequencies[var]  $\leftarrow$  frequencies[var] + 1
66   return generateJSONFrequencyTable(frequencies) // format the collection of variable
     $\hookrightarrow$  frequencies as JSON

```

Listing A.7: SPARQL-EVAL client

```

1 State:
2   storageEngine  $\leftarrow$  createStorageEngine()

```

```

3
4 // {sparqlResult, ⊥} ← SPARQL-EVAL.evaluateQuery(triplestoreID, plan):
5 function evaluateQuery(triplestoreID, plan):
6     worker ← createQueryWorker()
7     worker.setContext(triplestoreID, storageEngine)
8     execution ← createQueryExecution()
9     execution.exec(plan, worker)
10    return execution.getSPARQLResult()
11
12 function saveTriples(triplestoreID, triples):
13     return storageEngine.save(triplestoreID, triples)
14
15 function deleteTriples(triplestoreID, triples):
16     return storageEngine.delete(triplestoreID, triples)
17
18 function delete(triplestoreID):
19     return storageEngine.delete(triplestoreID)

```

Listing A.8: SPARQL-EVAL triplestore

```

1 State:
2     keywordschema ← generateRandomString()
3
4 function save(triplestoreID, triples, schema):
5     if schema == true:
6         delete(triplestoreID, schema)
7         foreach t in triples:
8             saveSchemaTriple(t)
9     else:
10        foreach t in triples:
11            saveTriple(t)
12    return "SUCCESS"
13
14 function delete(triplestoreID, schema):
15     if schema == true:
16         triplestores ← triplestores / { triplestores[triplestoreID][schema] }
17     else:
18         triplestores ← triplestores / { triplestores[triplestoreID] }
19    return "SUCCESS"
20
21 function saveSchemaTriple(triple = (s, p, o)):
22     parseds ← parseNode(s)
23     parsedp ← parseNode(p)
24     parsedo ← parseNode(o)
25     triple ← parseds || separatorIII || parsedp || separatorIII || parsedo
26     triplestores[triplestoreID][keywordschema] ← triplestores[triplestoreID][keywordschema] ∪ {
27         ↪ triple }
28
29 function searchSchema(triplestoreID):
30     triplesschema ← {}
31     foreach t in triplestores[triplestoreID][keywordschema]:
32         (parseds, parsedp, parsedo) ← t.split(separatorIII)
33         triplesschema ← triplesschema ∪ { (generateNode(parseds), generateNode(parsedp),
34             ↪ generateNode(parsedo)) }
35    return triplesschema

```

Listing A.9: Storage of schema triples

```

1 State:
2   storageEngine ← createStorageEngine()
3
4 function saveTriples(triplestoreID, triples, schema):
5   return storageEngine.save(triplestoreID, triples)
6
7 function delete(triplestoreID, schema):
8   return storageEngine.delete(triplestoreID, schema)
9
10 function findSchema(triplestoreID)
11   return storageEngine.searchSchema(triplestoreID)

```

Listing A.10: SPARQL-EVAL-II triplestore (extra operations)

```

1 State:
2   subClasses ← {}
3   equivalentClasses ← {}
4   classRestrictions ← {}
5   intersectionClasses ← {}
6   intersectionsWhereClassIsOperand ← {}
7   subProperties ← {}
8   equivalentProperties ← {}
9   inverseProperties ← {}
10  symmetricProperties ← {}
11  transitiveProperties ← {}
12  transitivityDepth ← ⊥
13  expansionDepth ← ⊥
14
15 function execInference(schema, e, t):
16   expansionDepth ← e
17   transitivityDepth ← t
18   tbox = JENAModelFactory.createOntologyModel(JENAspec.OWL_MEM)
19   JENAGraphUtil.add(tbox.getGraph(), schema)
20   model = JENAModelFactory.createOntologyModel(JENAspec.OWL_MEM_TRANS_INF, tbox)
21   execClassInference(model.listClasses())
22   execPropertyInference(model.listAllOntProperties())
23
24 function execClassInference(classes, restrictions, intersections):
25   foreach class = (subclasses, equivalentClasses) in classes:
26     subClasses[convertToNode(class)] ← subclasses
27     equivalentClasses[convertToNode(class)] ← equivalentClasses
28   foreach restriction = (type, property, value) in restrictions:
29     type ∈ {"owl:hasValue", "owl:isSomeValuesFrom", "owl:allValuesFrom"}:
30       classRestrictions[convertToNode(restriction)] ← restriction
31   foreach intersection = (id, operands) in intersections: // Intersections are anonymous
32     ↳ classes with a set of other classes (operands)
33     foreach operand in operands:
34       intersectionsDirectSuperclasses ← intersectionsWhereClassIsOperand[convertToNode(
35         ↳ operand)]
36       if intersectionsDirectSuperclasses == ⊥:
37         intersectionsDirectSuperclasses ← {}
38       intersectionsDirectSuperclasses ← intersectionsDirectSuperclasses ∪ intersection
39       intersectionsWhereClassIsOperand[convertToNode(operand)] ←
40         ↳ intersectionsDirectSuperclasses
41       intersectionClasses[convertToNode(intersection)] ← operands
42
43 function execPropertyInference(properties):
44   foreach p = (type, value, subproperties, equivalentProperties, inverseProperties) in properties:

```

```

42     if type == "owl:SymmetricProperty":
43         symmetricProperties ← symmetricProperties ∪ { p }
44     elif type == "owl:TransitiveProperty":
45         transitiveProperties ← transitiveProperties ∪ { p }
46         subProperties[convertToNode(p)] ← subproperties
47         equivalentProperties[convertToNode(p)] ← equivalentProperties
48         inverseProperties[convertToNode(p)] ← inverseProperties
49
50 function getEquivalentClasses(node):
51     classes ← equivalentClasses[node]
52     return classes == ⊥:
53         return {}
54     return classes
55
56 function getSubClasses(node):
57     classes ← subClasses[node]
58     return classes == ⊥:
59         return {}
60     return classes
61
62 function getRestriction(node):
63     return classRestrictions[node]
64
65 function getIntersection(node):
66     classes ← intersectionClasses[node]
67     return classes == ⊥:
68         return {}
69     return classes
70
71 function getIntersectionWhereClassIsOperand(node):
72     classes ← intersectionsWhereClassIsOperand[node]
73     return classes == ⊥:
74         return {}
75     return classes
76
77 function getSubProperties(node):
78     properties ← subProperties[node]
79     return properties == ⊥:
80         return {}
81     return properties
82
83 function getEquivalentProperties(node):
84     properties ← equivalentProperties[node]
85     if properties == ⊥:
86         return {}
87     return properties
88
89 function getInverseOf(node):
90     properties ← inverseProperties[node]
91     if properties == ⊥:
92         return {}
93     return properties
94
95 function isSymmetric(node):
96     return node ∈ symmetricProperties
97
98 function isTransitive(node):
99     return node ∈ transitiveProperties
100
101 function getTransitivityDepth():

```

```

102     return transitivityDepth
103
104 function getMaximumExpansionDepth():
105     return expansionDepth

```

Listing A.11: Ontology

```

1  State:
2      operationsparsed ← {}
3      uploadtemplate ← {}
4      deletetemplate ← {}
5      constructtemplate ← {}
6      queryType ← :bot
7      plan ← (jobs = {}, execOrder = createLinkedList(), vars = {})
8      variables ← {}
9      ontology ← ⊥
10
11 function setOntology(ontologynew):
12     ontology ← ontologynew
13
14 function extractVarsOfBinaryJob(jobs, job = (jobID, jobIDl, jobIDr)):
15     vars ← {}
16     jobsqueue ← createQueue()
17     jobsqueue.push(jobs[jobIDl])
18     jobsqueue.push(jobs[jobIDr])
19     while #jobsqueue > 0:
20         jobnext ← jobsqueue.poll()
21         if jobnext is SEARCH:
22             vars ← vars ∪ extractVars(searchJob)
23         elif jobnext = (jobIDnext, l, r) is BINARY:
24             jobsqueue.push(jobs[l])
25             jobsqueue.push(jobs[r])
26     return vars
27
28 function pushUnion(left, right, jobs, jobIDs):
29     unionPattern ← "union" || left || right
30     jobID ← jobIDs[unionPattern]
31     if jobID == ⊥:
32         jobMINUS ← createMinusJob(generateID(), right, left)
33         jobID ← generateID()
34         jobUNION ← createUnionJob(jobID, left, jobMINUS.jobID)
35         jobs[jobID] ← jobUNION
36         jobIDs[jobID] ← unionPattern
37         pushJob(jobMINUS)
38         pushJob(jobUNION)
39     return jobID
40
41 private String pushJoin(left, right, jobs, jobIDs):
42     joinPattern ← "join" || left || right
43     jobID ← jobIDs[joinPattern]
44     if jobID == ⊥:
45         jobID = generateID()
46         jobJOIN = createJoinJob(jobID, left, right)
47         jobs[jobID] ← jobJOIN
48         jobIDs[joinPattern] ← jobID
49         pushJob(jobJOIN)
50     return jobID
51
52 function generateSearchJobs(opOpBGP):

```

```

53  patterns = opOpBGP.getPattern()
54  jobIDs = {}
55  jobs = {}
56  jobVars = {}
57  foreach t=(s,p,o) in patterns:
58      jobID = pushSearch(s, p, o, jobs, jobIDs)
59      if ontology ≠ ⊥:
60          if (p == "rdf:type") and (o is not VARIABLE):
61              jobID = expandClass(jobID, s, o, 0, jobs, jobIDs)
62          elif p is not VARIABLE:
63              jobID = expandProperty(jobID, s, p, o, 0, jobs, jobIDs, true, true)
64      job ← jobs[jobID]
65      if job is SEARCH:
66          varsjob ← extractVars(job)
67          variables ← variables ∪ varsjob
68          jobVars[jobID] ← varsjob
69      elif job is BINARY:
70          jobVars[jobID] ← extractVarsOfBinaryJob(jobs, job)
71  if #patterns > 1:
72      generateRandomJoinPipeline(opOpBGP, jobVars)
73  else:
74      operationsparsed.put(opOpBGP, jobID)
75
76  function expandClass(jobID, s, o, depth, jobs, jobsIDs):
77      if depth == ontology.getMaximumExpansionDepth():
78          return jobID
79      subclasses ← ontology.getSubClasses(o)
80      equivalentClasses ← ontology.getEquivalentClasses(o)
81      restriction ← ontology.getRestriction(o)
82      intersectionClasses ← ontology.getIntersectionWhereClassIsOperand(o)
83      jobID ← expandClassDisjunction(subclasses, jobID, s, depth, jobs, jobsIDs)
84      jobID ← expandClassDisjunction(equivalentClasses, jobID, s, depth, jobs, jobsIDs)
85      jobID ← expandRestriction(restriction, jobID, s, depth, jobs, jobsIDs)
86      jobID ← expandClassDisjunction(intersectionClasses, jobID, s, depth, jobs, jobsIDs)
87      intersection ← ontology.getIntersection(o)
88      if #intersection == 0:
89          jobID ← expandClassConjunction(intersection, jobID, s, depth, jobs, jobsIDs)
90      return jobID
91
92  function expandRestriction(restriction=(type,property,value), jobID, s, depth, jobs, jobsIDs):
93      if restriction ≠ ⊥:
94          if type == "owl:hasValue":
95              var ← createVariable(parse(property))
96              left ← pushSearch(var, "rdf:type", value, jobs, jobsIDs)
97              right ← expandProperty(pushSearch(s, property, value, jobs, jobsIDs), s, property,
98                  ↪ , value, depth + 1, jobs, jobsIDs, true, true)
99              jobID ← pushUnion(right, left, jobs, jobsIDs)
100          elif type ∈ {"owl:isSomeValuesFrom", "owl:allValuesFrom"}:
101              var ← createVariable(parse(value))
102              left ← expandClass(pushSearch(var, "rdf:type", value, jobs, jobsIDs), var, value,
103                  ↪ depth + 1, jobs, jobsIDs)
104              right ← expandProperty(pushSearch(s, property, var, jobs, jobsIDs), s, property,
105                  ↪ var, depth + 1, jobs, jobsIDs, true, true)
106              jobID ← pushUnion(jobID, pushJoin(right, left, jobs, jobsIDs), jobs, jobsIDs)
107      return jobID
108
109  function expandClassConjunction(classes, jobID, s, depth, jobs, jobsIDs):
110      foreach c in classes:
111          search ← pushSearch(s, "rdf:type", c, jobs, jobsIDs)
112          current = expandClass(search, s, c, depth + 1, jobs, jobsIDs)

```

```

110     if next  $\neq$   $\perp$ :
111         next  $\leftarrow$  pushJoin(next, current, jobs, jobsIDs)
112     else:
113         next  $\leftarrow$  current
114     if next ==  $\perp$ :
115         return pushUnion(jobID, next, jobs, jobsIDs)
116     return jobID
117
118 function expandClassDisjunction(classes, jobID, s, depth, jobs, jobsIDs):
119     foreach c in classes:
120         search  $\leftarrow$  pushSearch(s, "rdf:type", c, jobs, jobsIDs)
121         current = expandClass(search, s, c, depth + 1, jobs, jobsIDs)
122         if next  $\neq$   $\perp$ :
123             next  $\leftarrow$  pushUnion(next, current, jobs, jobsIDs)
124         else:
125             next  $\leftarrow$  current
126     if next ==  $\perp$ :
127         return pushUnion(jobID, next, jobs, jobsIDs)
128     return jobID
129
130 function expandProperty(jobID, s, p, o, depth, jobs, jobsIDs, symmetric, inverseOf):
131     if depth == ontology.getMaximumExpansionDepth():
132         return jobID
133     if symmetric == true and ontology.isSymmetric(p):
134         jobID  $\leftarrow$  pushUnion(jobID, pushSearch(o, p, s, jobs, jobsIDs), jobs, jobsIDs)
135         jobID  $\leftarrow$  expandProperty(jobID, o, p, s, depth + 1, jobs, jobsIDs, false, inverseOf)
136     if ontology.isTransitive(p):
137         var  $\leftarrow$  createVariable(parse(p) || 0)
138         left  $\leftarrow$  pushSearch(s, p, var, jobs, jobsIDs)
139         join  $\leftarrow$  pushJoin(left, pushSearch(var, p, o, jobs, jobsIDs), jobs, jobsIDs)
140         jobID  $\leftarrow$  pushUnion(jobID, join, jobs, jobsIDs)
141         for i in range(0, ontology.getTransitivityDepth())
142             for j in range(0, i + 1):
143                 if j == i:
144                     varnext = createVariable(parse(p) || j)
145                     left  $\leftarrow$  pushJoin(left, pushSearch(var, p, varnext, jobs, jobsIDs), jobs,
146                                      $\hookrightarrow$  jobsIDs)
147                     var = varnext
148                     join  $\leftarrow$  pushJoin(left, pushSearch(var, p, o, jobs, jobsIDs), jobs, jobsIDs)
149                     jobID  $\leftarrow$  pushUnion(jobID, join, jobs, jobsIDs)
150     subproperties  $\leftarrow$  ontology.getSubProperties(p)
151     equivalentProperties  $\leftarrow$  ontology.getEquivalentProperties(p)
152     jobID  $\leftarrow$  expandProperties(subproperties, jobID, s, o, depth, jobs, jobsIDs, symmetric,
153                              $\hookrightarrow$  inverseOf)
154     jobID  $\leftarrow$  expandProperties(equivalentProperties, jobID, s, o, depth, jobs, jobsIDs,
155                              $\hookrightarrow$  symmetric, inverseOf)
156     if inverseOf == true:
157         inverseProperty  $\leftarrow$  ontology.getInverseOf(p)
158         jobID  $\leftarrow$  expandProperties(inverseProperty, jobID, o, s, depth, jobs, jobsIDs,
159                              $\hookrightarrow$  symmetric, false)
160     return jobID
161
162 function expandProperties(properties, jobID, s, o, depth, jobs, jobsIDs, symmetric, inverseOf
163                              $\hookrightarrow$  ):
164     for p in properties:
165         jobID  $\leftarrow$  pushUnion(jobID, pushSearch(s, p, o, jobs, jobsIDs), jobs, jobsIDs)
166         jobID  $\leftarrow$  expandProperty(jobID, s, p, o, depth + 1, jobs, jobsIDs, symmetric,
167                              $\hookrightarrow$  inverseOf)
168     return jobID

```

Listing A.12: Query Plan Generation with Query Rewriting

```

1 State:
2   queryEngine ← createSPARQLQueryEngine()
3   planner ← ⊥
4
5 function executeSPARQLQuery(triplestoreID, query, inference, expansionDepth, transitivityDepth):
6   try:
7     (plan, construct, upload, delete) ← planQuery(query, triplestoreID, inference,
8       ↪ expansionDepth, transitivityDepth)
9     queryType ← planner.getQueryType()
10    if queryType ∈ { SELECT, CONSTRUCT, ASK, DESCRIBE, MODIFY,
11      ↪ DELETE_WHERE }:
12      sparqlResult ← Triplestore.evaluateQuery(triplestoreID, plan)
13      return generateFinalResult(triplestoreID, queryType, plan.vars, sparqlResult.bindings,
14        ↪ construct, upload, delete)
15    catch Exception:
16      return ⊥
17
18 // {(plan, construct, upload, delete), ⊥} ← SPARQL-EVAL-II.planQuery(query, triplestoreID,
19   ↪ inference, expansionDepth, transitivityDepth)
20 function planQuery(query, triplestoreID, inference, expansionDepth, transitivityDepth):
21   planner ← createSPARQLQueryPlanner()
22   if inference == true:
23     schema ← Triplestore.fetchSchema(triplestoreID)
24     ontology ← createOntology()
25     ontology ← execInference(schema, inference, expansionDepth, transitivityDepth)
26     planner.setOntology(ontology)
27   return (queryEngine.getQueryPlan(planner, query), planner.getConstructTemplate(),
28     ↪ planner.getUploadTemplate(), planner.getDeleteTemplate())

```

Listing A.13: SPARQL-EVAL-II client (modified operations)

SSE-TP LISTINGS

```

1
2 // (KMASTER, KRND, ivDET, keywordschema) ← SSE-TP.Setup()
3 State:
4   KMASTER ← generateKey()
5   KRND ← generateKey()
6   ivDET ← generateRandomIV()
7   keywordschema ← generateRandomString()
8   ivZERO ← generateZeroFilledIV()
9   encryptedNodes ← {}
10  keywordFrequencies ← {}
11  derivedKeys ← {}
12
13 // E(T) ← SSE-TP.ReInitialize(KMASTER, KRND, keywordschema)
14 function ReInitialize(master, rnd, iv, schema):
15   KMASTER ← master
16   KRND ← rnd
17   ivZERO ← iv
18   keywordschema ← schema
19   ivZERO ← newZeroFilledIV()
20   encryptedNodes ← {}
21   keywordFrequencies ← {}
22   derivedKeys ← {}
23
24 // E(T) ← SSE-TP.Encrypt([t0, t1, ..., tn], schema)
25 function encrypt(triples, schema):
26   if (schema == true)
27     encryptSchemaTriples(triples)
28   else
29     encryptTriples(triples)
30   return encryptContextualData()
31
32 function EncryptSchemaTriples(triples):
33   foreach t = (s, p, o) in triples:
34     encryptSchemaNode(parseNode(s))
35     encryptSchemaNode(parseNode(p))
36     encryptSchemaNode(parseNode(o))
37
38 function encryptSchemaNode(node):
39   freq ← incrementKeywordFrequency(keywordschema)
40   trapdoor ← E(deriveKey(keywordschema), keywordschema, freq)
41   encryptedNodes[trapdoor] ← E(KRND, node)
42

```

```

43
44 //  $E(T) \leftarrow \text{SSE-TP.Decrypt}(E(\text{node}), \text{schema})$ 
45 function decrypt( $ct, \text{schema}$ ):
46     if ( $\text{schema} == \text{true}$ )
47         return  $D(K_{RND}, ct)$ 
48     else
49         return decryptNode( $ct$ )
50
51 function decryptContextualData( $ct$ ):
52     return  $D(K_{RND}, ct)$ 
53
54 function encryptTriples( $\text{triples}$ ):
55     foreach  $t = (s, p, o)$  in  $\text{triples}$ :
56          $\text{freq} \leftarrow \{\}$ 
57          $\text{parsed}_s \leftarrow \text{parseNode}(s)$ 
58          $\text{parsed}_p \leftarrow \text{parseNode}(p)$ 
59          $\text{parsed}_o \leftarrow \text{parseNode}(o)$ 
60          $\text{keyword}_s \leftarrow \text{parseKeyword}(s)$ 
61          $\text{keyword}_p \leftarrow \text{parseKeyword}(p)$ 
62          $\text{keyword}_o \leftarrow \text{parseKeyword}(o)$ 
63          $\text{freq} \leftarrow \text{freq} \cup \{ \text{encryptNode}(\text{parsed}_p, "P0" \parallel \text{keyword}_s) \}$ 
64          $\text{freq} \leftarrow \text{freq} \cup \{ \text{encryptNode}(\text{parsed}_o, "P0" \parallel \text{keyword}_s) \}$ 
65          $\text{freq} \leftarrow \text{freq} \cup \{ \text{encryptNode}(\text{parsed}_s, "S0" \parallel \text{keyword}_p) \}$ 
66          $\text{freq} \leftarrow \text{freq} \cup \{ \text{encryptNode}(\text{parsed}_o, "S0" \parallel \text{keyword}_p) \}$ 
67          $\text{freq} \leftarrow \text{freq} \cup \{ \text{encryptNode}(\text{parsed}_s, "SP" \parallel \text{keyword}_o) \}$ 
68          $\text{freq} \leftarrow \text{freq} \cup \{ \text{encryptNode}(\text{parsed}_p, "SP" \parallel \text{keyword}_o) \}$ 
69          $\text{freq} \leftarrow \text{freq} \cup \{ \text{encryptNode}(\text{parsed}_s, "S" \parallel \text{keyword}_p \parallel \text{keyword}_o) \}$ 
70          $\text{freq} \leftarrow \text{freq} \cup \{ \text{encryptNode}(\text{parsed}_p, "P" \parallel \text{keyword}_s \parallel \text{keyword}_o) \}$ 
71          $\text{freq} \leftarrow \text{freq} \cup \{ \text{encryptNode}(\text{parsed}_o, "O" \parallel \text{keyword}_s \parallel \text{keyword}_p) \}$ 
72          $\text{encryptTriple}("SP0" \parallel \text{keyword}_s \parallel \text{keyword}_p \parallel \text{keyword}_o, \text{freq})$ 
73
74 function encryptTriplePattern( $\text{keyword}, \text{frequencies}$ ):
75      $i \leftarrow 0$ 
76     foreach  $\text{freq}$  in  $\text{frequencies}$ :
77          $\text{trapdoor} \leftarrow E(\text{deriveKey}(\text{keyword}), \text{keyword}, i)$ 
78          $\text{encryptedNodes}[\text{trapdoor}] \leftarrow E(K_{RND}, \text{freq})$ 
79          $i++$ 
80
81 function deriveKey( $\text{context}$ ):
82      $\text{key} \leftarrow \text{derivedKeys}[\text{context}]$ 
83     if  $\text{key} == \perp$ :
84          $\text{key} \leftarrow \text{generateKey}(K_{MASTER}, \text{context})$ 
85          $\text{derivedKeys}[\text{context}] \leftarrow \text{key}$ 
86     return  $\text{key}$ 
87
88 function incrementKeywordFrequency( $\text{keyword}$ ):
89      $\text{freq} \leftarrow \text{keywordFrequencies}[\text{keyword}]$ 
90     if  $\text{freq} == \perp$ :
91          $\text{freq} \leftarrow 1$ 
92          $\text{keywordFrequencies}[\text{keyword}] \leftarrow \text{freq}$ 
93     else:
94          $\text{keywordFrequencies}[\text{keyword}] \leftarrow \text{freq} + 1$ 
95     return  $\text{freq}$ 
96
97 //  $\text{trapdoor}_{n,i} \leftarrow \text{SSE-TP.Trapdoor}(\text{keyword}_n, i)$ 
98 function generateTrapdoor( $\text{keyword}, i$ ):
99     return  $E(\text{deriveKey}(\text{keyword}), \text{keyword}, i)$ 
100
101 function generateTrapdoor( $\text{keyword}$ ):
102     return  $E(\text{deriveKey}(\text{keyword}), \text{keyword}, \text{ivZERO})$ 

```

```

103
104 function generateTrapdoorAndIncrementIV(keyword):
105     return E(deriveKey(keyword), keyword, incrementKeywordFrequency(keyword))
106
107 function generateTriplePatternTrapdoors(triples):
108     keywordPatternTrapdoors ← {}
109     foreach t=(s,p,o) in triples:
110         keywords ← {}
111         i ← 0
112         keywords = parseKeyword(s)
113         keywordp = parseKeyword(p)
114         keywordo = parseKeyword(o)
115         keywordt = "SP0" || keywords || keywordp || keywordo
116         keywords ← keywords ∪ { "P0" || keywords }
117         keywords ← keywords ∪ { "P0" || keywords }
118         keywords ← keywords ∪ { "S0" || keywordp }
119         keywords ← keywords ∪ { "S0" || keywordp }
120         keywords ← keywords ∪ { "SP" || keywordo }
121         keywords ← keywords ∪ { "SP" || keywordo }
122         keywords ← keywords ∪ { "S" || keywordp || keywordo }
123         keywords ← keywords ∪ { "P" || keywords || keywordo }
124         keywords ← keywords ∪ { "O" || keywords || keywordp }
125         foreach keyword in keywords:
126             trapdoors ← keywordPatternTrapdoors[keyword]
127             if (trapdoors == ⊥):
128                 trapdoors ← {}
129             trapdoors ← trapdoors ∪ { E(deriveKey(keywordt), keywordt, i) }
130             keywordPatternTrapdoors[keyword] ← trapdoors
131             i++
132     return keywordPatternTrapdoors
133
134 function setKeywordFrequencies(values):
135     foreach keyword in values:
136         freq ← values[keyword]
137         if (freq > 0)
138             keywordFrequencies[keyword] ← freq
139
140 function getSchemaKeyword():
141     return keywordschema
142
143 function getMasterKey():
144     return KMASTER
145
146 function getRNDKey():
147     return KRND
148
149 function getEncryptedNodes():
150     return encryptedNodes
151
152 function clearNodes():
153     encryptedNodes ← {}
154
155 function clearFrequencies():
156     keywordFrequencies ← {}

```

Listing B.1: Base logic for the encryption scheme

```

1 State:
2 KDET ← generateKey()

```

```

3
4 function ReInitialize(master, rnd, det, iv, schema):
5     loadState(master, rnd, iv, schema)
6      $K_{DET} \leftarrow det$ 
7
8 function getSecrets():
9     return ( $K_{MASTER}$ ,  $K_{RND}$ ,  $K_{DET}$ ,  $iv_{DET}$ ,  $keyword_{schema}$ )
10
11 function encryptContextualData():
12     foreach (keyword, frequency) in keywordFrequencies:
13         trapdoor  $\leftarrow E(\text{derivedKeys}[keyword], keyword, iv_{ZERO})$ 
14         encryptedNodes[trapdoor]  $\leftarrow E(K_{RND}, frequency)$ 
15     return encryptedNodes
16
17 function encryptNode(node, keyword):
18     freq  $\leftarrow \text{incrementKeywordFrequency}(keyword)$ 
19     trapdoor  $\leftarrow E(\text{deriveKey}(keyword), keyword, freq)$ 
20     encryptedNodes[trapdoor]  $\leftarrow E(K_{RND}, E(K_{DET}, node, iv_{DET}))$ 
21     return freq
22
23 function decryptNode(node):
24     return  $D(K_{RND}, D(K_{DET}, node, iv_{DET}))$ 
25
26 function decryptBinding(ct):
27     return  $D(K_{DET}, ct)$ 
28
29 function getDETKey():
30     return  $K_{DET}$ 
31
32 function getIvDET():
33     return  $iv_{DET}$ 

```

Listing B.2: Triple Pattern Encryption for SEnTriS V1

```

1 State:
2     ( $K_{DGK_{priv}}$ ,  $K_{DGK_{pub}}$ )  $\leftarrow \text{generateDGKKeyPair}()$ 
3     eqTags  $\leftarrow \{\}$ 
4     eqIlast  $\leftarrow 0$ 
5
6 function ReInitialize(master, rnd, KeyPairDGK, iv, schema):
7     loadState(master, rnd, iv, schema)
8     ( $K_{DGK_{priv}}$ ,  $K_{DGK_{pub}}$ )  $\leftarrow \text{KeyPair}_{DGK}$ 
9
10 function getSecrets():
11     return ( $K_{MASTER}$ ,  $K_{RND}$ ,  $K_{DGK_{priv}}$ ,  $K_{DGK_{pub}}$ ,  $iv_{DET}$ ,  $keyword_{schema}$ )
12
13 function encryptContextualData():
14     foreach (keyword, frequency) in keywordFrequencies:
15         trapdoor  $\leftarrow E(\text{derivedKeys}[keyword], keyword, iv_{ZERO})$ 
16         encryptedNodes[trapdoor]  $\leftarrow E(K_{RND}, \text{keywordFrequencies}[keyword])$ 
17     foreach (node, eqTag) in eqTags:
18         trapdoor  $\leftarrow E(\text{derivedKeys}[node], node, iv_{ZERO})$ 
19         encryptedNodes[trapdoor]  $\leftarrow E(K_{RND}, eqTag)$ 
20     return encryptedNodes
21
22 function encryptNode(node, keyword):
23     freq  $\leftarrow \text{incrementKeywordFrequency}(keyword)$ 
24     trapdoor  $\leftarrow E(\text{deriveKey}(keyword), keyword, freq)$ 
25      $enc_{eq_{tag}} \leftarrow E(K_{DGK_{pub}}, \text{getEqTag}(node))$ 

```

```

26   encryptedNodes[trapdoor] ← enceqtag
27   encryptedNodes[enceqtag] ← E(KRND, node)
28   return freq
29
30 function decryptNode(node):
31   return D(KRND, node)
32
33 function getEqTag(node):
34   eqtag ← eqTags[node]
35   if eqtag == ⊥:
36     eqtag ← eqlast + 1
37     eqTags[node] ← eqtag
38   if (eqtag > eqlast)
39     eqlast ← eqtag
40   return eqtag;
41
42 function setEqTags(value):
43   eqTags ← value
44
45 function getDETKey():
46   return KDET
47
48 function getIvDET():
49   return ivDET
50
51 function getLastEqTag():
52   return eqlast
53
54 function getEqKey():
55   return (KDGKpriv.p, KDGKpriv.vp, KDGKpub.n, KDGKpub.u)
56
57 function getPubDGK()
58   return (KDGKpub.n, KDGKpub.g, KDGKpub.h, KDGKpub.u, KDGKpub.k, KDGKpub.t, KDGKpub.l)
59
60 function clearEqTags():
61   eqTags ← {}

```

Listing B.3: Triple Pattern Encryption for SEnTriS V2

```

1  State:
2    TIMEOUT ← fetchConfig(TIMEOUT)
3    triplestores ← {}
4    uploads ← {}
5    deletions ← {}
6
7  function delete(triplestoreID):
8    triplestores ← triplestores / { triplestores[triplestoreID] }
9    return "SUCCESS"
10
11 function commitUpload(triplestoreID, encryptedNodes):
12   uploadID ← generateID()
13   uploads[uploadID] ← encryptedNodes
14   schedule(expireUpload(uploadID), TIMEOUT) // schedules function to execute after a
15     ↳ period specified by TIMEOUT
16   return uploadID
17
18 function commitDelete(triplestoreID, trapdoors):
19   deletionID ← generateID()

```

```

19  deletions[deletionID] ← trapdoors
20  schedule(expireDeletion(deletionID), TIMEOUT)
21  return deletionID
22
23  function update(triplestoreID, deletionIDs, uploadIDs):
24    foreach deletionID in deletionIDs:
25      trapdoors ← deletions[deletionID]
26      if trapdoors == ⊥:
27        return ⊥
28      foreach trapdoor in trapdoors:
29        triplestores[triplestoreID][trapdoor] ← ⊥
30    foreach uploadID in uploadIDs:
31      encryptedNodes ← uploads[uploadID]
32      if encryptedNodes == ⊥:
33        return ⊥
34      foreach (trapdoor, encryptedNode) in encryptedNodes:
35        triplestores[triplestoreID][trapdoor] ← encryptedNode
36    return "SUCCESS"
37
38  function search(triplestoreID, trapdoors):
39    result ← createArrayList(#trapdoors)
40    for i in range(0, #trapdoor):
41      result.add(triplestores[triplestoreID][trapdoors[i]])
42    return result
43
44  function expireUpload(uploadID):
45    if uploads[uploadID] ≠ ⊥:
46      uploads[uploadID] ← ⊥
47
48  function expireDeletion(deletionID):
49    if deletions[deletionID] ≠ ⊥:
50      deletions[deletionID] ← ⊥

```

Listing B.4: SSE-TP Storage Engine

```

1  State:
2    storageEngine ← createStorageEngine()
3
4  // [E(node0),...,E(noden)] ← SSE-TP.Search(triplestoreID, [trapdoor0,...,trapdoorn])
5  function search(triplestoreID, trapdoors):
6    return storageEngine.search(triplestoreID, trapdoors)
7
8  // {uploadID, ⊥} ← SSE-TP.PrepareUpload(triplestoreID, E(nodes))
9  function prepareUpload(triplestoreID, encryptedNodes):
10   return storageEngine.commitUpload(triplestoreID, encryptedNodes)
11
12  // {deleteID, ⊥} ← SSE-TP.PrepareDelete(triplestoreID, [trapdoor0,...,trapdoorn])
13  function prepareDelete(triplestoreID, trapdoors):
14   return storageEngine.commitDelete(triplestoreID, trapdoors)
15
16  // {SUCCESS, ⊥} ← SSE-TP.Update(triplestoreID, [updateID0,...,updateIDn],
17   ↪ [deletionID0,...,deletionIDn])
18  function update(triplestoreID, uploadIDs, deletionIDs):
19   return storageEngine.update(triplestoreID, uploadIDs, deletionIDs)
20
21  // {SUCCESS, ⊥} ← SSE-TP.DeleteAll(triplestoreID):
22  function deleteAll(triplestoreID)
23   return storageEngine.delete(triplestoreID)

```

Listing B.5: SSE-TP Triplestore Component

```

1 State:
2   secrets ← {}
3
4 function fetchKeywordsFrequencies(triplestoreID, keywords, scheme):
5   frequencies ← {}
6   trapdoors ← [#keywords]
7   for keyword in keywords:
8     trapdoors.add(scheme.generateTrapdoor(keyword))
9   encryptedFrequencies ← Triplestore.search(triplestoreID, trapdoors)
10  for i in range(0, #encryptedFrequencies)
11    frequency ← encryptedFrequencies.get(i)
12    if frequency ≠ ⊥:
13      frequencies[keywords[i]] ← scheme.decryptContextualData(frequency)
14    else
15      frequencies[keywords[i]] ← 0
16  return frequencies
17
18 // {[node0, ..., noden], ⊥} ← SSE-TP.SearchTriplestore(triplestoreID, triplePatterns)
19 function search(triplestoreID, triplePatterns):
20   if secrets[triplestoreID] == ⊥
21     return ⊥
22   scheme ← createEncryptionScheme()
23   scheme.ReInitialize(secrets[triplestoreID])
24   keywords ← [#triplePatterns]
25   i ← 0
26   foreach t in triplePatterns:
27     keyword ← extractKeyword(triplePattern) // Parses the pattern and, based on its
28     // ↳ type, generates a keyword like the encryption scheme would. Outputs ⊥ if
29     // ↳ pattern not supported.
30     if keyword ≠ ⊥:
31       keywords[i] ← keyword
32       i++
33   frequencies ← fetchKeywordsFrequencies(triplestoreID, keywords, scheme)
34   trapdoors ← createLinkedList()
35   for keyword in keywords:
36     for i in range(0, frequencies[keyword]):
37       trapdoors.add(scheme.generateTrapdoorAndIncrementIV(keyword))
38   encryptedNodes ← Triplestore.search(triplestoreID, trapdoors)
39   results ← {}
40   foreach node in encryptedNodes:
41     if node ≠ ⊥:
42       results ← results ∪ { scheme.decrypt(node, false) }
43   return results
44
45 // {[t0, t1, ..., tn], ⊥} ← SSE-TP.SearchSchema(triplestoreID)
46 function searchSchema(triplestoreID):
47   if secrets[triplestoreID] == ⊥
48     return ⊥
49   scheme ← createEncryptionScheme()
50   scheme.ReInitialize(secrets[triplestoreID])
51   keyword ← scheme.getSchemaKeyword()
52   frequencies ← fetchKeywordsFrequencies(triplestoreID, [keyword], scheme)
53   trapdoors ← createLinkedList()
54   for i in range(0, frequencies[keyword]):
55     trapdoors.add(scheme.generateTrapdoorAndIncrementIV(keyword))
56   encryptedNodes ← Triplestore.search(triplestoreID, trapdoors)
57   schema ← {}
58   i ← 0
59   while i < frequencies[keyword]:
60     s ← scheme.decrypt(encryptedNodes.get(i), true)

```

```

59      $p \leftarrow \text{scheme.decrypt}(\text{encryptedNodes.get}(i + 1), \text{true})$ 
60      $o \leftarrow \text{scheme.decrypt}(\text{encryptedNodes.get}(i + 2), \text{true})$ 
61      $\text{schema} \leftarrow \text{schema} \cup \{ (s, p, o) \}$ 
62      $i \leftarrow i + 3$ 
63     return schema

```

Listing B.6: Searches in SSE-TP Client (SEnTriS V1)

```

1  State:
2      secrets  $\leftarrow \{\}$ 
3
4  //  $\{[node_0, \dots, node_n], \perp\} \leftarrow \text{SSE-TP.SearchTriplestore}(\text{triplestoreID}, \text{triplePatterns})$ 
5  function search(triplestoreID, triplePatterns):
6      if secrets[triplestoreID]  $= \perp$ 
7          return  $\perp$ 
8      scheme  $\leftarrow \text{createEncryptionScheme}()$ 
9      scheme.ReInitialize(secrets[triplestoreID])
10     keywords  $\leftarrow [\#triplePatterns]$ 
11      $i \leftarrow 0$ 
12     foreach t in triplePatterns:
13         keyword  $\leftarrow \text{extractKeyword}(\text{triplePattern})$ 
14         if keyword  $\neq \perp$ :
15             keywords[i]  $\leftarrow \text{keyword}$ 
16              $i++$ 
17     frequencies  $\leftarrow \text{fetchKeywordsFrequencies}(\text{triplestoreID}, \text{keywords}, \text{scheme})$ 
18     trapdoors  $\leftarrow \text{createLinkedList}()$ 
19     for keyword in keywords:
20         for i in range(0, frequencies[keyword]):
21             trapdoors.add(scheme.generateTrapdoorAndIncrementIV(keyword))
22     encryptedEqTags  $\leftarrow \text{Triplestore.search}(\text{triplestoreID}, \text{trapdoors})$ 
23     encryptedNodes  $\leftarrow \text{Triplestore.search}(\text{triplestoreID}, \text{encryptedEqTags})$ 
24     results  $\leftarrow \{\}$ 
25     foreach node in encryptedNodes:
26         if node  $\neq \perp$ :
27             results  $\leftarrow \text{results} \cup \{ \text{scheme.decrypt}(\text{node}, \text{false}) \}$ 
28     return results

```

Listing B.7: Searches in SSE-TP Client (SEnTriS V2)

```

1  State:
2      secrets  $\leftarrow \{\}$ 
3      BATCH_SIZE  $\leftarrow \text{fetchConfig}(\text{BATCH\_SIZE})$ 
4
5  //  $\{\text{SUCCESS}, \perp\} \leftarrow \text{SSE-TP.CreateTriplestore}(\text{triplestoreID}, [t_0, t_1, \dots, t_n])$ 
6  function createTriplestore(triplestoreID, triplesABox):
7      scheme  $\leftarrow \text{createEncryptionScheme}()$ 
8      secrets[triplestoreID]  $\leftarrow \text{scheme.getSecrets}()$ 
9      return executeUpdate(triplestoreID,  $\{\}$ , triplesABox, scheme)
10
11 //  $\{\text{SUCCESS}, \perp\} \leftarrow \text{SSE-TP.DeleteTriplestore}(\text{triplestoreID}, \text{schema})$ 
12 function deleteTriplestore(triplestoreID, schema):
13     if secrets[triplestoreID]  $= \perp$ 
14         return  $\perp$ 
15     else:
16         if schema == true:
17             return deleteTBox(triplestoreID, schema)
18         return Triplestore.delete(triplestoreID)
19

```

```

20 // {SUCCESS,  $\perp$ }  $\leftarrow$  SSE-TP.Upload(triplestoreID,  $[t_0, t_1, \dots, t_n]$ , schema)
21 function upload(triplestoreID, triples, schema):
22     if secrets[triplestoreID] ==  $\perp$ 
23         return  $\perp$ 
24     scheme  $\leftarrow$  createEncryptionScheme()
25     scheme.ReInitialize(secrets[triplestoreID])
26     if schema == true:
27         return uploadTBox(triplestoreID, triples, scheme)
28     return executeUpdate(triplestoreID, {}, triples, scheme)
29
30 // {SUCCESS,  $\perp$ }  $\leftarrow$  SSE-TP.Delete(triplestoreID,  $[t_0, t_1, \dots, t_n]$ )
31 function delete(triplestoreID, triples):
32     if secrets[triplestoreID] ==  $\perp$ 
33         return  $\perp$ 
34     scheme  $\leftarrow$  createEncryptionScheme()
35     scheme.ReInitialize(secrets[triplestoreID])
36     return executeUpdate(triplestoreID, triples, {}, scheme)
37
38 // {SUCCESS,  $\perp$ }  $\leftarrow$  SSE-TP.UpdateTriplestore(triplestoreID, triplesdelete, triplesupload)
39 function updateTriplestore(triplestoreID, triplesdelete, triplesupload):
40     if secrets[triplestoreID] ==  $\perp$ 
41         return  $\perp$ 
42     scheme  $\leftarrow$  createEncryptionScheme()
43     scheme.ReInitialize(secrets[triplestoreID])
44     return executeUpdate(triplestoreID, triplesdelete, triplesupload)
45
46 function executeUpdate(triplestoreID, triplesdelete, triplesupload):
47     deleteIDs  $\leftarrow$  batchExecute(triplestoreID, scheme, triplesdelete, "DELETION")
48     uploadIDs  $\leftarrow$  batchExecute(triplestoreID, scheme, triplesupload, "UPLOAD")
49     return Triplestore.update(triplestoreID, uploadIDs, deleteIDs)
50
51 function batchExecute(triplestoreID, scheme, triples, opType):
52     batchIDs  $\leftarrow$  createLinkedList()
53     while #triples > 0:
54         batch  $\leftarrow$  createLinkedList()
55         for i in range(0, BATCH_SIZE):
56             if #triples > 0:
57                 batch.add(triples.removeFirst())
58         if #batch > 0:
59             batchID  $\leftarrow$  batch(triplestoreID, scheme, batch, opType)
60             if batchID ==  $\perp$ :
61                 throw BatchException()
62             batchIDs.add(batchID)
63         scheme.clearNodes()
64     scheme.clearFrequencies()
65     return batchIDs
66
67 function batch(triplestoreID, scheme, batch, opType):
68     if opType == "UPLOAD":
69         return prepareUpload(triplestoreID, batch, scheme)
70     elif opType == "DELETION":
71         return prepareDelete(triplestoreID, batch, scheme)
72     else:
73         throw IllegalStateException()
74
75 function prepareUpload(triplestoreID, triples, scheme):
76     keywords  $\leftarrow$  [#triples]
77     i  $\leftarrow$  0
78     foreach t in triples:
79         keyword  $\leftarrow$  extractKeyword(triplePattern)

```

```

80     if keyword  $\neq$   $\perp$ :
81         keywords[i]  $\leftarrow$  keyword
82         i++
83     frequencies  $\leftarrow$  fetchKeywordsFrequencies(triplestoreID, keywords, scheme)
84     scheme.setKeywordFrequencies(frequencies)
85     encryptedNodes  $\leftarrow$  scheme.encrypt(triples, false)
86     return Triplestore.prepareUpload(triplestoreID, encryptedNodes)
87
88 function prepareDelete(triplestoreID, triples, scheme):
89     keywordsTrapdoors  $\leftarrow$  scheme.generateTriplePatternTrapdoors(triples)
90     encryptedFrequencies  $\leftarrow$  Triplestore.search(triplestoreID, triplePatternTrapdoors)
91     keywords  $\leftarrow$  createArrayList(#keywordsTrapdoors.getKeys())
92     totalTrapdoors  $\leftarrow$  0
93     foreach keywordTrapdoors in keywordsTrapdoors.getValues():
94         totalTrapdoors  $\leftarrow$  totalTrapdoors + #keywordTrapdoors
95     trapdoors  $\leftarrow$  createArrayList(totalTrapdoors)
96     foreach keyword in keywords:
97         foreach keywordTrapdoors in keywordsTrapdoors[keyword]:
98             foreach trapdoor in keywordTrapdoors:
99                 trapdoors.add(trapdoor)
100     encryptedFrequencies  $\leftarrow$  Triplestore.search(triplestoreID, trapdoors)
101     offset  $\leftarrow$  0
102     deletions  $\leftarrow$  createLinkedList()
103     foreach keyword in keywords:
104         length  $\leftarrow$  #keywordsTrapdoors[keyword]
105         for i in range(offset, offset + length)
106             encryptedFrequency  $\leftarrow$  encryptedFrequencies.get(i)
107             if encryptedFrequency  $\neq$   $\perp$ :
108                 frequency  $\leftarrow$  scheme.decryptContextualData(encryptedFrequency)
109                 deletions.add(scheme.generateTrapdoor(keyword, frequency))
110                 deletions.add(trapdoors.get(i))
111             offset  $\leftarrow$  offset + length
112     return Triplestore.prepareDelete(triplestoreID, trapdoors)
113
114 function uploadTBox(triplestoreID, triples, scheme):
115     keyword  $\leftarrow$  scheme.getSchemaKeyword()
116     frequencies  $\leftarrow$  fetchKeywordsFrequencies(triplestoreID, [keyword], scheme)
117     deletionIDs  $\leftarrow$  {}
118     if frequencies[keyword] > 0:
119         deletionIDs  $\leftarrow$  batchExecuteDeleteTBox(triplestoreID, frequencies[keyword], scheme)
120     uploadIDs  $\leftarrow$  batchExecuteUploadTBox(triplestoreID, triples, frequencies[keyword], scheme)
121     return Triplestore.update(triplestoreID, uploadIDs, deleteIDs)
122
123 function deleteTBox(triplestoreID, triples, scheme):
124     keyword  $\leftarrow$  scheme.getSchemaKeyword()
125     frequencies  $\leftarrow$  fetchKeywordsFrequencies(triplestoreID, [keyword], scheme)
126     deletionIDs  $\leftarrow$  {}
127     if frequencies[keyword] > 0:
128         deletionIDs  $\leftarrow$  batchExecuteDeleteTBox(triplestoreID, keyword, frequencies[keyword],
129              $\hookrightarrow$  scheme)
130     return Triplestore.update(triplestoreID, uploadIDs, deleteIDs)
131
132 function batchExecuteDeleteTBox(triplestoreID, frequency, keyword, scheme):
133     batchIDs  $\leftarrow$  createLinkedList()
134     while #triples > 0:
135         batch  $\leftarrow$  createLinkedList()
136         for i in range(0, BATCH_SIZE):
137             if #triples > 0:
138                 batch.add(triples.removeFirst())
139         if #batch > 0:

```

```

139     offset ← 0
140     while true:
141         trapdoors ← createLinkedList()
142         for i in range(offset, BATCH_SIZE + offset):
143             if i < frequency:
144                 trapdoors.add(scheme.generateTrapdoorAndIncrementIV(keyword))
145             if #trapdoors > 0:
146                 batchID ← Triplestore.prepareDelete(triplestoreID, encryptedNodes)
147                 if batchID == ⊥:
148                     throw BatchException()
149                 batchIDs.add(batchID)
150                 offset ← offset + BATCH_SIZE
151             else:
152                 break
153     return batchIDs
154
155 function batchExecuteUploadTBox(triplestoreID, triples, scheme):
156     batchIDs ← createLinkedList()
157     while #triples > 0:
158         batch ← createLinkedList()
159         for i in range(0, BATCH_SIZE):
160             if #triples > 0:
161                 batch.add(triples.removeFirst())
162         if #batch > 0:
163             encryptedNodes ← scheme.encrypt(triples, true)
164             batchID ← Triplestore.prepareUpload(triplestoreID, encryptedNodes)
165             if batchID == ⊥:
166                 throw BatchException()
167             batchIDs.add(batchID)
168             scheme.clearNodes()
169     return batchIDs

```

Listing B.8: Write operations in SSE-TP Client (SEnTriS V1)

```

1 State:
2     secrets ← {}
3     BATCH_SIZE ← fetchConfig(BATCH_SIZE)
4
5 function batchExecute(triplestoreID, scheme, triples, opType):
6     batchIDs ← createLinkedList()
7     while #triples > 0:
8         batch ← createLinkedList()
9         for i in range(0, BATCH_SIZE):
10             if #triples > 0:
11                 batch.add(triples.removeFirst())
12         if #batch > 0:
13             batchID ← batch(triplestoreID, scheme, batch, opType)
14             if batchID == ⊥:
15                 throw BatchException()
16             batchIDs.add(batchID)
17             scheme.clearNodes()
18             scheme.clearEqTags()
19             scheme.clearFrequencies()
20     return batchIDs
21
22 function prepareUpload(triplestoreID, triples, scheme):
23     keywords ← [#triples]
24     i ← 0
25     foreach t in triples:

```

```

26     keyword ← extractKeyword(triplePattern)
27     if keyword ≠ ⊥:
28         keywords[i] ← keyword
29         i++
30 frequencies ← fetchKeywordsFrequencies(triplestoreID, keywords, scheme)
31 scheme.setKeywordFrequencies(frequencies)
32 eqTags ← fetchEqTags(triplestoreID, triples, scheme)
33 scheme.setEqTags(eqTags)
34 encryptedNodes ← scheme.encrypt(triples, false)
35 return Triplestore.prepareUpload(triplestoreID, encryptedNodes)
36
37 function fetchEqTags(triplestoreID, triples, scheme):
38     eqTags ← {}
39     eqTagTrapdoors ← createLinkedList()
40     nodes ← createLinkedList()
41     processed ← {}
42     foreach (s,p,o) in triples
43         parseds ← parseNode(s)
44         parsedp ← parseNode(p)
45         parsedo ← parseNode(o)
46         if s ∉ processed:
47             nodes.add(parseds)
48             eqTagTrapdoors.add(scheme.generateTrapdoor(parseds))
49             processed.add(parseds)
50         if p ∉ processed:
51             nodes.add(parsedp)
52             eqTagTrapdoors.add(scheme.generateTrapdoor(parsedp))
53             processed.add(parsedp)
54         if o ∉ processed:
55             nodes.add(parsedo)
56             eqTagTrapdoors.add(scheme.generateTrapdoor(parsedo))
57             processed.add(parsedo)
58 encryptedEqTags ← Triplestore.search(triplestoreID, eqTagTrapdoors)
59 i ← 0
60 foreach node in nodes:
61     encryptedEqTag = encryptedEqTags.get(i)
62     if encryptedEqTag ≠ ⊥:
63         eqTags[node] ← scheme.decryptContextualData(encryptedEqTag)
64     i++
65 return eqTags
66
67 function prepareDelete(triplestoreID, triples, scheme):
68     keywordsTrapdoors ← scheme.generateTriplePatternTrapdoors(triples)
69     encryptedFrequencies ← Triplestore.search(triplestoreID, triplePatternTrapdoors)
70     keywords ← createArrayList(#keywordsTrapdoors.getKeys())
71     totalTrapdoors ← 0
72     foreach keywordTrapdoors in keywordsTrapdoors.getValues():
73         totalTrapdoors ← totalTrapdoors + #keywordTrapdoors
74     trapdoors ← createArrayList(totalTrapdoors)
75     foreach keyword in keywords:
76         foreach keywordTrapdoors in keywordsTrapdoors[keyword]:
77             foreach trapdoor in keywordTrapdoors:
78                 trapdoors.add(trapdoor)
79     encryptedFrequencies ← Triplestore.search(triplestoreID, trapdoors)
80     offset ← 0
81     deletions ← createLinkedList()
82     eqTagTrapdoors ← createLinkedList()
83     foreach keyword in keywords:
84         length ← #keywordsTrapdoors[keyword]
85         for i in range(offset, offset + length)

```

```
86   encryptedFrequency ← encryptedFrequencies.get(i)
87   if encryptedFrequency ≠ ⊥:
88       frequency ← scheme.decryptContextualData(encryptedFrequency)
89       eqTagTrapdoor ← scheme.generateTrapdoor(keyword, frequency)
90       deletions.add(eqTagTrapdoor)
91       deletions.add(trapdoors.get(i))
92       eqTagTrapdoors.add(eqTagTrapdoor)
93   offset ← offset + length
94   encryptedEqTags ← Triplestore.search(triplestoreID, eqTagTrapdoors)
95   foreach encryptedEqTag in encryptedEqTags:
96       if encryptedEqTag ≠ ⊥:
97           deletions.add(encryptedEqTag)
98   return Triplestore.prepareDelete(triplestoreID, trapdoors)
```

Listing B.9: Write operations in SSE-TP Client (SEnTriS V2)

SSE-SPARQL-EVAL LISTINGS

```

1  State:
2      subClasses ← {}
3      equivalentClasses ← {}
4      classRestrictions ← {}
5      intersectionClasses ← {}
6      intersectionsWhereClassIsOperand ← {}
7      subProperties ← {}
8      equivalentProperties ← {}
9      inverseProperties ← {}
10     symmetricProperties ← createLinkedList()
11     transitiveProperties ← createLinkedList()
12     transitivityDepth ← ⊥
13     expansionDepth ← ⊥
14
15 function execInference(schema, e, t):
16     expansionDepth ← e
17     transitivityDepth ← t
18     tbox = JENAModelFactory.createOntologyModel(JENAspec.OWL_MEM)
19     JENAGraphUtil.add(tbox.getGraph(), schema)
20     model = JENAModelFactory.createOntologyModel(JENAspec.OWL_MEM_TRANS_INF, tbox)
21     execClassInference(model.listClasses())
22     execPropertyInference(model.listAllOntProperties())
23
24 function execClassInference(classes, restrictions, intersections):
25     foreach class = (subclasses, equivalentClasses) in classes:
26         subClasses[convertToNode(class)] ← shuffle(subclasses)
27         equivalentClasses[convertToNode(class)] ← shuffle(equivalentClasses)
28     foreach restriction = (type, property, value) in restrictions:
29         type ∈ {"owl:hasValue", "owl:isSomeValuesFrom", "owl:allValuesFrom"}:
30         classRestrictions[convertToNode(restriction)] ← restriction
31     foreach intersection = (id, operands) in intersections: // Intersections are anonymous
32         ↪ classes with a set of other classes (operands)
33     foreach operand in operands:
34         intersectionsDirectSuperclasses ← intersectionsWhereClassIsOperand[convertToNode(
35             ↪ operand)]
36         if intersectionsDirectSuperclasses == ⊥:
37             intersectionsDirectSuperclasses ← createLinkedList()
38         intersectionsDirectSuperclasses.add(intersection)
39         intersectionsWhereClassIsOperand[convertToNode(operand)] ←
40             ↪ intersectionsDirectSuperclasses
41         intersectionClasses[convertToNode(intersection)] ← shuffle(operands)
42     foreach (node, intersectionsDirectSuperclasses) in intersectionsWhereClassIsOperand:

```

```

40     intersectionsWhereClassIsOperand[node] ← shuffle(intersectionsDirectSuperclasses)
41
42 function execPropertyInference(properties):
43     foreach p=(type, value, subproperties, equivalentProperties, inverseProperties) in properties:
44         if type == "owl:SymmetricProperty":
45             symmetricProperties.add(p)
46         elif type == "owl:TransitiveProperty":
47             transitiveProperties.add(p)
48         subProperties[convertToNode(p)] ← shuffle(subproperties)
49         equivalentProperties[convertToNode(p)] ← shuffle(equivalentProperties)
50         inverseProperties[convertToNode(p)] ← shuffle(inverseProperties)

```

Listing C.1: SSE-SPARQL-EVAL ontology

```

1  State:
2      operationsparsed ← {}
3      uploadtemplate ← {}
4      deletetemplate ← {}
5      constructtemplate ← {}
6      queryType ← :bot
7      plan ← (jobs = {}, execOrder = createLinkedList(), vars = {})
8      ontology ← ⊥
9      classExpansionOrder ← ["SUBCLASS", "EQUIVALENT_CLASS", "INTERSECTION_OPERAND", "
      ↳ RESTRICTION", "INTERSECTION"]
10     propertyExpansionOrder ← ["SUBPROPERTY", "EQUIVALENT_PROPERTY", "SYMMETRIC", "
      ↳ TRANSITIVE", "INVERSE_OF"]
11     obfuscationMap ← {}
12     deobfuscationMap ← {}
13     keywords ← {}
14     searchJobIDs ← {}
15
16 function generatePlan(op)
17     JENA.OpWalker.walk(op, this)
18     if #plan.vars == 0:
19         plan.vars ← obfuscationMap.getKeys()
20     return plan
21
22 function getKeywords():
23     return keywords
24
25 function getSearchJobIDs():
26     return searchJobIDs
27
28 function getDeobfuscationMap():
29     return deobfuscationMap
30
31 function obfuscate(var):
32     varobfuscated ← obfuscationMap[var]
33     if varobfuscated == ⊥:
34         varobfuscated ← generateID()
35         obfuscationMap[var] ← varobfuscated
36         deobfuscationMap[varobfuscated] ← var
37     return varobfuscated
38
39 function generateSecureSearchJob(var, node1, node2, pattern):
40     searches ← {}
41     var ← obfuscate(var)
42     jobID ← generateID()
43     keyword ← pattern || parseKeyword(node1) || parseKeyword(node2)

```

```

44     searches[var] ← keyword
45     keywords ← keywords ∪ { keyword }
46     searchJobsIDs ← searchJobsIDs ∪ { jobID }
47     return createSecureSearchJob(jobID, [var], searches)
48
49 function generateSecureSearchJob(var1, var2, node, pattern):
50     searches ← {}
51     jobID ← generateID()
52     keyword ← pattern || parseKeyword(node)
53     keywords ← keywords ∪ { keyword }
54     searchJobsIDs ← searchJobsIDs ∪ { jobID }
55     var1 = obfuscate(var1);
56     var2 = obfuscate(var2);
57     searches[var1] ← keyword
58     searches[var2] ← keyword
59     return createSecureSearchJob(jobID, [var1, var2], searches);
60
61 function pushSecureSearch(s, p, o, jobs, jobIDs)
62     searchPattern ← parsePattern(s, p, o)
63     jobID = jobIDs.get(searchPattern)
64     if jobID == ⊥:
65         jobID ← generateID()
66         pattern ← extractVariablesPattern(s, p, o)
67         jobSEC_SEARCH ← ⊥
68         switch (pattern):
69             case "S": jobSEC_SEARCH ← generateSecureSearchJob(s, p, o, pattern)
70             case "P": jobSEC_SEARCH ← generateSecureSearchJob(p, s, o, pattern)
71             case "O": jobSEC_SEARCH ← generateSecureSearchJob(o, s, p, pattern)
72             case "SP": jobSEC_SEARCH ← generateSecureSearchJob(s, p, o, pattern)
73             case "SO": jobSEC_SEARCH ← generateSecureSearchJob(s, o, p, pattern)
74             case "PO": jobSEC_SEARCH ← generateSecureSearchJob(p, o, s, pattern)
75             default: throw NotImplementedException()
76         jobs[jobID] ← jobSEC_SEARCH
77         jobIDs[searchPattern] ← jobID
78         pushJob(jobSEC_SEARCH)
79     return jobID
80
81 function generateSearchJobs(opOpBGP):
82     patterns = opOpBGP.getPattern()
83     jobIDs = {}
84     jobs = {}
85     jobVars = {}
86     foreach t=(s,p,o) in patterns:
87         jobID = pushSecureSearch(s, p, o, jobs, jobIDs)
88         if ontology ≠ ⊥:
89             if (p == "rdf:type") and (o is not VARIABLE):
90                 jobID = expandClass(jobID, s, o, 0, jobs, jobIDs)
91             elif p is not VARIABLE:
92                 jobID = expandProperty(jobID, s, p, o, 0, jobs, jobIDs, true, true)
93         job ← jobs[jobID]
94         if job is SEC_SEARCH:
95             varsjob ← job.getVariables()
96             variables ← variables ∪ varsjob
97             jobVars[jobID] ← varsjob
98         elif job is BINARY:
99             jobVars[jobID] ← extractVarsOfBinaryJob(jobs, job)
100     if #patterns > 1:
101         generateRandomJoinPipeline(opOpBGP, jobVars)
102     else:
103         operationsparsed.put(opOpBGP, jobID)

```

```

104
105 function expandClass(jobID, s, o, depth, jobs, jobsIDs):
106     if depth == ontology.getMaximumExpansionDepth():
107         return jobID
108     foreach expansionType in shuffle(classExpansionOrder):
109         if expansionType == "SUBCLASS":
110             subclasses ← ontology.getSubClasses(o)
111             jobID ← expandClassDisjunction(subclasses, jobID, s, depth, jobs, jobsIDs)
112         elif expansionType == "EQUIVALENT_CLASS":
113             equivalentClasses ← ontology.getEquivalentClasses(o)
114             jobID ← expandClassDisjunction(equivalentClasses, jobID, s, depth, jobs, jobsIDs)
115         elif expansionType == "RESTRICTION":
116             restriction ← ontology.getRestriction(o)
117             jobID ← expandRestriction(restriction, jobID, s, depth, jobs, jobsIDs)
118         elif expansionType == "INTERSECTION_OPERAND":
119             intersectionClasses ← ontology.getIntersectionWhereClassIsOperand(o)
120             jobID ← expandClassDisjunction(intersectionClasses, jobID, s, depth, jobs, jobsIDs)
121             ↪ )
122         elif expansionType == "INTERSECTION":
123             intersection ← ontology.getIntersection(o)
124             if #intersection == 0:
125                 jobID ← expandClassConjunction(intersection, jobID, s, depth, jobs, jobsIDs)
126     return jobID
127
128 function expandProperty(jobID, s, p, o, depth, jobs, jobsIDs, symmetric, inverseOf):
129     if depth == ontology.getMaximumExpansionDepth():
130         return jobID
131     foreach expansionType in shuffle(classExpansionOrder):
132         if expansionType == "SYMMETRIC":
133             if symmetric == true and ontology.isSymmetric(p):
134                 jobID ← pushUnion(jobID, pushSearch(o, p, s, jobs, jobsIDs), jobs, jobsIDs)
135                 jobID ← expandProperty(jobID, o, p, s, depth + 1, jobs, jobsIDs, false,
136                 ↪ inverseOf)
137             elif expansionType == "TRANSITIVE":
138                 if ontology.isTransitive(p):
139                     var ← createVariable(parse(p) || 0)
140                     left ← pushSearch(s, p, var, jobs, jobsIDs)
141                     join ← pushJoin(left, pushSearch(var, p, o, jobs, jobsIDs), jobs, jobsIDs)
142                     jobID ← pushUnion(jobID, join, jobs, jobsIDs)
143                     for i in range(0, ontology.getTransitivityDepth())
144                         for j in range(0, i + 1):
145                             if j == i:
146                                 varnext = createVariable(parse(p) || j)
147                                 left ← pushJoin(left, pushSearch(var, p, varnext, jobs, jobsIDs),
148                                 ↪ jobs, jobsIDs)
149                                 var = varnext
150                             join ← pushJoin(left, pushSearch(var, p, o, jobs, jobsIDs), jobs, jobsIDs)
151                             jobID ← pushUnion(jobID, join, jobs, jobsIDs)
152             elif expansionType == "SUBPROPERTY":
153                 subproperties ← ontology.getSubProperties(p)
154                 jobID ← expandProperties(subproperties, jobID, s, o, depth, jobs, jobsIDs,
155                 ↪ symmetric, inverseOf)
156             elif expansionType == "EQUIVALENT_PROPERTY":
157                 equivalentProperties ← ontology.getEquivalentProperties(p)
158                 jobID ← expandProperties(equivalentProperties, jobID, s, o, depth, jobs, jobsIDs,
159                 ↪ symmetric, inverseOf)
160             elif expansionType == "INVERSE_OF":
161                 if inverseOf == true:
162                     inverseProperty ← ontology.getInverseOf(p)

```

```

158         jobID ← expandProperties(inverseProperty, jobID, o, s, depth, jobs, jobsIDs,
159             ↳ symmetric, false)
        return jobID

```

Listing C.2: SSE-SPARQL-EVAL SPARQL Planner

```

1  State:
2      queryEngine ← createSPARQLQueryEngine()
3      planner ← ⊥
4      secrets ← {}
5
6  function executeSPARQLQuery(triplestoreID, query, inference, expansionDepth, transitivityDepth):
7      if secrets[triplestoreID] == ⊥
8          return ⊥
9      try:
10         scheme ← createEncryptionScheme()
11         scheme.ReInitialize(secrets[triplestoreID])
12         (planrnd, construct, upload, delete) ← planQuery(query, triplestoreID, inference,
13             ↳ expansionDepth, transitivityDepth)
14         queryType ← planner.getQueryType()
15         if queryType ∈ { SELECT, CONSTRUCT, ASK, DESCRIBE, MODIFY,
16             ↳ DELETE_WHERE }:
17             preparedPlan ← prepareQuery(triplestoreID, planrnd, scheme, ⊥)
18             sparqlResult ← Proxy.evaluateQuery(preparedPlan, scheme.getRNDKey())
19             deobfuscationMap ← planner.getDeobfuscationMap()
20             bindings ← decryptBindings(sparqlResult.bindings, deobfuscationMap, scheme)
21             bindings ← applyModifiers(sparqlResult, deobfuscationMap, bindings)
22             return generateFinalResult(triplestoreID, queryType, plan.vars, bindings, construct,
23                 ↳ upload, delete)
24         catch Exception:
25             return ⊥
26
27  function prepareQuery(triplestoreID, planner, planrnd = (jobs, rndExecOrder), scheme, mask):
28      keywordsFrequency ← fetchKeywordsFrequencies(triplestoreID, planner.getKeywords(),
29          ↳ scheme)
30      searchIDCollector ← {}
31      foreach jobID in planner.getSearchJobIDs():
32          jobSEC_SEARCH ← jobs[jobID]
33          jobvars = jobSEC_SEARCH.getVariables()
34          keyword = jobSEC_SEARCH.getSearches()[jobvars[0]]
35          searchIDs ← searchIDCollector.get(keyword);
36          if searchIDs == ⊥:
37              searchIDs ← createArrayList(#jobvars)
38              keywordFrequency ← keywordsFrequency[keyword]
39              trapdoors ← createArrayList(jobvars)
40              for i in range(0, #jobvars)
41                  trapdoors.add(createArrayList(keywordFrequency / #jobvars))
42              for i in range(0, keywordFrequency)
43                  trapdoors.get(i mod #jobvars).add(scheme.generateTrapdoorAndIncrementIV(
44                      ↳ keyword))
45              for i in range(0, #jobvars)
46                  searchID = Triplestore.prepareSearch(triplestoreID, mask, trapdoors.get(i))
47                  searchIDs.add(searchID)
48                  jobSEC_SEARCH.getSearches()[jobvars[i]] ← searchID
49                  searchIDCollector[keyword] ← searchIDs
50          else:
51              for i in range(0, #jobvars)
52                  jobSEC_SEARCH.getSearches()[jobvars[i]] ← searchIDs.get(i)
53      return planrnd

```

Listing C.3: Query Preparation in SSE-SPARQL-EVAL Client (SEnTriS V1)

```

1 State:
2   queryEngine ← createSPARQLQueryEngine()
3   planner ← ⊥
4   secrets ← {}
5
6 function executeSPARQLQuery(triplestoreID, query, inference, expansionDepth, transitivityDepth):
7   if secrets[triplestoreID] == ⊥
8     return ⊥
9   try:
10    scheme ← createEncryptionScheme()
11    scheme.ReInitialize(secrets[triplestoreID])
12    (planrnd, construct, upload, delete) ← planQuery(query, triplestoreID, inference,
13      ↪ expansionDepth, transitivityDepth)
14    queryType ← planner.getQueryType()
15    if queryType ∈ { SELECT, CONSTRUCT, ASK, DESCRIBE, MODIFY,
16      ↪ DELETE_WHERE }:
17      mask ← generateMask(scheme.getPubDGK())
18      preparedPlan ← prepareQuery(triplestoreID, planrnd, scheme, mask)
19      sparqlResult ← Proxy.evaluateQuery(preparedPlan, scheme.getEqKey())
20      inverseMask ← generateInverseMask(mask, scheme.getPubDGK())
21      deobfuscationMap ← planner.getDeobfuscationMap()
22      bindings ← fetchAndDecryptBindings(triplestoreID, sparqlResult.bindings,
23        ↪ deobfuscationMap, scheme)
24      bindings ← applyModifiers(sparqlResult, deobfuscationMap, bindings)
25      return generateFinalResult(triplestoreID, queryType, plan.vars, bindings, construct,
26        ↪ upload, delete)
27    catch Exception:
28      return ⊥
29
30 function generateMask( $K_{DGK_{pub}} = (n, g, h, u, k, t, l)$ ):
31   mask ← generateDGKRandom(t) // Generate random 2.5t-bit integer
32   return  $h^{mask} \bmod n$ 
33
34 function generateInverseMask(mask,  $K_{DGK_{pub}} = (n, g, h, u, k, t, l)$ ):
35   return  $h^{-mask} \bmod n$ 

```

Listing C.4: Query Preparation in SSE-SPARQL-EVAL Client (SEnTriS V2)

```

1 State:
2   storageEngine ← createStorageEngine()
3
4 // {searchID, ⊥} ← SSE-SPARQL.PrepareSearch(triplestoreID, mask, [trapdoor0, ..., trapdoorn])
5 function prepareSearch(triplestoreID, mask, trapdoors):
6   if mask == ⊥:
7     return Proxy.prepareBindings(storageEngine.search(triplestoreID, trapdoors))
8   return Proxy.prepareBindings(storageEngine.search(triplestoreID, mask, trapdoors))

```

Listing C.5: SSE-SPARQL-EVAL Triplestore

```

1 State:
2   TIMEOUT ← fetchConfig(TIMEOUT)
3   triplestores ← {}

```

```

4  uploads ← {}
5  deletions ← {}
6
7  function search(triplestoreID, mask, trapdoors):
8      result ← createArrayList(#trapdoors)
9      for i in range(0, #trapdoor):
10         binding ← triplestores[triplestoreID][trapdoors[i]]
11         if binding ≠ ⊥
12             result.add(binding * mask)
13     return result

```

Listing C.6: SSE-SPARQL-EVAL Triplestore Storage Engine

```

1  State:
2      storageEngine ← createStorageEngine()
3
4  // {searchID, ⊥} ← SSE-SPARQL-EVAL.PrepareBindings([E(binding0), ..., E(bindingn)])
5  function prepareBindings(bindings):
6      return storageEngine.save(bindings)

```

Listing C.7: Query Preparation in SSE-SPARQL-EVAL Proxy

```

1  State:
2      TIMEOUT ← fetchConfig(TIMEOUT)
3      bindings ← {}
4
5  function save(bindings):
6      searchID ← generateID()
7      bindings[searchID] ← bindings
8      schedule(expireBindings(searchID), TIMEOUT)
9      return searchID
10
11 function delete(searchIDs):
12     foreach searchID in searchIDs:
13         expireBindings(searchID)
14
15 function expireBindings(searchID):
16     if bindings[searchID] ≠ ⊥:
17         bindings[searchID] ← ⊥

```

Listing C.8: Persisting Bindings in SSE-SPARQL-EVAL Proxy Storage Engine

```

1  State:
2      storageEngine ← createStorageEngine()
3
4  // {E(sparqlResult), ⊥} ← SSE-SPARQL-EVAL.EvaluateQuery(preparedPlanrnd, K)
5  function evaluateQuery(plan, \; key):
6      worker ← createQueryWorker()
7      worker.setContext(key, storageEngine)
8      execution ← createQueryExecution()
9      execution.exec(plan, worker)
10     storageEngine.delete(worker.getSearchIDs())
11     return execution.getSPARQLResult()

```

Listing C.9: Query Evaluation in SSE-SPARQL-EVAL Proxy

```

1  State:
2      searchIDs ← {}
3      key ← ⊥
4      sparqlResult ← ({}, false, false, false, ⊥, ⊥)
5      storageEngine ← ⊥
6
7  function setContext(k, storage):
8      key ← k
9      storageEngine ← storage
10
11 function exec(job):
12     if job is SEC_SEARCH:
13         searches ← job.getSearches()
14         foreach (var, searchID) in searches
15             searchIDs ← searchIDs ∪ { searchID }
16         return storageEngine.search(key, job.getVariables(), searches)
17     elif job is EMPTY:
18         return createBindingsTable(job.vars)
19     throw JobInstanceException()
20
21 function generateResults(jobResults):
22     bindings ← ⊥
23     if sparqlResult.distinct == true:
24         bindings ← generateBindings(createHashSet(), jobResults)
25     else
26         bindings ← generateBindings(createLinkedList(), jobResults)
27     sparqlResult.bindings ← bindings
28     return sparqlResult.bindings
29
30 function generateBindings(collector, jobResults):
31     vars = jobResults.getVariables()
32     values ← {}
33     foreach patterns in jobResults.getPatterns():
34         binding ← {}
35         i ← 0
36         foreach encryptedBinding in patterns:
37             if encryptedBinding ≠ ⊥:
38                 binding[vars[i]] ← encryptedBinding
39             i++
40         collector.add(binding)
41     return collector
42
43 function getSearchIDs():
44     return searchIDs

```

Listing C.10: SSE-SPARQL-EVAL Query Worker

```

1  State:
2      TIMEOUT ← fetchConfig(TIMEOUT)
3      bindings ← {}
4
5  function search(KRND, variables, searches):
6      table ← createBindingsTable()
7      table.setVariables(variables)
8      searchResults ← {}
9      foreach variable in variables:
10         searchResults[variable] ← bindings[searches[variable]]
11     for i in range(0, #searchResults[variables[0]]):
12         patternID ← generateID()

```

```

13     foreach variable in variables:
14         table.add(patternID, variable, D( $K_{RND}$ , searchResults[variable][i]))
15     return table

```

Listing C.11: Constructing Bindings Tables (SEnTriS V1)

```

1  State:
2      searchIDs  $\leftarrow$  {}
3      key  $\leftarrow$   $\perp$ 
4      sparqlResult  $\leftarrow$  ({}, false, false, false,  $\perp$ ,  $\perp$ )
5      storageEngine  $\leftarrow$   $\perp$ 
6
7  function setContext( $k$ , storage):
8      key  $\leftarrow$   $k$ 
9      storageEngine  $\leftarrow$  storage
10
11 function execJoin(left, right):
12     return left.join(right, key)
13
14 function execOptional(left, right):
15     return left.leftOuterJoin(right, key)
16
17 function execMinus(left, right):
18     return left.minus(right, key)

```

Listing C.12: Keyed set operations in SSE-SPARQL-EVAL Query Worker (SEnTriS V2)

```

1  State:
2      TIMEOUT  $\leftarrow$  fetchConfig(TIMEOUT)
3      bindings  $\leftarrow$  {}
4
5  function search( $K_{DGK_{eq}} = (p, vp, n, u)$ , variables, searches):
6      table  $\leftarrow$  createBindingsTable()
7      table.setVariables(variables)
8      searchResults  $\leftarrow$  {}
9      foreach variable in variables:
10         searchBindingsvar  $\leftarrow$  bindings[searches[variable]]
11         for  $i$  in range(0, #bindingsvar):
12             searchResults[variable][ $i$ ]  $\leftarrow$  searchBindingsvar
13     foreach variable in variables:
14         for  $i$  in range(0, #searchResults[variable]):
15             eqTag  $\leftarrow$  searchResults[variable][ $i$ ]
16             searchResults[variable][ $i$ ]  $\leftarrow$  eqTag mod  $n$ 
17     groupedEqTags  $\leftarrow$  {}
18     foreach variable in variables:
19         groupedEqTags[variable]  $\leftarrow$  {}
20     for  $i$  in range(0, #searchResults[variables[0]]):
21         patternID  $\leftarrow$  generateID()
22         foreach variable in variables:
23             eqTagGroup  $\leftarrow$  findEqTagGroup( $K_{DGK_{eq}}$ , groupedEqTags[variable], searchResults[
24                  $\hookrightarrow$  variable][ $i$ ])
25             table.add(patternID, variable, eqTagGroup)
26     return table
27
28 function findEqTagGroup( $K_{DGK_{eq}}$ , groupedEqTags, newEqTag):
29     eqTagGroup  $\leftarrow$   $\perp$ 
30     foreach eqTag in groupedEqTags:

```

```

30     if equals( $K_{DGK_{eq}}$ , eqTag, newEqTag)
31         eqTagGroup  $\leftarrow$  eqTag
32     if eqTagGroup  $\neq \perp$ :
33         return eqTagGroup
34     groupedEqTags  $\leftarrow$  groupedEqTags  $\cup$  { newEqTag }
35     return newEqTag
36
37 function equals( $K_{DGK_{eq}} = (p, vp, n, u)$ ,  $ct_1$ ,  $ct_2$ ):
38      $\Delta \leftarrow ct_1 * ct_2^{(u-1)} \bmod n$ 
39     if  $\Delta^{vp} \bmod p = 1$ :
40         return true
41     return false

```

Listing C.13: Constructing Bindings Tables (SEnTriS V2)

```

1  State:
2      bindings  $\leftarrow$  {}
3      patterns  $\leftarrow$  {}
4
5  function equals( $K_{DGK_{eq}} = (p, vp, n, u)$ ,  $ct_1$ ,  $ct_2$ ):
6       $\Delta \leftarrow ct_1 * ct_2^{(u-1)} \bmod n$ 
7      if  $\Delta^{vp} \bmod p = 1$ :
8          return true
9      return false
10
11 function join(other,  $K_{DGK_{eq}}$ ):
12     vars_mutual  $\leftarrow$  copy(getVariables())
13     vars_mutual.retain(other.getVariables())
14     return join(vars_mutual, this, other, "INNER",  $K_{DGK_{eq}}$ )
15
16 function leftOuterJoin(other,  $K_{DGK_{eq}}$ ):
17     vars_mutual  $\leftarrow$  copy(getVariables())
18     vars_mutual.retain(other.getVariables())
19     return join(vars_mutual, this, other, "LEFT",  $K_{DGK_{eq}}$ )
20
21 function minus(other,  $K_{DGK_{eq}}$ ):
22     vars_mutual  $\leftarrow$  copy(getVariables())
23     vars_mutual.retain(other.getVariables())
24     diff  $\leftarrow$  {}
25     table_minus  $\leftarrow$  createBindingsTable()
26     table_minus.setVariables(vars)
27     foreach var : vars_mutual:
28         bindings_l  $\leftarrow$  getBindings(var)
29         bindings_r  $\leftarrow$  other.getBindings(var)
30         foreach (binding, patternIDs) in bindings_l:
31             if findPatterns( $K_{DGK_{eq}}$ , bindings_r, binding) ==  $\perp$ :
32                 diff  $\leftarrow$  diff  $\cup$  patternIDs
33     varBindings_RND  $\leftarrow$  getBindings(sample(vars_mutual))
34     foreach (binding, patternIDs) in varBindings_RND:
35         foreach patternID in patternIDs:
36             if patternID  $\in$  diff:
37                 copyBindings(patternID, patternID, vars_mutual, this, table_minus)
38     return table_minus
39
40 function join(vars_mutual, table_l, table_r, type,  $K_{DGK_{eq}}$ )
41     vars_l  $\leftarrow$  copy(table_l.getVariables())
42     vars_r  $\leftarrow$  copy(table_r.getVariables())
43     vars  $\leftarrow$  vars_l  $\cup$  vars_r

```

```

44  varsl ← varsl / varsr
45  varsr ← varsr / varsl
46  tablejoin ← createBindingsTable()
47  tablejoin.setVariables(vars)
48  var ← sample(varsmutual)
49  bindingsl = tablel.getBindings(var)
50  bindingsr = tabler.getBindings(var)
51  foreach (binding, patternIDsl) in bindingsl:
52    patternIDsr ← findPatterns(KDGKeq, bindingsr, binding)
53    if patternIDsr ≠ ⊥ or type == "LEFT":
54      joinPatterns(varsmutual, tablel, varsl, patternIDsl, tabler, varsr, patternIDsr,
        ↪ tablejoin, type, KDGKeq)
55  return tablejoin
56
57  function joinPatterns(varsmutual, tablel, varsl, patternIDsl, tabler, varsr, patternIDsr,
    ↪ tablejoin, type, KDGKeq)
58    foreach patternIDl in patternIDsl:
59      match ← false
60      if patternIDsr ≠ ⊥:
61        foreach patternIDr in patternIDsr:
62          if equalPatterns(varsmutual, tablel, patternIDl, tabler, patternIDr, KDGKeq):
63            match ← true
64            patternID = generateID()
65            copyBindings(patternID, patternIDl, varsmutual, tablel, tablejoin)
66            copyBindings(patternID, patternIDl, varsl, tablel, tablejoin)
67            copyBindings(patternID, patternIDr, varsr, tabler, tablejoin)
68          if match == false and type == "LEFT":
69            copyBindings(patternIDl, patternIDl, varsmutual, tablel, tablejoin)
70            copyBindings(patternIDl, patternIDl, varsl, tablel, tablejoin)
71
72  function equalPatterns(varsmutual, tablel, patternIDl, tabler, patternIDr, KDGKeq = (p, vp, n, u))
73    foreach var in varsmutual:
74      bindingl = tablel.getPatternIDs(var)[patternIDl]
75      bindingr = tabler.getPatternIDs(var)[patternIDr]
76      if bindingl == ⊥ and bindingr ≠ ⊥
77        return false
78      elif bindingl ≠ ⊥ and bindingr == ⊥
79        return false
80      elif bindingl ≠ ⊥ and bindingl ≠ bindingr
81        return equals(KDGKeq, bindingl, bindingr)
82    return true
83
84  function findPatterns(KDGKeq = (p, vp, n, u), bindings, target):
85    res ← ⊥
86    foreach eqTag in bindings:
87      if equals(KDGKeq, eqTag, target) = true:
88        res ← eqTag
89    if eqTagGroup ≠ ⊥:
90      return ⊥
91    return bindings[res]

```

Listing C.14: Bindings Tables (SEnTriS V2)

```

1  State:
2    queryEngine ← createSPARQLQueryEngine()
3    secrets ← {}
4    planner ← ⊥
5

```

```

6 function decryptBindings(bindings, deobfuscationMap, scheme):
7   decryptedBindings ← createLinkedList()
8   foreach binding in bindings:
9     bindingCollector ← {}
10    foreach (var, ct) in binding:
11      bindingCollector ← bindingCollector ∪ { (deobfuscationMap[var], scheme.
        ↳ decryptBinding(ct)) }
12    decryptedBindings.add(bindingCollector)
13  return decryptedBindings
14
15 function applyModifiers(sparqlResult, deobfuscationMap, bindings):
16   if sparqlResult.ordered = true:
17     bindings = orderResults(sparqlResult.distinct, sparqlResult.sortConditions,
        ↳ deobfuscationMap, bindings)
18   if sparqlResult.sliced = true:
19     bindings = slice(sparqlResult.offset, sparqlResult.length, bindings)
20  return bindings
21
22 function orderResults(distinct, obfuscatedSortConditions, deobfuscationMap, bindings):
23   sortConditions ← createLinkedList()
24   foreach (obfuscatedVar, dir) in obfuscatedSortConditions:
25     sortConditions.add((deobfuscationMap[obfuscatedVar], dir))
26   comparator ← JENABindingComparator.create(sortConditions)
27   if distinct = true:
28     sortedBindings ← createTreeSet(comparator)
29     foreach binding in bindings:
30       sortedBindings.add(binding)
31   else:
32     sortedBindings ← createLinkedList()
33     foreach binding in bindings:
34       sortedBindings.add(binding)
35     sort(sortedBindings, comparator)
36  return sortedBindings
37
38 // {sparqlFinalResult, ⊥} ← SSE-SPARQL-EVAL.generateFinalResult(triplestoreID, variables,
        ↳ queryType, bindings, construct, upload, delete):
39 function generateFinalResult(triplestoreID, queryType, variables, bindings, constructtemplate,
        ↳ uploadtemplate, deletetemplate, scheme)
40   if queryType == "SELECT":
41     return generateSELECTResults(variables, bindings)
42   elif queryType == "CONSTRUCT":
43     return generateCONSTRUCTResults(constructtemplate, bindings)
44   elif queryType == "ASK":
45     return generateASKResults(bindings)
46   elif queryType == "DESCRIBE":
47     return generateDESCRIBEResults(variables, sparqlResult)
48   else:
49     return applyModifications(triplestoreID, bindings, uploadtemplate, deletetemplate)
50
51 function applyModifications(triplestoreID, bindings, uploadtemplate, deletetemplate):
52   triplesupload ← {}
53   triplesdelete ← {}
54   if queryType == "MODIFY":
55     triplesupload ← generateTriplesFromBindings(uploadtemplate, bindings)
56     triplesdelete ← generateTriplesFromBindings(deletetemplate, bindings)
57   else if (queryType == INSERT_DATA)
58     triplesupload ← uploadtemplate
59   else if (queryType == DELETE_DATA)
60     triplesdelete ← deletetemplate
61  return updateTriplestore(triplestoreID, triplesdelete, triplesupload)

```

Listing C.15: Generating Final Results in SSE-SPARQL-EVAL (SEnTriS V1)

```

1 State:
2   queryEngine  $\leftarrow$  createSPARQLQueryEngine()
3   secrets  $\leftarrow$  {}
4   planner  $\leftarrow \perp$ 
5
6 function fetchAndDecryptBindings(triplestoreID, inverseMask, bindings, deobfuscationMap,
    $\hookrightarrow$  scheme):
7   decryptedBindings  $\leftarrow$  createLinkedList()
8   eqTagsOrder  $\leftarrow$  {}
9   eqTags  $\leftarrow$  createArrayList()
10  i  $\leftarrow$  0
11  KDGKpub  $\leftarrow$  scheme.getPubDGK()
12  foreach binding in bindings:
13    foreach (obfuscatedVar, eqTag) in binding:
14      if eqTag  $\notin$  eqTagsOrder:
15        eqTagsOrder[eqTag]  $\leftarrow$  i
16        eqTags.add(unmask(KDGKpub, inverseMask, eqTag))
17      i++
18  encryptedBindings  $\leftarrow$  Triplestore.search(triplestoreID, eqTags)
19  decryptedNodes  $\leftarrow$  {}
20  foreach binding in bindings:
21    bindingCollector  $\leftarrow$  {}
22    foreach (obfuscatedVar, eqTag) in binding:
23      i  $\leftarrow$  eqTagsOrder[eqTag]
24      decryptedNode  $\leftarrow$  decryptedNodes[i]
25      if decryptedNode =  $\perp$ :
26        decryptedNode  $\leftarrow$  scheme.decryptNode(encryptedBindings.get(i))
27        decryptedNodes[i]  $\leftarrow$  decryptedNode
28        bindingCollector  $\leftarrow$  bindingCollector  $\cup$  { (deobfuscationMap[var], decryptedNode) }
29      decryptedBindings.add(bindingCollector)
30  return decryptedBindings
31
32 function unmask(KDGKpub = (n, g, h, u, k, t, l), mask, eqTag):
33  return eqTag * mask mod n

```

Listing C.16: Generating Final Results in SSE-SPARQL-EVAL (SEnTriS V2)

IAM LISTINGS

```

1 function registerUser(username, password):
2     return IAM.saveUser(username, password)
3
4 function auth(username, password):
5     response ← IAM.createSession(username, password)
6     if response ≠ ⊥:
7         return extractCookie(response)
8     return response
9
10 function unregisterUser(cookie, username):
11     return IAM.deleteUser(cookie, username)

```

Listing D.1: Authentication in Client

```

1 State:
2     storageEngine ← createStorageEngine()
3     locksClient ← createLocksClient()
4
5 function createSession(username, password):
6     cookie ← ⊥
7     if checkPassword(username, password):
8         cookie ← storageEngine.cacheSession(username);
9     return cookie
10
11 function saveUser(username, password):
12     try:
13         lockID ← locksClient.acquireUserLock(username)
14         if storageEngine.userExists(username) = false:
15             lockID ← locksClient.releaseUserLock(username)
16             return ⊥
17         hash ← hash(password) // Generates a hash using a one-way key derivation function,
18             ↪ and a randomly generated salt
19         storageEngine.saveUser(username, hash, "BASIC")
20         locksClient.releaseUserLock(username, lockID)
21         return "SUCCESS"
22     catch Exception:
23         return ⊥
24
25 function deleteUser(cookie, username):
26     try:
27         authCheck(cookie, username)

```

```

27     lockID ← locksClient.acquireUserLock(username)
28     if storageEngine.checkIfOwnsAny(username): // Triplestores owned by the user must
        ↳ be deleted beforehand
29         locksClient.releaseUserLock(username, lockID)
30         return ⊥
31     storageEngine.deleteUser(username)
32     locksClient.deleteAllUserLocks(username)
33     locksClient.releaseUserLock(username, lockID)
34     return "SUCCESS"
35 catch Exception:
36     return ⊥
37
38 function checkPassword(username, password):
39     hash ← storageEngine.getPassword(username)
40     if hash ← ⊥
41         salt ← extractSalt(hash)
42         hash2 ← hash(password, salt) // Generates a hash a one-way key derivation function,
        ↳ and the extracted salt
43         return hash2 = hash // We will assume a secure implementation of the comparison
44     return false
45
46 function authCheck(cookie, username):
47     if cookie ≠ ⊥:
48         return sessionCheck(cookie.getValue(), username)
49     throw AuthException()
50
51 function sessionCheck(sessionIDprovided, username)
52     if sessionIDprovided ≠ ⊥:
53         sessionID ← storageEngine.getSession(username)
54         if sessionID ≠ ⊥:
55             return sessionIDprovided = sessionID

```

Listing D.2: Session Validation in IAM Provider

```

1 State:
2     SESSION_LIFETIME ← fetchConfig(SESSION_LIFETIME)
3     roles ← {}
4     passwords ← {}
5     sessions ← {}
6     owned ← {}
7
8 function cacheSession(username):
9     sessionID ← generateID()
10    sessions[username] ← sessionID
11    schedule(expireSession(sessionID, username), SESSION_LIFETIME)
12    return createCookie(sessionID)
13
14 function getSession(username):
15     return sessions[username]
16
17 function userExists(username):
18     return roles[username] ≠ ⊥ and passwords[username] ≠ ⊥
19
20 function saveUser(username, password, role)
21     roles[username] ← role
22     passwords[username] ← password
23     owned[username] ← {}
24
25 function getPassword(username)

```

```

26     return passwords[username]
27
28 function getRole(username)
29     return roles[username]
30
31 function checkIfOwnsAny(username)
32     return #owned[username] > 0
33
34 function deleteUser(username)
35     roles[username] ← ⊥
36     passwords[username] ← ⊥
37     owned[username] ← ⊥
38     sessions[username] ← ⊥
39
40 function expireSession(sessionID, username):
41     if sessions[username] = sessionID:
42         sessions[username] ← ⊥

```

Listing D.3: User Data Storage

```

1  State:
2  SEPARATOR ← fetchConfig(SEPARATOR)
3  LOCK_LIFETIME ← fetchConfig(LOCK_LIFETIME)
4  LOCK_TRIES ← fetchConfig(LOCK_TRIES)
5  LOCK_SLEEP ← fetchConfig(LOCK_SLEEP)
6  locks ← {}
7
8  function acquireUserLock(username):
9      key = "UL" || SEPARATOR || username
10     lockID ← generateID()
11     tries ← LOCK_TRIES
12     while tries > 0:
13         tries--
14         try:
15             locks[key] ← lockID
16             expire(releaseUserLock(username, lockID), LOCK_LIFETIME)
17             return lockID
18         catch Exception:
19             sleep(LOCK_SLEEP)
20     throw TooManyLockRetriesException()
21
22 function releaseUserLock(username, lockID)
23     key = "UL" || SEPARATOR || username
24     if locks[key] = lockID:
25         locks[key] ← ⊥
26
27 function acquireTriplestoreLock(username, triplestoreID):
28     triplestoreLockKey ← "LT" || triplestoreID
29     userTriplestoreKey ← "UTL" || username || SEPARATOR || triplestoreID || SEPARATOR
30         ↳ || lockID
31     lockID ← generateID()
32     tries ← LOCK_TRIES
33     while tries > 0:
34         tries--
35         try:
36             locks[triplestoreLockKey] ← lockID
37             locks[userTriplestoreKey] ← 0
38             expire(releaseTriplestoreLock(username, triplestoreID, lockID), LOCK_LIFETIME)
39             ↳

```

```

38     return lockID
39     catch Exception:
40         sleep(LOCK_SLEEP)
41     throw TooManyLockRetriesException()
42
43 function checkIfTriplestoreLockExists(triplestoreID, lockID)
44     triplestoreLockKey ← "LT" || triplestoreID
45     return locks[triplestoreLockKey] = lockID
46
47 function releaseTriplestoreLock(username, triplestoreID, lockID):
48     triplestoreLockKey ← "LT" || triplestoreID
49     userTriplestoreKey ← "UTL" || username || SEPARATOR || triplestoreID || SEPARATOR
        ↳ || lockID
50     if locks[triplestoreLockKey] = lockID:
51         locks[triplestoreLockKey] ← ⊥
52         locks[userTriplestoreKey] ← ⊥
53
54 function deleteUserTriplestoreLock(username, triplestoreID):
55     pattern ← "UTL" || username || SEPARATOR || triplestoreID || SEPARATOR || "*"
56     deleteUserLocks(username, pattern)
57
58 function deleteAllUserLocks(username):
59     pattern ← "UTL" || username || SEPARATOR || "*"
60     deleteUserLocks(username, pattern)
61
62 function deleteUserLocks(username, pattern):
63     lockKeys ← {}
64     foreach (key, _) in locks:
65         if patternMatch(key, pattern) = true:
66             values = key.split(SEPARATOR)
67             releaseTriplestoreLock(username, values[1], values[2])

```

Listing D.4: Locks Client

```

1 State:
2     storageEngine ← createStorageEngine()
3     locksClient ← createLocksClient()
4
5 function createAccessToken(cookie, triplestoreID, target):
6     try:
7         authCheck(cookie, target)
8         if storageEngine.triplestoreAccessListExists(triplestoreID) = false:
9             return ⊥
10        role ← storageEngine.getRole(target)
11        if role ≠ "ADMIN" and storageEngine.checkIfUserHasReadAccess(target, triplestoreID) =
            ↳ false and storageEngine.checkIfUserHasWriteAccess(target, triplestoreID) =
            ↳ false and storageEngine.checkIfOwns(target, triplestoreID) = false:
12            return ⊥
13        return storageEngine.saveToken(target, cookie.getValue(), triplestoreID)
14    catch Exception:
15        return ⊥
16
17 function deleteAccessToken(cookie, triplestoreID, headers):
18     try:
19         (tokenID, username, sessionID, triplestoreIDtoken, lockID) ← fetchAccessToken(headers)
20         sessionCheck(sessionID, username)
21         if storageEngine.triplestoreAccessListExists(triplestoreID) = false:
22             return ⊥
23         if triplestoretoken ≠ triplestoreID:

```

```

24         return  $\perp$ 
25         storageEngine.deleteAccessToken(tokenID, lockID, username, triplestoreIDtoken)
26         return "SUCCESS"
27     catch Exception:
28         return  $\perp$ 
29
30 function checkIfActive(headers)
31     try:
32         (tokenID, username, sessionID, triplestoreIDtoken, lockID)  $\leftarrow$  fetchAccessToken(headers)
33         sessionCheck(sessionID, username)
34         return "OK"
35     catch Exception:
36         return  $\perp$ 
37
38 function checkReadAccess(triplestoreID, headers):
39     try:
40         (tokenID, username, sessionID, triplestoreIDtoken, lockID)  $\leftarrow$  fetchAccessToken(headers)
41         sessionCheck(sessionID, username)
42         if storageEngine.triplestoreAccessListExists(triplestoreID) = false:
43             return  $\perp$ 
44         if triplestoretoken  $\neq$  triplestoreID:
45             return  $\perp$ 
46         if storageEngine.checkIfUserHasReadAccess(username, triplestoreID) = false and
47              $\hookrightarrow$  storageEngine.checkIfUserHasWriteAccess(username, triplestoreID) = false and
48              $\hookrightarrow$  storageEngine.checkIfOwns(username, triplestoreID):
49             return  $\perp$ 
50         return "OK"
51     catch Exception:
52         return  $\perp$ 
53
54 function checkWriteAccess(triplestoreID, headers)
55     try:
56         (tokenID, username, sessionID, triplestoreIDtoken, lockID)  $\leftarrow$  fetchAccessToken(headers)
57         sessionCheck(sessionID, username)
58         if storageEngine.triplestoreAccessListExists(triplestoreID) = false:
59             return  $\perp$ 
60         if triplestoretoken  $\neq$  triplestoreID:
61             return  $\perp$ 
62         if lockID =  $\perp$  or locksClient.checkIfTriplestoreLockExists(triplestoreID, lockID) =
63              $\hookrightarrow$  false:
64             return  $\perp$ 
65         if storageEngine.checkIfUserHasWriteAccess(username, triplestoreID) = false and
66              $\hookrightarrow$  storageEngine.checkIfOwns(username, triplestoreID) = false:
67             return  $\perp$ 
68         return "OK"
69     catch Exception:
70         return  $\perp$ 
71
72 function checkOwnerAccess(triplestoreID, headers)
73     try:
74         (tokenID, username, sessionID, triplestoreIDtoken, lockID)  $\leftarrow$  fetchAccessToken(headers)
75         sessionCheck(sessionID, username)
76         if storageEngine.triplestoreAccessListExists(triplestoreID) = false:
77             return  $\perp$ 
78         if triplestoretoken  $\neq$  triplestoreID:
79             return  $\perp$ 
80         if storageEngine.checkIfOwns(username, triplestoreID) = false
81             return  $\perp$ 
82         if lockID =  $\perp$  or locksClient.checkIfTriplestoreLockExists(triplestoreID, lockID) =
83              $\hookrightarrow$  false:

```

```

79     return  $\perp$ 
80     return "OK"
81 catch Exception:
82     return  $\perp$ 
83
84 function acquireTriplestoreLock(cookie, triplestoreID, headers)
85     try:
86         (tokenID, username, sessionID, triplestoreIDtoken, lockID)  $\leftarrow$  fetchAccessToken(headers)
87         authCheck(cookie, username)
88         if cookie.getValue()  $\neq$  sessionID:
89             return  $\perp$ 
90         if storageEngine.triplestoreAccessListExists(triplestoreID) = false:
91             return  $\perp$ 
92         if triplestoretoken  $\neq$  triplestoreID:
93             return  $\perp$ 
94         role  $\leftarrow$  storageEngine.getRole(username)
95         if role  $\neq$  "ADMIN" and storageEngine.checkIfUserHasWriteAccess(username, triplestoreID)
96              $\hookrightarrow$  = false and storageEngine.checkIfOwns(username, triplestoreID) = false:
97             return  $\perp$ 
98         lockID  $\leftarrow$  locksClient.acquireTriplestoreLock(username, triplestoreID)
99         storageEngine.addLockToToken(tokenID, lockID)
100         return lockID
101     catch Exception:
102         return  $\perp$ 
103
104 function releaseTriplestoreLock(cookie, triplestoreID, headers)
105     try:
106         (tokenID, username, sessionID, triplestoreIDtoken, lockID)  $\leftarrow$  fetchAccessToken(headers)
107         authCheck(cookie, username)
108         if cookie.getValue()  $\neq$  sessionID:
109             return  $\perp$ 
110         if storageEngine.triplestoreAccessListExists(triplestoreID) = false:
111             return  $\perp$ 
112         if triplestoretoken  $\neq$  triplestoreID:
113             return  $\perp$ 
114         if lockID =  $\perp$ :
115             return  $\perp$ 
116         storageEngine.deleteLockFromToken(tokenID)
117         locksClient.releaseTriplestoreLock(username, triplestoreID, lockID)
118         return "SUCCESS"
119     catch Exception:
120         return  $\perp$ 
121
122 function fetchAccessToken(headers):
123     tokenID  $\leftarrow$  extractAccessToken(headers) // Extracts the value of the access token from
124      $\hookrightarrow$  the headers
125     if tokenID =  $\perp$ 
126         throw TokenException
127     token  $\leftarrow$  storageEngine.getToken(tokenID)
128     if token =  $\perp$ : // Expired or invalid
129         throw TokenException
130     return (tokenID, token.username, token.sessionID, token.triplestoreID, token.lockID)

```

Listing D.5: Permission Checks

```

1 State:
2     locksClient  $\leftarrow$  createLocksClient()
3     SESSION_LIFETIME  $\leftarrow$  fetchConfig(SESSION_LIFETIME)
4     ACCESS_TOKEN_LIFETIME  $\leftarrow$  fetchConfig(ACCESS_TOKEN_LIFETIME)

```

```
5  roles ← {}
6  passwords ← {}
7  sessions ← {}
8  owned ← {}
9  acls ← {}
10 pendingAccessRequests ← {}
11 accessTokens ← {}
12
13 function checkIfOwns(username, triplestoreID):
14     return triplestoreID ∈ owned[username]
15
16 function getOwner(triplestoreID):
17     (owner, read, write, acl) ← acls[triplestore]
18     return owner
19
20 function checkIfUserHasReadAccess(username, triplestoreID):
21     (owner, read, write, acl) ← acls[triplestore]
22     return username ∈ read
23
24 function checkIfUserHasWriteAccess(username, triplestoreID):
25     (owner, read, write, acl) ← acls[triplestore]
26     return username ∈ write
27
28 function triplestoreAccessListExists(triplestoreID):
29     return acls[triplestore] ≠ ⊥
30
31 function saveTriplestoreAccessList(triplestoreID, owner, read, write):
32     owned[owner] ← owned[owner] ∪ { triplestoreID }
33     acls[triplestore] ← (owner, read, write, {})
34
35 function updateTriplestoreOwner(triplestoreID, currentOwner, newOwner):
36     owned[currentOwner] ← owned[currentOwner] / { triplestoreID }
37     owned[newOwner] ← owned[currentOwner] ∪ { triplestoreID }
38     (owner, read, write, acl) ← acls[triplestore]
39     acls[triplestore] ← (newOwner, read / { newOwner }, write / { newOwner }, acl)
40
41 function deleteTriplestoreAccessList(triplestoreID):
42     (owner, read, write, acl) ← acls[triplestore]
43     pendingAccessRequests ← pendingAccessRequests / acl
44     owned[owner] ← owned[owner] / { triplestoreID }
45     acls[triplestore] ← ⊥
46
47 function getTriplestores(username, write, read, owns):
48     if write = false and read = false and owns = false:
49         return {}
50     res ← owned[username]
51     if write = false and read = false and owns = true:
52         return res
53     foreach (triplestoreID, (owner, read, write, acl)) in acls:
54         if write = true:
55             if username ∈ write:
56                 res ← res ∪ { triplestoreID }
57         else:
58             if username ∈ read:
59                 res ← res ∪ { triplestoreID }
60     return res
61
62 function getUserWithReadAccess(triplestoreID):
63     (owner, read, write, acl) ← acls[triplestore]
64     return read
```

```

65
66 function getUserWithWriteAccess(triplestoreID):
67     (owner, read, write, acrl) ← acls[triplestore]
68     return write
69
70 function saveToken(username, session, triplestoreID):
71     tokenID ← generateID()
72     accessTokens[tokenID] ← (tokenID, username, sessionID, triplestoreID, ⊥)
73     schedule(deleteAccessToken(tokenID, username, triplestoreID, ⊥),
74               ↪ ACCESS_TOKEN_LIFETIME)
75     return tokenID
76
77 function addLockToToken(tokenID, lockIDnew):
78     (tokenID, username, sessionID, triplestoreID, lockID) ← accessTokens[tokenID]
79     accessTokens[tokenID] ← (tokenID, username, sessionID, triplestoreID, lockIDnew)
80
81 function deleteLockFromToken(tokenID):
82     (tokenID, username, sessionID, triplestoreID, lockID) ← accessTokens[tokenID]
83     accessTokens[tokenID] ← (tokenID, username, sessionID, triplestoreID, ⊥)
84
85 function getToken(tokenID):
86     return accessTokens[tokenID]
87
88 function deleteAccessToken(tokenID, username, triplestoreID, lockID):
89     if lockID ≠ ⊥:
90         locksClient.releaseTriplestoreLock(username, triplestoreID, lockID)
91     accessTokens[tokenID] ← ⊥

```

Listing D.6: ACL & Access Tokens Data Storage

```

1 function createTriplestore(cookie, triplestoreID, issuer):
2     if cookie = ⊥
3         return ⊥
4     result ← IAM_Provider.createTriplestoreAccessList(cookie, triplestoreID, issuer)
5     if response = ⊥:
6         IAM_Provider.deleteAccessToken(cookie, triplestoreID, accessToken)
7         return ⊥
8     scheme ← createEncryptionScheme()
9     response ← Vault.saveSecrets(triplestoreID, scheme.getSecrets(), accessToken)
10    if response = ⊥:
11        reponse2 ← IAM_Provider.deleteTriplestoreAccessList(cookie, triplestoreID, issuer,
12                                                              ↪ accessToken)
13        if reponse2 = ⊥
14            IAM_Provider.deleteAccessToken(cookie, triplestoreID, accessToken)
15            return ⊥
16    IAM_Provider.deleteAccessToken(cookie, triplestoreID, accessToken)
17    return response
18
19 function deleteTriplestore(cookie, triplestoreID, issuer):
20     if cookie = ⊥
21         return ⊥
22     accessToken ← IAM_Provider.createAccessToken(cookie, issuer, triplestoreID)
23     if accessToken = ⊥:
24         return ⊥
25     response ← IAM_Provider.acquireTriplestoreLock(cookie, triplestoreID, accessToken)
26     if response = ⊥:
27         IAM_Provider.deleteAccessToken(cookie, triplestoreID, accessToken)
28         return ⊥
29     reponse ← Vault.deleteSecrets(triplestoreID, accessToken)

```

```
29   if response =  $\perp$ :
30       IAM_Provider.deleteAccessToken(cookie, triplestoreID, accessToken)
31       return  $\perp$ 
32   reponse  $\leftarrow$  IAM_Provider.deleteTriplestoreAccessList(cookie, triplestoreID, issuer,
33        $\hookrightarrow$  accessToken)
34   if response =  $\perp$ :
35       IAM_Provider.deleteAccessToken(cookie, triplestoreID, accessToken)
36       return  $\perp$ 
37   return response
38
39 function listTriplestores(cookie, issuer, write, read, owns):
40     if cookie =  $\perp$ 
41         return  $\perp$ 
42     return IAM_Provider.listTriplestores(cookie, issuer, write, read, owns)
43
44 function updateTriplestoreOwner(cookie, triplestoreID, issuer, target):
45     if cookie =  $\perp$ 
46         return  $\perp$ 
47     accessToken  $\leftarrow$  IAM_Provider.createAccessToken(cookie, issuer, triplestoreID)
48     if accessToken =  $\perp$ :
49         return  $\perp$ 
50     response  $\leftarrow$  IAM_Provider.acquireTriplestoreLock(cookie, triplestoreID, accessToken)
51     if response =  $\perp$ :
52         IAM_Provider.deleteAccessToken(cookie, triplestoreID, accessToken)
53         return  $\perp$ 
54     response  $\leftarrow$  IAM_Provider.updateTriplestoreOwner(cookie, triplestoreID, target, accessToken)
55     IAM_Provider.deleteAccessToken(cookie, triplestoreID, accessToken)
56     return response
57
58 function requestAccess(cookie, triplestoreID, issuer, write):
59     if cookie =  $\perp$ 
60         return  $\perp$ 
61     return IAM_Provider.issueAccessRequest(cookie, triplestoreID, issuer, write)
62
63 function listPendingAccessRequests(cookie, triplestoreID, issuer):
64     if cookie =  $\perp$ 
65         return  $\perp$ 
66     accessToken  $\leftarrow$  IAM_Provider.createAccessToken(cookie, issuer, triplestoreID)
67     if accessToken =  $\perp$ :
68         return  $\perp$ 
69     response  $\leftarrow$  IAM_Provider.listPendingAccessRequests(cookie, triplestoreID, accessToken)
70     IAM_Provider.deleteAccessToken(cookie, triplestoreID, accessToken)
71     return response
72
73 function issueAccessDecision(cookie, triplestoreID, issuer, requestID, accept):
74     if cookie =  $\perp$ 
75         return  $\perp$ 
76     accessToken  $\leftarrow$  IAM_Provider.createAccessToken(cookie, issuer, triplestoreID)
77     if accessToken =  $\perp$ :
78         return  $\perp$ 
79     response  $\leftarrow$  IAM_Provider.processAccessRequest(cookie, triplestoreID, requestID, accept,
80        $\hookrightarrow$  accessToken)
81     IAM_Provider.deleteAccessToken(cookie, triplestoreID, accessToken)
82     return response
83
84 function grantAccess(cookie, triplestoreID, issuer, target, write):
85     if cookie =  $\perp$ 
86         return  $\perp$ 
87     accessToken  $\leftarrow$  IAM_Provider.createAccessToken(cookie, issuer, triplestoreID)
88     if accessToken =  $\perp$ :
```

```

87     return  $\perp$ 
88 response  $\leftarrow$  IAM_Provider.acquireTriplestoreLock(cookie, triplestoreID, accessToken)
89 if response =  $\perp$ :
90     IAM_Provider.deleteAccessToken(cookie, triplestoreID, accessToken)
91     return  $\perp$ 
92 response  $\leftarrow$  IAM_Provider.grantTriplestoreAccess(cookie, triplestoreID, target, write,
93      $\hookrightarrow$  accessToken)
94 IAM_Provider.deleteAccessToken(cookie, triplestoreID, accessToken)
95 return response
96
97 function revokeAccess(cookie, triplestoreID, issuer, target, write):
98     if cookie =  $\perp$ 
99         return  $\perp$ 
100     accessToken  $\leftarrow$  IAM_Provider.createAccessToken(cookie, issuer, triplestoreID)
101     if accessToken =  $\perp$ :
102         return  $\perp$ 
103     response = acquireTriplestoreLock(cookie, triplestoreID, accessToken)
104     if response =  $\perp$ :
105         IAM_Provider.deleteAccessToken(cookie, triplestoreID, accessToken)
106         return  $\perp$ 
107     response  $\leftarrow$  IAM_Provider.revokeTriplestoreAccess(cookie, triplestoreID, target, write,
108          $\hookrightarrow$  accessToken)
109     IAM_Provider.deleteAccessToken(cookie, triplestoreID, accessToken)
110     return response
111
112 function listUsersWithAccess(cookie, triplestoreID, issuer, write):
113     if cookie =  $\perp$ 
114         return  $\perp$ 
115     accessToken  $\leftarrow$  IAM_Provider.createAccessToken(cookie, issuer, triplestoreID)
116     if accessToken =  $\perp$ :
117         return  $\perp$ 
118     response  $\leftarrow$  IAM_Provider.listUsersWithAccess(cookie, triplestoreID, write, accessToken)
119     IAM_Provider.deleteAccessToken(cookie, triplestoreID, accessToken)
120     return response

```

Listing D.7: Access Management in Client

```

1 State:
2     storageEngine  $\leftarrow$  createStorageEngine()
3     locksClient  $\leftarrow$  createLocksClient()
4
5 function createTriplestoreAccessList(cookie, issuer, triplestoreID):
6     try:
7         authCheck(cookie, issuer)
8         role  $\leftarrow$  storageEngine.getRole(issuer)
9         if role = "BASIC":
10             return  $\perp$ 
11         lockID  $\leftarrow$  locksClient.acquireTriplestoreLock(issuer, triplestoreID)
12         if storageEngine.triplestoreAccessListExists(triplestoreID) = true:
13             locksClient.releaseTriplestoreLock(issuer, triplestoreID, lockID)
14             return  $\perp$ 
15         storageEngine.saveTriplestoreAccessList(triplestoreID, issuer, {}, {})
16         locksClient.releaseTriplestoreLock(issuer, triplestoreID, lockID)
17         return "SUCCESS"
18     catch Exception:
19         return  $\perp$ 
20
21 function getTriplestores(cookie, issuer, write, read, owns):
22     try:

```

```

23     authCheck(cookie, issuer)
24     return storageEngine.getTriplestores(issuer, write, read, owns)
25 catch Exception:
26     return ⊥
27
28 function getUsersWithAccess(cookie, triplestoreID, write, headers):
29     try:
30         (tokenID, username, sessionID, triplestoreIDtoken, lockID) ← fetchAccessToken(headers)
31         authCheck(cookie, username)
32         if cookie.getValue() ≠ sessionID:
33             return ⊥
34         if storageEngine.triplestoreAccessListExists(triplestoreID) = false:
35             return ⊥
36         if triplestoretoken ≠ triplestoreID:
37             return ⊥
38         role ← storageEngine.getRole(issuer)
39         isOwner ← storageEngine.checkIfOwns(issuer, triplestoreID)
40         if role = "BASIC" or (role = "PRIVILEGED" and isOwner = false):
41             return ⊥
42         usersread ← storageEngine.getUserWithReadAccess(triplestoreID)
43         userswrite ← {}
44         if write = true:
45             userswrite ← storageEngine.getUserWithWriteAccess(triplestoreID)
46             usersread ← usersread ∪ userswrite
47         if isOwner = false:
48             return {
49                 "owner": storageEngine.getOwner(triplestoreID),
50                 "write": userswrite,
51                 "read": usersread
52             }
53         return {
54             "owner": issuer,
55             "write": userswrite,
56             "read": usersread
57         }
58     catch Exception:
59         return ⊥
60
61 function changeTriplestoreOwner(cookie, triplestoreID, target, headers):
62     try:
63         (tokenID, username, sessionID, triplestoreIDtoken, lockID) ← fetchAccessToken(headers)
64         authCheck(cookie, username)
65         if cookie.getValue() ≠ sessionID:
66             return ⊥
67         if storageEngine.triplestoreAccessListExists(triplestoreID) = false:
68             return ⊥
69         if triplestoretoken ≠ triplestoreID:
70             return ⊥
71         if storageEngine.userExists(target) = false:
72             return ⊥
73         roleissuer ← storageEngine.getRole(issuer)
74         if roleissuer = "BASIC" or (roleissuer = "PRIVILEGED" and storageEngine.checkIfOwns(
75             ↪ target, triplestoreID) = false):
76             return ⊥
77         owner ← storageEngine.getOwner(triplestoreID)
78         if target = owner:
79             return ⊥
80         role ← storageEngine.getRole(target)
81         if role = "BASIC":
82             return ⊥

```

```

82     if lockID =  $\perp$  or locksClient.checkIfTriplestoreLockExists(triplestoreID, lockID) =
83          $\hookrightarrow$  false:
84         return  $\perp$ 
85     storageEngine.updateTriplestoreOwner(triplestoreID, owner, target)
86     return "OK"
87 catch Exception:
88     return  $\perp$ 
89
90 function deleteTriplestoreAccessList(cookie, triplestoreID, headers):
91     try:
92         (tokenID, username, sessionID, triplestoreIDtoken, lockID)  $\leftarrow$  fetchAccessToken(headers)
93         authCheck(cookie, username)
94         if cookie.getValue()  $\neq$  sessionID:
95             return  $\perp$ 
96         if storageEngine.triplestoreAccessListExists(triplestoreID) = false:
97             return  $\perp$ 
98         if triplestoretoken  $\neq$  triplestoreID:
99             return  $\perp$ 
100         role = storageEngine.getRole(username)
101         if role = "BASIC" or (role = "PRIVILEGED" and storageEngine.checkIfOwns(username,
102              $\hookrightarrow$  triplestoreID) = false):
103             return  $\perp$ 
104         if lockID =  $\perp$  or locksClient.checkIfTriplestoreLockExists(triplestoreID, lockID) =
105              $\hookrightarrow$  false:
106             return  $\perp$ 
107         storageEngine.deleteTriplestoreAccessList(owner, triplestoreID)
108         storageEngine.deleteAccessToken(tokenID, lockID, username, triplestoreIDtoken)
109         return "OK"
110     catch Exception:
111         return  $\perp$ 
112
113 function grantTriplestoreAccess(cookie, triplestoreID, target, write, headers):
114     try:
115         (tokenID, username, sessionID, triplestoreIDtoken, lockID)  $\leftarrow$  fetchAccessToken(headers)
116         authCheck(cookie, username)
117         if cookie.getValue()  $\neq$  sessionID:
118             return  $\perp$ 
119         if storageEngine.triplestoreAccessListExists(triplestoreID) = false:
120             return  $\perp$ 
121         if triplestoretoken  $\neq$  triplestoreID:
122             return  $\perp$ 
123         if storageEngine.userExists(target) = false:
124             return  $\perp$ 
125         roleissuer  $\leftarrow$  storageEngine.getRole(issuer)
126         if issuer = target:
127             return  $\perp$ 
128         elif roleissuer = "BASIC" or (roleissuer = "PRIVILEGED" and storageEngine.checkIfOwns(
129              $\hookrightarrow$  issuer, triplestoreID) = false):
130             return  $\perp$ 
131         if lockID =  $\perp$  or locksClient.checkIfTriplestoreLockExists(triplestoreID, lockID) =
132              $\hookrightarrow$  false:
133             return  $\perp$ 
134         if storageEngine.getRole(target) = "ADMIN" or storageEngine.checkIfOwns(target,
135              $\hookrightarrow$  triplestoreID) = true or (storageEngine.checkIfUserHasWriteAccess(target,
136              $\hookrightarrow$  triplestoreID) = write and storageEngine.checkIfUserHasReadAccess(target,
137              $\hookrightarrow$  triplestoreID) = true):
138             return  $\perp$ 
139         storageEngine.grantAccess(triplestoreID, target, write)
140         return "OK"
141     catch Exception:

```

```

134     return  $\perp$ 
135
136 function revokeTriplestoreAccess(cookie, triplestoreID, target, write, headers):
137     try:
138         (tokenID, username, sessionID, triplestoreIDtoken, lockID)  $\leftarrow$  fetchAccessToken(headers)
139         authCheck(cookie, username)
140         if cookie.getValue()  $\neq$  sessionID:
141             return  $\perp$ 
142         if storageEngine.triplestoreAccessListExists(triplestoreID) = false:
143             return  $\perp$ 
144         if triplestoretoken  $\neq$  triplestoreID:
145             return  $\perp$ 
146         if storageEngine.userExists(target) = false:
147             return  $\perp$ 
148         roleissuer  $\leftarrow$  storageEngine.getRole(issuer)
149         if issuer  $\neq$  target:
150             if roleissuer = "BASIC" or (roleissuer = "PRIVILEGED" and storageEngine.checkIfOwns(
151                  $\hookrightarrow$  issuer, triplestoreID) = false):
152                 return  $\perp$ 
153             if lockID =  $\perp$  or locksClient.checkIfTriplestoreLockExists(triplestoreID, lockID) =
154                  $\hookrightarrow$  false:
155                 return  $\perp$ 
156             if storageEngine.checkIfOwns(target, triplestoreID) = true:
157                 return  $\perp$ 
158             locksClient.deleteUserTriplestoreLock(target, triplestoreID)
159             if roleissuer = "ADMIN":
160                 storageEngine.revokeAccess(triplestoreID, target, write)
161             return "OK"
162         catch Exception:
163             return  $\perp$ 
164
165 function issueAccessRequest(cookie, triplestoreID, target, write):
166     try:
167         authCheck(cookie, target)
168         if storageEngine.triplestoreAccessListExists(triplestoreID) = false:
169             return  $\perp$ 
170         role  $\leftarrow$  storageEngine.getRole(target)
171         if role = "ADMIN" or (role = "PRIVILEGED" and storageEngine.checkIfOwns(target,
172              $\hookrightarrow$  triplestoreID) = true)
173             return  $\perp$ 
174         if storageEngine.checkIfUserHasWriteAccess(target, triplestoreID) = write and
175              $\hookrightarrow$  storageEngine.checkIfUserHasReadAccess(target, triplestoreID):
176             return  $\perp$ 
177         storageEngine.saveAccessRequest(triplestoreID, target, write)
178         return "OK"
179     catch Exception:
180         return  $\perp$ 
181
182 function getPendingAccessRequests(cookie, triplestoreID, headers):
183     try:
184         (tokenID, username, sessionID, triplestoreIDtoken, lockID)  $\leftarrow$  fetchAccessToken(headers)
185         authCheck(cookie, username)
186         if cookie.getValue()  $\neq$  sessionID:
187             return  $\perp$ 
188         if storageEngine.triplestoreAccessListExists(triplestoreID) = false:
189             return  $\perp$ 
190         if triplestoretoken  $\neq$  triplestoreID:
191             return  $\perp$ 
192         role = storageEngine.getRole(issuer)

```

```

189     if role = "BASIC" or (role = "PRIVILEGED" and storageEngine.checkIfOwns(issuer,
190         ↪ triplestoreID) = false):
191         return ⊥
192     return "OK"
193 catch Exception:
194     return ⊥
195
196 function processAccessRequest(cookie, triplestoreID, requestID, decision, headers):
197     try:
198         (tokenID, username, sessionID, triplestoreIDtoken, lockID) ← fetchAccessToken(headers)
199         authCheck(cookie, username)
200         if cookie.getValue() ≠ sessionID:
201             return ⊥
202         if storageEngine.triplestoreAccessListExists(triplestoreID) = false:
203             return ⊥
204         if triplestoretoken ≠ triplestoreID:
205             return ⊥
206         role = storageEngine.getRole(issuer)
207         if role = "BASIC" or (role = "PRIVILEGED" and storageEngine.checkIfOwns(issuer,
208             ↪ triplestoreID) = false):
209             return "OK"
210         accessRequest ← storageEngine.getPendingAccessRequest(requestID)
211         if accessRequest = ⊥:
212             return ⊥
213         if storageEngine.userExists(accessRequest.username) = false:
214             return ⊥
215         if decision = true:
216             storageEngine.grantAccess(triplestoreID, username, accessRequest.write)
217             storageEngine.deleteAccessRequest(triplestoreID, requestID)
218             return "OK"
219         catch Exception:
220             return ⊥

```

Listing D.8: Access Management in IAM Provider

```

1 State:
2     locksClient ← createLocksClient()
3     SESSION_LIFETIME ← fetchConfig(SESSION_LIFETIME)
4     ACCESS_TOKEN_LIFETIME ← fetchConfig(ACCESS_TOKEN_LIFETIME)
5     roles ← {}
6     passwords ← {}
7     sessions ← {}
8     owned ← {}
9     acls ← {}
10    pendingAccessRequests ← {}
11    accessTokens ← {}
12
13 function revokeAccess(triplestoreID, username, write):
14     (owner, read, write, acl) ← acls[triplestore]
15     if write = false:
16         acls[triplestore] ← (owner, read / { username }, write / { username }, acl)
17     else:
18         acls[triplestore] ← (owner, read, write / { username }, acl)
19
20 function grantAccess(triplestoreID, username, write):
21     (owner, read, write, acl) ← acls[triplestore]
22     if write:
23         acls[triplestore] ← (owner, read ∪ { username }, write ∪ { username }, acl)
24     else:

```

```

25     acls[triplestore] ← (owner, read ∪ { username }, write, acrl)
26
27 function saveAccessRequest(triplestoreID, username, write):
28     requestID ← generateID()
29     (owner, read, write, acrl) ← acls[triplestore]
30     acls[triplestore] ← (owner, read, write, acrl ∪ { (username, write) })
31     pendingAccessRequests[requestID] ← (username, write)
32
33 function getPendingAccessRequest(requestID):
34     return pendingAccessRequests[requestID]
35
36 function getPendingAccessRequests(triplestoreID):
37     (owner, read, write, acrl) ← acls[triplestore]
38     return acrl
39
40 function deleteAccessRequest(triplestoreID, requestID):
41     (owner, read, write, acrl) ← acls[triplestore]
42     acls[triplestore] ← (owner, read, write, acrl / { (username, write) })
43     pendingAccessRequests[requestID] ← ⊥

```

Listing D.9: Storage of Access Requests

```

1 State:
2     secrets ← {}
3
4 function saveSecrets(triplestoreID, secrets, accessToken):
5     tokenID ← extractAccessToken(headers)
6     if tokenID = ⊥ or IAM_Provider.checkOwnerAccess(triplestoreID, accessToken) = false:
7         return ⊥
8     if secrets[triplestoreID] ≠ ⊥:
9         return ⊥
10    secrets[triplestoreID] ← secrets
11    return "SUCCESS"
12
13 function getSecrets(triplestoreID, accessToken):
14    tokenID ← extractAccessToken(headers)
15    if tokenID = ⊥ or IAM_Provider.checkReadAccess(triplestoreID, accessToken) = false:
16        return ⊥
17    return secrets[triplestoreID]
18
19 function deleteSecrets(triplestoreID, accessToken):
20    tokenID ← extractAccessToken(headers)
21    if tokenID = ⊥ or IAM_Provider.checkOwnerAccess(triplestoreID, accessToken) = false:
22        return ⊥
23    if secrets[triplestoreID] = ⊥:
24        return ⊥
25    secrets[triplestoreID] ← ⊥
26    return "SUCCESS"

```

Listing D.10: Secret Management in Vault

```

1 function issueUpgradeRequest(cookie, username):
2     if cookie = ⊥:
3         return ⊥
4     return IAM_Provider.issueGrantRoleRequest(cookie, username, username, "PRIVILEGED")
5
6 function issueDowngradeRequest(cookie, username):
7     if cookie = ⊥:

```

```

8     return  $\perp$ 
9     return IAM_Provider.issueGrantRoleRequest(cookie, username, username, "BASIC")
10
11 function listPendingRoleRequests(cookie, username)
12     if cookie =  $\perp$ :
13         return  $\perp$ 
14     return IAM_Provider.getPendingRoleRequests(cookie, username)
15
16 function issueRoleDecision(cookie, username, requestID, target, decision)
17     if cookie =  $\perp$ :
18         return  $\perp$ 
19     return IAM_Provider.processRoleRequest(cookie, username, requestID, target, decision)

```

Listing D.11: Role Management in Client

```

1 State:
2     storageEngine  $\leftarrow$  createStorageEngine()
3     locksClient  $\leftarrow$  createLocksClient()
4
5 function issueGrantRoleRequest(cookie, username, issuer, role)
6     try:
7         authCheck(cookie, issuer)
8         if storageEngine.userExists(username) = false:
9             return  $\perp$ 
10        roleissuer  $\leftarrow$  storageEngine.getRole(issuer)
11        if storageEngine.getRole(username) = role:
12            return  $\perp$ 
13        if roleissuer = "BASIC" or roleissuer = "PRIVILEGED":
14            if username = issuer:
15                storageEngine.saveRoleRequest(username, role)
16                return "SUCCESS"
17            else:
18                return  $\perp$ 
19        else:
20            lockID  $\leftarrow$  locksClient.acquireUserLock(username)
21            storageEngine.setRole(username, role)
22            locksClient.releaseUserLock(username, lockID)
23            return "SUCCESS"
24        catch Exception:
25            return  $\perp$ 
26
27 function getPendingRoleRequests(cookie, username)
28     try:
29         authCheck(cookie, username)
30         role  $\leftarrow$  storageEngine.getRole(username)
31         if role  $\neq$  "ADMIN":
32             return  $\perp$ 
33         return storageEngine.getPendingRoleRequests()
34     catch Exception:
35         return  $\perp$ 
36
37 function processRoleRequest(cookie, username, requestID, target, decision):
38     try:
39         authCheck(cookie, username)
40         roleissuer  $\leftarrow$  storageEngine.getRole(username)
41         if roleissuer  $\neq$  "ADMIN":
42             return  $\perp$ 
43         if storageEngine.userExists(target) = false:
44             return  $\perp$ 

```

```

45     lockID ← locksClient.acquireUserLock(target)
46     (usernamerequest, role) ← storageEngine.getPendingRoleRequest(requestID)
47     if request = ⊥:
48         locksClient.releaseUserLock(target, lockID)
49         return ⊥
50     if target ≠ usernamerequest:
51         locksClient.releaseUserLock(target, lockID)
52         return ⊥
53     if decision = true:
54         if storageEngine.getRole(target) = role:
55             storageEngine.deleteRoleRequest(requestID)
56             locksClient.releaseUserLock(target, lockID)
57             return ⊥
58         storageEngine.setRole(target, role)
59     storageEngine.deleteRoleRequest(requestID)
60     locksClient.releaseUserLock(target, lockID)
61     return "SUCCESS"
62 catch Exception:
63     return ⊥

```

Listing D.12: Role Management in IAM Provider

```

1  State:
2     SESSION_LIFETIME ← fetchConfig(SESSION_LIFETIME)
3     roles ← {}
4     passwords ← {}
5     sessions ← {}
6     owned ← {}
7     roleRequests ← {}
8
9  function setRole(username, role):
10     roles[username] ← role
11
12  function deleteUser(username)
13     roles[username] ← ⊥
14     passwords[username] ← ⊥
15     owned[username] ← ⊥
16     sessions[username] ← ⊥
17
18  function expireSession(sessionID, username):
19     if sessions[username] = sessionID:
20         sessions[username] ← ⊥
21
22  function saveRoleRequest(username, role):
23     requestID ← generateID()
24     roleRequests[requestID] ← (username, role)
25
26  function getPendingRoleRequest(requestID):
27     return roleRequests[requestID]
28
29  function getPendingRoleRequests():
30     pendingRequests ← {}
31     foreach (requestID, (username, role)) in roleRequests:
32         pendingRequests ← pendingRequests ∪ { (requestID, username, role) }
33     return pendingRequests
34
35  function deleteRoleRequest(requestID):
36     roleRequests[requestID] ← ⊥

```

Listing D.13: Storage of Role Requests